When Do Local Income Taxes Cause Residential Sorting?

Evidence from Pennsylvania's Piecewise EIT System

Matthew Hockert*

May 16, 2026

Abstract

Do local income taxes cause residential sorting across sub-county boundaries? I use LEHD Origin-Destination data linking each worker's residence block to their workplace block at annual frequency, covering approximately 2,500 Pennsylvania municipalities and 500 school districts over 2002–2019.

Pennsylvania's Earned Income Tax operates a piecewise max-rule: each worker's bill is $T = \max(\tau^{\text{home}}, \tau^{\text{nr,work}})$, where τ^{home} stacks the municipal and school-district resident rates and $\tau^{\text{nr,work}}$ is the work-location nonresident rate. A statutory rate change only moves the bill when the changed side of the max binds. I exploit this structure with a continuous stacked difference-in-differences design that decomposes each rate event by direction and binding-state transition, generating a within-design placebo: rate hikes whose post-event home stack stays below the nonresident rate (Inc-nobind) do not change the bill and should not move residents.

The preferred causal estimate is at the census block, where both rate-setting jurisdictions enter at their native geography. A 1 percentage-point municipal resident-rate

*University of Minnesota. Email: hocke043@umn.edu

increase in the binding cell reduces the log count of workers living in the block by -0.21 log points (approximately 19%), with a clean realized-dose pre-trend ($\beta_{\text{pre}} = -0.027$, $|t| < 1$). The Inc-nobind within-design placebo is statistically null on the same count outcome.

The municipality-level Inc-bind coefficient is larger (-0.40 log points on residents working in their own municipality) but pre-trended; I report it as policy-jurisdiction magnitude calibration rather than as the causal headline. School-district rate changes do not produce robust evidence of sorting and are treated as an aggregation diagnostic. Industry decomposition shows consumer-facing jobs fall in the binding cell while tradeable jobs are insulated, consistent with a local-services demand spillover from departing residents rather than direct nonresident-tax incidence.

1 Introduction

Do local income taxes cause residential sorting? The Tiebout hypothesis predicts that households "vote with their feet" by sorting across jurisdictions offering different bundles of taxes and public services (Tiebout, 1956), and a large literature documents tax-driven migration at the state level for high earners (Kleven et al., 2020; Moretti & Wilson, 2017; Young et al., 2016) and at sub-national borders for sales taxes (Agrawal, 2015). Yet sub-state income-tax sorting—the canonical Tiebout margin, where moving costs are small and tax boundaries are crossed by a short relocation—is essentially untested in the United States. Recent work on Indiana local income taxes (Yang & Heim, 2017) uses county-level IRS data and finds county-level migration responses, but cannot resolve the municipality-level variation that defines the actual tax boundary; the canonical U.S. sub-state place-of-taxation studies (Coomes & Hoyt, 2008; Rohlin et al., 2014) likewise operate at the state-border / county level. A literature survey by Agrawal (2026) explicitly identifies sub-state local-income-tax sorting as an open question, and Agrawal et al. (2022) characterize the empirical evidence as "thin."

The gap persists for two reasons. First, the jurisdictions that levy local income taxes are sub-county, so standard migration data sources do not observe the correct margins: the IRS Statistics of Income tracks county-to-county migration, and the American Community Survey measures place-level population without linking workers to employers. Second, local income-tax systems feature credit mechanisms, overlapping school-district taxes, and uneven enforcement that complicate the mapping from statutory rates to effective tax burden. A naive comparison of rates and population changes conflates tax incidence with amenity valuation, fiscal capacity, and enforcement quality.

I use Pennsylvania (2002–2019) and the LEHD Origin-Destination Employment Statistics (LODES), the only administrative data source that links each worker's residence to their workplace at annual frequency and at sub-municipal geography. Pennsylvania's Earned Income Tax (EIT) is unusual along two dimensions that the empirical design exploits. First, the EIT is set independently by two overlapping rate-setting jurisdictions—municipalities

and school districts. Each Pennsylvania address sits in both a municipality and a school district, and a resident's home-side tax bill is the sum of the municipality resident rate and the school-district rate. The paper's primary focus is the municipality resident rate; the school-district rate enters the empirical design but is treated as a secondary, aggregation-diagnostic channel because SD rate changes do not recover a robust sorting response. Second, the EIT operates a piecewise max-rule: each worker's bill is $T = \max(\tau^{\text{home}}, \tau^{\text{nr,work}})$, where $\tau^{\text{home}} = \tau^{\text{res,muni}} + \tau^{\text{SD}}$ is the sum of the municipal resident rate and the school-district resident rate at the worker's address (the home stack, and the term I use throughout the paper), and $\tau^{\text{nr,work}}$ is the nonresident rate at the work location. A statutory rate change on the home side only moves the worker's bill if the home stack is on the binding (i.e., higher) side of the max; the work-side nonresident rate is always paid by nonresidents, so the home-side rate change is the marginal margin only when the home stack exceeds it. Act 32 of 2008 (effective 2012) consolidated EIT collection into county-based Tax Collection Districts and mandated employer withholding at the higher of the two rates, which made the max-rule mechanically enforced and provides an enforcement-break heterogeneity test.

Section 3 formalizes the binding-state cell decomposition—six cells per rate-setter, partitioning each statutory rate event by direction (Increase/Decrease) and by the home-stack-vs-nonresident-rate binding transition—and lays out the worker-bill differencing function each cell implies.

At the block level, the preferred causal estimate, a 1 percentage-point municipal resident-rate increase in the Inc-bind cell reduces the log count of workers living in the block by -0.21 log points (approximately 19%); the Inc-nobind within-design placebo is null. The municipality-level analogue is larger (-0.40 log points on live-and-work residents, $\approx 33\%$ per pp) and is reported as policy-jurisdiction magnitude calibration. School-district rate changes do not produce robust evidence of sorting and are treated as an aggregation diagnostic. An industry decomposition shows consumer-facing jobs fall in the binding cell while tradeable jobs are insulated, consistent with a local-services demand spillover from departing

residents rather than direct nonresident-tax incidence on jobs. The paper's contribution is the binding-state decomposition: local income taxes cause residential sorting precisely when the institutional rule makes them mechanically change the worker's bill.

2 Institutional Background

Among U.S. states that authorize sub-state income taxation, Pennsylvania has the deepest cross-section of jurisdictions setting their own rate: approximately 2,500 municipalities and 500 school districts each levy a separate Earned Income Tax (EIT) on essentially all earned income earned in or by Pennsylvania residents (Agrawal, 2026; Bagchi, 2022). Unlike property taxes, which are paid by owners and capitalized into home values, the EIT falls directly on workers and varies across municipal boundaries within the same metropolitan area, creating sharp spatial discontinuities in tax burden that can influence where workers choose to live and work.

2.1 Pennsylvania Earned Income Tax

Pennsylvania's EIT is mandatory for all municipalities under the Local Tax Enabling Act of 1965 (Act 511; Pennsylvania General Assembly, 1965). It applies to all earned income with no minimum threshold; retirees are exempt. The standard combined cap is 1%, typically split as 0.5% municipal and 0.5% school district. Home Rule municipalities may exceed this cap (Pennsylvania General Assembly, 1972), and Philadelphia operates separately under the Sterling Act of 1932 (Pennsylvania General Assembly, 1932) at approximately 3.9%. For most non-Home-Rule boroughs and townships the EIT is the primary discretionary revenue lever, because property-tax millage limits and reassessment lags constrain the property side; quantifying that constraint would require a 2002–2019 municipality-level millage panel that is not currently available from DCED, and I leave the comparison to property-tax sorting as a separate exercise. Across 2002–2019, the panel contains 210 municipality resident-rate

increases, 119 municipality decreases, 611 school-district increases, and 52 school-district decreases; the binding-state breakdown of these events is in Table 6.

Each Pennsylvania municipality has three EIT rates relevant to a worker: the municipality resident rate $\tau^{\text{res,muni}}$ levied on residents of the municipality; the municipality nonresident rate $\tau^{\text{nr,muni}}$ levied on workers employed in the municipality who live elsewhere; and the school-district rate τ^{SD} levied on residents only (school districts cannot tax nonresidents). Each address is associated with a six-digit Political Subdivision (PSD) code identifying its (municipality, school district) pair. Workers report this code on the Residency Certification Form, and employers use it to determine withholding obligations.

Under Act 32 (effective January 1, 2012), a worker's total EIT bill is determined by the maximum of the home-side rate stack and the work-location nonresident rate, and the bill is split piecewise between home and work jurisdictions:

$$\left(\tau^{\text{home}},\ \tau^{\text{work}}\right) = \begin{cases} \left(\tau_j^{\text{res,muni}} + \tau_j^{\text{SD}} - \tau_i^{\text{nr,muni}},\ \tau_i^{\text{nr,muni}}\right) & \text{if } \tau_j^{\text{res,muni}} + \tau_j^{\text{SD}} > \tau_i^{\text{nr,muni}} \\ \left(0,\ \tau_i^{\text{nr,muni}}\right) & \text{if } \tau_j^{\text{res,muni}} + \tau_j^{\text{SD}} \leq \tau_i^{\text{nr,muni}} \end{cases} \tag{1}$$

where j is the worker's home municipality and i is the work municipality. The work jurisdiction always collects its full nonresident rate first; the home municipality and school district then split whatever portion of the worker's bill is left over after the work claim. If the work nonresident rate exceeds the home stack, the home jurisdictions get nothing. The worker's total bill is the higher of the two stacks:

$$\text{Bill}_{ji} = \max\left(\tau_j^{\text{res,muni}} + \tau_j^{\text{SD}},\ \tau_i^{\text{nr,muni}}\right). \tag{2}$$

Equation (1) implies which margin a worker can use to respond to a rate increase. When the home stack binds ($\tau_j^{\text{home}} > \tau_i^{\text{nr,muni}}$), the worker can lower the bill only by moving home to a lower-rate municipality or school district. When the work nonresident rate binds, the worker can lower the bill only by changing job to a lower-nonres-rate work location; moving

home without changing job has no effect. For workers who live and work in the same municipality, only the home stack applies—they pay $\tau_i^{\text{res,muni}} + \tau_i^{\text{SD}}$.

A worker who lives in municipality j and moves to a lower-rate municipality j' while keeping the same job at i saves at least the SD rate at j: the new bill is $\max(\tau_{j'}^{\text{home}}, \tau_i^{\text{nr}})$, which is at most the old bill τ_j^{home} minus τ_j^{SD}. Residence moves therefore do save tax under Pennsylvania's piecewise structure, because the home-side stack includes the SD rate, which the work jurisdiction cannot collect from nonresidents. The empirical design follows this logic directly: both the treatment $\Delta\tau^{\text{muni}}$ and the binding-state classification are built from the worker's home stack. The regressions therefore measure sorting in response to the same tax a mover would actually face.

Philadelphia is the one exception to the Act 32 framework. It operates under the Sterling Act of 1932: Philadelphia residents pay the city resident rate regardless of work location with no offset, and Philadelphia nonresidents pay the city nonresident rate, with no residual flowing to home jurisdictions in other Pennsylvania municipalities. My main estimates retain Philadelphia in the sample for completeness; results are robust to dropping it.

Act 32 of 2008 (Pennsylvania General Assembly, 2008) is central to identification in this paper. Before its 2012 effective date, EIT collection was radically decentralized: each of Pennsylvania's 2,500+ municipalities appointed its own tax officer. Nonresident EIT enforcement was weak because no centralized system linked employers, workers, and municipalities. A worker who moved out of a rate-raising municipality could, in practice, avoid paying the nonresident portion to the work municipality even though they legally owed it. The piecewise rule described in Equation 1 existed on paper but was not effectively enforced.

Act 32 transformed this system by consolidating collection into approximately 67 county-based Tax Collection Districts (TCDs), each with a single appointed collector. The law mandated standardized employer withholding at the higher of the employee's resident and work rate. This had two consequences. First, it made the EIT mechanically binding: rate changes now appear in every worker's paycheck through employer withholding, regardless of

whether the worker resides in the taxing municipality. Second, it eliminated the possibility of de facto avoidance of the work-location nonresident claim through weak enforcement. After Act 32, a worker could no longer effectively duck the nonresident portion owed to the work municipality.

Importantly, Act 32 did not eliminate all tax savings from residential relocation. The piecewise rule in Equation (1) continues to imply that a worker whose original home stack binds can reduce the bill by moving to a sufficiently lower-stack home jurisdiction: the new bill is $\max(\tau_{j'}^{\text{home}}, \tau_i^{\text{nr}})$, which falls below the old home stack τ_j^{home} whenever $\tau_{j'}^{\text{home}} < \tau_j^{\text{home}}$ and the original home stack was the binding side. Act 32 thus changed enforcement and made the max-rule mechanically binding through withholding; it did not eliminate the residential-relocation margin that the mechanism in this paper exploits. Bagchi (2022) documents the quantitative consequences of Act 32 directly: consolidated employer withholding increased PA EIT compliance, raised reported earned income in the rate-raising municipality, and shifted the de facto incidence toward workers. The present paper complements that work by isolating the residential-sorting response to the same withholding-enforcement regime, exploiting the cross-sectional variation in whether a given rate change actually moves the worker's bill (the binding-state decomposition); the pre/post-Act-32 split in Section 7 shows the binding-cell sorting response strengthens after Act 32, complementing the compliance-response finding.

School districts in Pennsylvania also levy an EIT under Act 511 (Pennsylvania General Assembly, 1965), typically at 0.5% (the other half of the standard 1% cap). Act 1 of 2006 (Pennsylvania General Assembly, 2006) (the Taxpayer Relief Act) additionally authorized school districts to exceed the 0.5% cap specifically to fund property tax relief. This generated a wave of school district rate increases in 2006–2011, creating rate variation that is largely independent of municipal fiscal decisions: the correlation between municipal and school district EIT rates across my panel is -0.12. Since every census tract sits in both a municipality and a school district, workers face two overlapping EIT rates that can change

independently, providing a natural decomposition of the residential sorting response into municipal and school district components.

3 Conceptual Framework: Binding-State Decomposition

The piecewise tax allocation in Equation (1) is the analytical engine of the empirical design. A worker's bill is $T = \max(\tau^{\text{home}}, \tau^{\text{nr,work}})$. A rate change only moves the bill when its side of the max binds. The implied change in the worker's bill from a single rate event is

$$\Delta T \;=\; \max(\tau^{\text{home}}_{\text{post}}, \tau^{\text{nr}}) - \max(\tau^{\text{home}}_{\text{pre}}, \tau^{\text{nr}}),$$

which depends jointly on the direction of the rate change and on which side of the max is binding pre- and post-event. This generates a sharp prediction: the same statutory rate change can produce a strong behavioral response, a moderate response, or no response at all depending on the binding-state transition.

3.1 Six Cells per Rate-Setter

Two definitions structure the decomposition. First, the piecewise rule $T = \max(\tau^{\text{home}}, \tau^{\text{nr,work}})$ has two sides: the home side $\tau^{\text{home}} = \tau^{\text{res,muni}} + \tau^{\text{SD}}$, set at the worker's residence, and the work side $\tau^{\text{nr,work}}$, set at the worker's work municipality. Whichever side is larger is the binding side and determines the bill. Second, three rate-setters adjust statutory rates independently across municipality-years: the municipality resident rate $\tau^{\text{res,muni}}$, the school-district rate τ^{SD}, and the municipality nonresident rate $\tau^{\text{nr,muni}}$. The decomposition is built one rate-setter at a time.

For each rate-setter τ, a cell classifies a municipality by (i) the direction of its first rate event in τ (Increase or Decrease) and (ii) the binding-state transition of its home stack

across that event: stayed-binding, flipped, or stayed-non-binding. The cross gives six cells plus a seventh never-changed control pool (Table 1). Each municipality therefore carries three cell labels—one for each of $\tau^{\text{res,muni}}$, τ^{SD}, and $\tau^{\text{nr,muni}}$—and the binding state for each is determined by where its home stack sits relative to its nonresident rate at the time of that event.

Cell	Direction	Pre-binding?	Post-binding?	Bill effect
Inc-bind	$\uparrow$	yes	yes	rises by full $\Delta\tau$
Inc-becomes	$\uparrow$	no	yes	rises by part of $\Delta\tau$
Inc-nobind	$\uparrow$	no	no	no change (placebo)
Dec-bind	$\downarrow$	yes	yes	falls by full $\Delta\tau$
Dec-becomes-non	$\downarrow$	yes	no	falls only to τ^{nr}
Dec-nobind	$\downarrow$	no	no	no change (placebo)
NC	none	—	—	control pool

Table 1: Binding-state cells per rate-setter. For Inc-becomes the bill rises only by the portion of $\Delta\tau$ above the nonresident rate; for Dec-becomes-non the bill falls only until the home stack hits the nonresident floor.

The fundamental testable prediction of the institutional rule is an ordered ranking of cell magnitudes: Inc-bind effects exceed Inc-becomes effects, which in turn exceed Inc-nobind effects, with the last close to zero in absolute value and serving as the within-design placebo. A salience-based argument cuts the other way on Inc-bind versus Inc-becomes: workers in the Inc-becomes cell experience a discrete switch from paying zero home-side tax to paying a positive amount, which could trigger an outsized behavioral response on the threshold-crossing margin. The data favor the institutional prediction at both estimating units. The central empirical comparison of the paper is between Inc-bind and Inc-nobind: if statutory rates matter mechanically through tax liability, the residential response should appear in Inc-bind and not in Inc-nobind. If Inc-nobind moves, the mechanism interpretation fails—the response is something other than the binding rule, because Inc-nobind events do not

change the worker's bill. This makes the design self-falsifying in a way that a conventional pooled Inc/Dec specification is not: in the pooled specification, every increase is a treated event, while here the institutional rule predicts that only a subset of statutory increases should generate a response.

The increase cells carry the identification because the decrease cells are thin: eleven Dec-bind, four Dec-becomes-non, fifteen Dec-nobind municipalities in total. The empirical results below focus on the three Inc cells, with decrease cells reported alongside for completeness.

3.2 Three Treatment Margins

Pennsylvania's institutional structure separates three rate variables that change independently across municipality-years. The municipality resident rate $\tau^{\text{res,muni}}$ directly affects the bill of in-municipality residents through the home-stack channel. The school district rate τ^{SD} also affects the home stack, but is set by school districts rather than municipalities, so the predicted mechanism is the same binding-rule channel as the municipality rate even though the SD's institutional decision-making differs from the municipality's. The municipality nonresident rate $\tau^{\text{nr,muni}}$ affects nonresident workers commuting in, since the piecewise rule means it moves the bill of workers from low-home-stack origins.

Two predictions follow. First, each rate's binding-cell decomposition should show the same ranking of Inc-bind above Inc-becomes above Inc-nobind. Second, channel separation should hold: the municipality resident rate moves residents (live-and-work, total residents); the municipality nonresident rate moves nonresident workers (in-commuters); and the SD rate's effect depends on whether SD changes induce similar residential responses.

4 Data

4.1 Data Sources

The analysis links two administrative systems that do not share common identifiers: LEHD/LODES residence–workplace records, keyed by Census block, and Pennsylvania EIT rate records, keyed by DCED municipality ID. The bridge is the TIGER 2020 block geometry, which nests blocks within county subdivisions (cousubs) by construction. A municipality-name crosswalk maps each DCED `municipality_id` to its corresponding COUSUBFP, matching 2,496 of 2,505 PA municipalities (99.6%). The remaining 9 are small boroughs that dissolved or merged during 2002–2019 and are dropped from the analysis. Appendix E reports the full pipeline including scraper specs, PSD mapping, fiscal-control construction, and auxiliary cross-state-inflow handling.

The residence–workplace data come from LEHD LODES Version 8, 2002–2019 (U.S. Census Bureau, 2024c): block-level origin–destination links derived from state Unemployment Insurance wage records and the Quarterly Census of Employment and Wages (Abowd et al., 2009). The Census Bureau applies synthetic noise infusion at the block level to protect respondent confidentiality, calibrated so that aggregation to tract and larger geographies approximates ground truth (Graham et al., n.d.). The sample period of 2002–2019 reflects LODES coverage beginning in 2002 and the need to drop 2020–2023, when COVID disrupted both commuting patterns and AFR fiscal data through ARPA transfers.

Municipal and school-district EIT rates come from the Pennsylvania Department of Community and Economic Development (DCED) MunStats portal (Pennsylvania Department of Community and Economic Development, 2024a), scraped annually. Pre-2012 nonresident rates, which DCED recorded only sporadically before Act 32 made employer withholding mandatory, are backfilled from the earliest observed post-2012 value per municipality. Two cleaning rules apply: DCED reports the flat-dollar Local Services Tax in the same field as the EIT in some years, so rates above 5% are treated as data-entry errors and dropped (the highest standard PA EIT is $\sim$1%; only Philadelphia exceeds 3%), and rate values are aligned to April 1 of each year to match LODES's reference date.

Municipal fiscal controls come from the PA DCED Statewide Municipal Annual Financial Report database (Pennsylvania Department of Community and Economic Development, 2024b). School-district fiscal controls come from the Census Annual Survey of School System Finances (F-33) (U.S. Census Bureau, 2024a); the NCES `LEAID`-to-DCED `sd_id` crosswalk covers all 500 Pennsylvania districts (see Appendix E). Two auxiliary sources support the pre-treatment balance and robustness exercises: FBI Uniform Crime Reports, Offenses Known and Clearances by Arrest, in the concatenated agency-year files compiled by Kaplan (2024), and the Census Bureau Building Permits Survey (U.S. Census Bureau, 2024b), matched to PA municipalities at 93.5% and 93.3% rates respectively.

4.2 Descriptive Statistics

The primary outcome variable for the headline causal analysis is the block-level log count of workers living in the block. The count directly measures the residential location of active workers and avoids the denominator fragility of block-level shares, which are mechanically sensitive at sparse geography because both the numerator and denominator move. At the municipality and school-district aggregations I also report resident-worker counts, workplace counts, in-commuter flows, and share outcomes (*share_live_in*). Share outcomes at coarser geographies retain the interpretation of earlier literature; at the block level I treat them as supporting evidence rather than as the primary estimand. Fiscal outcomes used as falsification placebos and pre-treatment balance variables (log EIT revenue, log total tax revenue, log total revenue, log police expenditure, log population, log market value of real estate, plus crime and building permits) are reported in Appendices D and E.7.

This section reports summary statistics and key distributions of the estimation panels. Table 2 reports the municipality-year panel; Table 3 reports the block-year panel that underlies the headline causal estimate. Table 4 reports the school-district aggregation. Table 5 summarizes the binding-state cell classification on the municipality resident-rate side, including the dose magnitude in each cell.

Table 2: Muni-year summary statistics (2002–2019)

Variable	N	Mean	SD	p25	p50	p75	Min	Max
Resident rate τ^M (%)	43,374	0.525	0.163	0.500	0.500	0.500	0.000	4.331
Nonresident rate $\tau^{nr,M}$ (%)	34,080	0.883	0.237	1.000	1.000	1.000	0.000	3.820
SD rate τ^{SD} (%)	43,374	0.622	0.289	0.500	0.500	0.500	0.000	2.050
Home stack $\tau^M + \tau^{SD}$ (%)	43,374	1.147	0.316	1.000	1.000	1.000	0.000	4.331
log live-and-work	40,748	3.638	2.011	2.079	3.714	5.088	0.000	12.877
log total jobs	45,291	5.819	2.006	4.489	5.883	7.234	0.000	13.353
Share live-and-work in muni	45,291	0.095	0.081	0.037	0.080	0.133	0.000	1.000
Total jobs	45,291	2076.691	13031.730	89.000	359.000	1385.500	1.000	629698.000

Notes: Muni-year level; one observation per (municipality, year). Rates in percentage points. $\tau^{nr,M}$ has reporting gaps before Act 32; the N for that row is smaller than the full panel.

Table 3: Block-year summary statistics (2002–2019)

Variable	N	Mean	SD	p25	p50	p75	Min	Max
Home count (workers living)	5,030,958	18.822	31.677	4.000	10.000	22.000	0.000	1572.000
log home count	4,795,213	2.293	1.198	1.386	2.398	3.135	0.000	7.360
Work count (jobs in block)	5,030,958	18.826	138.957	0.000	0.000	4.000	0.000	20945.000
log work count	1,894,432	2.256	1.663	1.099	2.079	3.367	0.000	9.950
Share live-and-work in muni	4,795,213	0.154	0.222	0.000	0.067	0.203	0.000	1.000
Share work outside muni	4,795,213	0.846	0.222	0.797	0.933	1.000	0.000	1.000
Resident rate τ^M (%)	4,782,928	0.762	0.799	0.500	0.500	0.500	0.000	4.331
SD rate τ^{SD} (%)	4,782,928	0.607	0.376	0.500	0.500	0.500	0.000	2.050
Nonresident rate $\tau^{nr,M}$ (%)	4,181,214	1.063	0.657	1.000	1.000	1.000	0.000	3.820
Blocks per muni	2,561	120.922	384.585	41.000	73.000	132.000	1.000	16550.000

Notes: Block-year level; one observation per (census block, year). Panel covers 309,680 blocks in 2,561 municipalities.

Table 4: School-district-year summary statistics (2002–2019)

Variable	N	Mean	SD	p25	p50	p75	Min	Max
SD rate τ^{SD} (%)	49,226	0.610	0.278	0.500	0.500	0.500	0.000	2.050
Resident rate τ^{M} (%) (muni-year)	49,226	0.523	0.162	0.500	0.500	0.500	0.000	4.564
Nonres rate $\tau^{nr,M}$ (%) (muni-year)	38,536	0.881	0.237	1.000	1.000	1.000	0.000	3.967
share_state_rev_sd	9,278	0.416	0.166	0.279	0.414	0.555	0.000	0.821
share_local_rev_sd	9,278	0.536	0.186	0.376	0.536	0.691	0.103	1.000
share_federal_rev_sd	9,278	0.048	0.033	0.024	0.040	0.062	0.000	0.247
share_property_tax_sd	9,278	0.420	0.176	0.270	0.407	0.562	0.000	0.853
share_instructional_sd	9,278	0.522	0.070	0.492	0.536	0.567	0.000	1.000
share_support_sd	9,278	0.279	0.039	0.259	0.284	0.305	0.000	0.438

Notes: SD-year level. F-33 share controls are constructed from Census F-33 revenue/expenditure totals.

Table 5: Cell counts: muni resident-rate side, 2002–2019

Cell	Munis	Muni-years	Mean dose (pp)	Max dose (pp)
NA	1	18	0.000	0.000
NC	1,793	32,266	0.000	0.000
Inc-bind	40	720	0.191	1.500
Inc-becomes	58	1,044	0.339	4.331
Inc-nobind	26	468	0.273	0.975
Dec-bind	11	198	-0.104	0.350
Dec-becomes-non	4	72	-0.217	0.000
Dec-nobind	15	270	-0.283	0.700

Notes: Cells classify each muni's first resident-rate event by direction (Inc/Dec) and the binding-state transition relative to the muni's nonresident rate. NC is the never-changed control pool. Dose is the cumulative $\Delta\tau$ at event-time $k = +5$.

Figure 1 plots the empirical CDF of the three rate variables that enter the piecewise max-rule: the municipality resident rate τ^{M}, the municipality nonresident rate $\tau^{nr,M}$, and

the home stack $\tau^M + \tau^{SD}$. Under Act 511 of 1965 (Pennsylvania General Assembly, 1965), the combined home-side EIT is capped at 1% for ordinary boroughs and townships and is conventionally split 0.5% municipal plus 0.5% school district; the vast majority of Pennsylvania municipalities sit at this statutory default and never adjust. Approximately 121 of the ∼2,500 municipalities change their resident rate at any point in the sample, and most school districts likewise hold at 0.5%. The ECDF makes this concentration visible: roughly 95% of municipality-years sit at or below the 0.5% resident-rate cap, and the home-stack CDF reaches $\sim$ 78% by the 1% combined cap. The long right tails correspond to two policy carve-outs: Act 47 distressed cities (Pennsylvania General Assembly, 1987) (Pittsburgh-area municipalities, Wilkes-Barre, Hazleton, Johnstown, and others) that exceeded the municipality half of the 1% cap, and Act 1 school districts (Pennsylvania General Assembly, 2006) that converted property-tax authority into EIT authority. Philadelphia operates under the Sterling Act of 1932 (Pennsylvania General Assembly, 1932) rather than Act 511 and accordingly sits at the far right of the resident-rate distribution. These tail municipalities and school districts generate the identifying variation that the binding-state decomposition then splits by whether each statutory rate change moved the worker's bill.

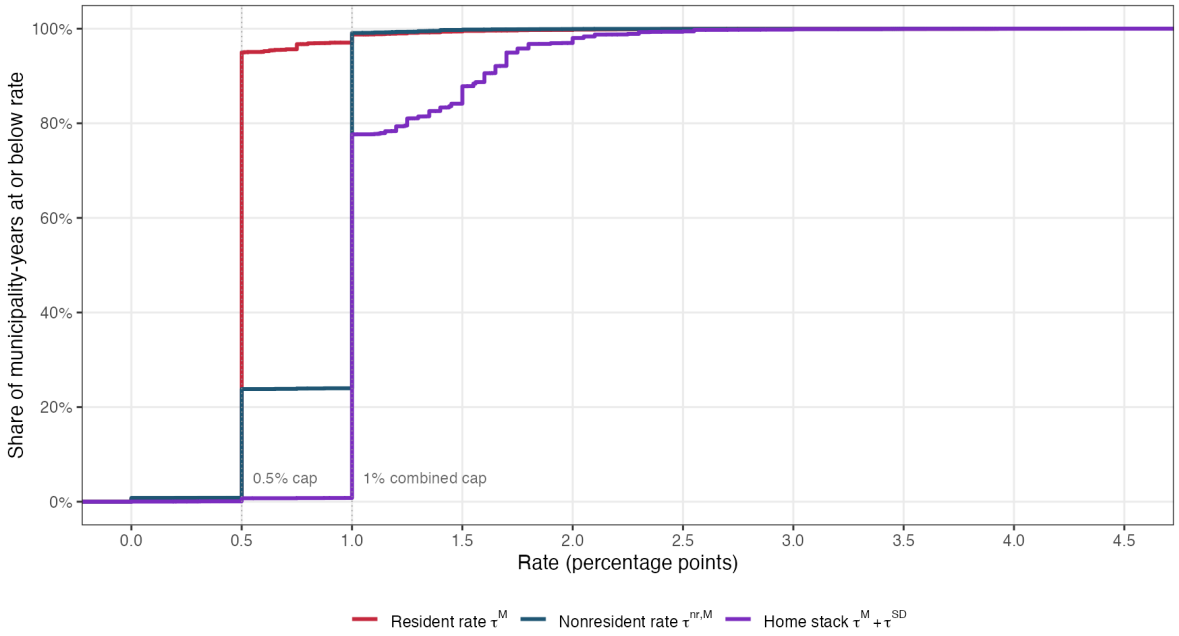


Figure 1: Empirical CDF of Pennsylvania EIT rates, 2002–2019. Each step shows the share of municipality-years at or below a given rate, separately for the municipality resident rate τ^{M} (red), the municipality nonresident rate $\tau^{nr,M}$ (blue), and the home stack $\tau^{M} + \tau^{SD}$ (purple). Dotted reference lines mark the 0.5% municipality cap and the 1% combined home-side cap under Act 511 of 1965. Approximately 95% of municipality-years sit at the 0.5% resident-rate cap and $\sim 78\%$ at the 1% combined home-stack cap; the long right tails are Act 47 distressed cities, Home-Rule cities, and Act 1 school districts.

Figure 2 reports the timing of municipal resident-rate changes across the panel. Increases outnumber decreases nearly two-to-one overall (210 increases vs. 119 decreases over 2002–2019), and the gap widens further after Act 32 took effect in 2012: roughly 80% of post-2012 rate-change events are increases, consistent with consolidated withholding raising the political cost of cutting a rate that is now mechanically enforced.

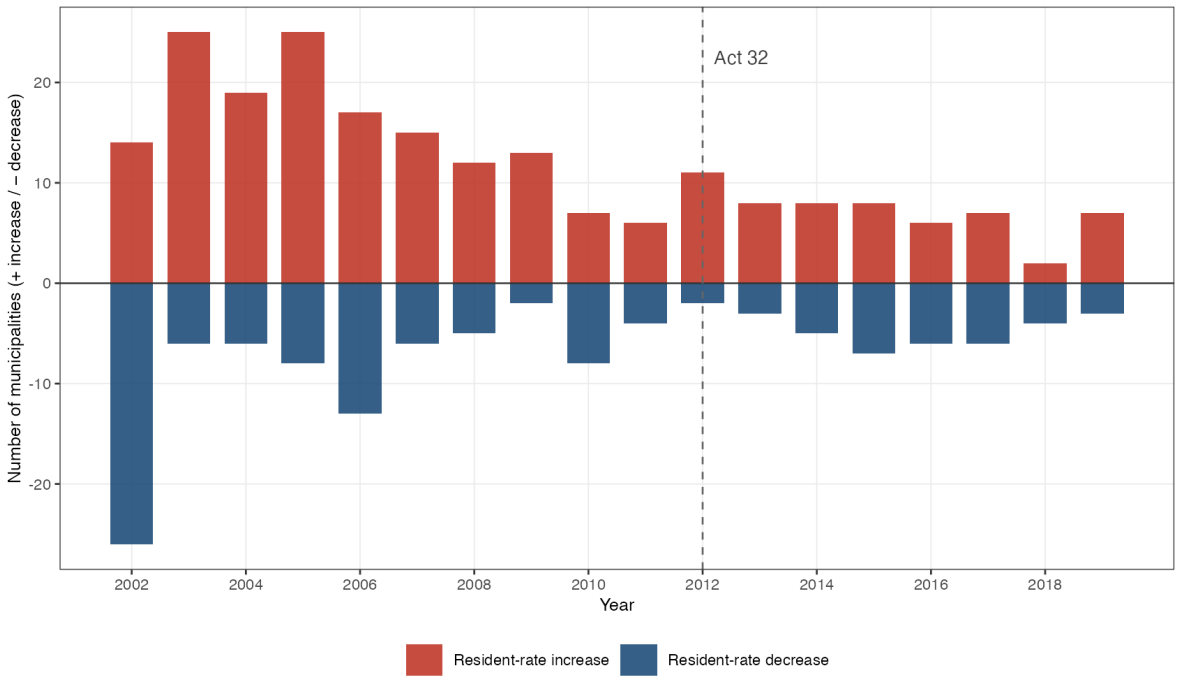


Figure 2: Number of PA municipalities changing their resident rate, by year. Increases shown above zero; decreases shown below. Dashed line: 2012, when Act 32 consolidated withholding.

Figure 3 documents the institutional break at Act 32. Before 2012, fewer than 60% of PA municipalities reported a nonresident rate to DCED; after 2012, when consolidated withholding made the max-rule mechanically binding, coverage rises sharply. This is the pre-/post-Act-32 enforcement break used in the heterogeneity test of Section 7.2.

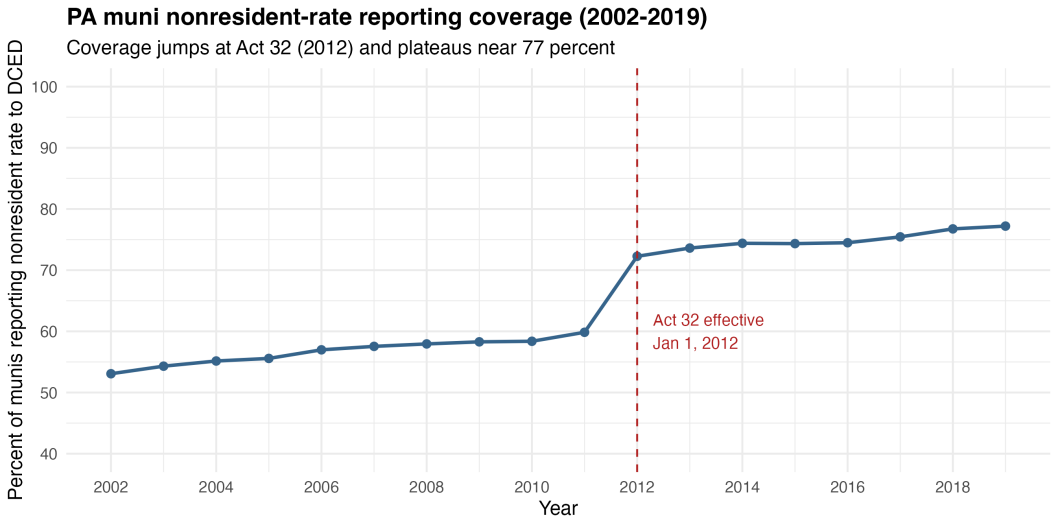


Figure 3: Share of PA municipalities reporting a nonresident rate to DCED, by year. The break at 2012 reflects the Act 32 requirement that employers withhold at the higher of the home resident and work nonresident rates.

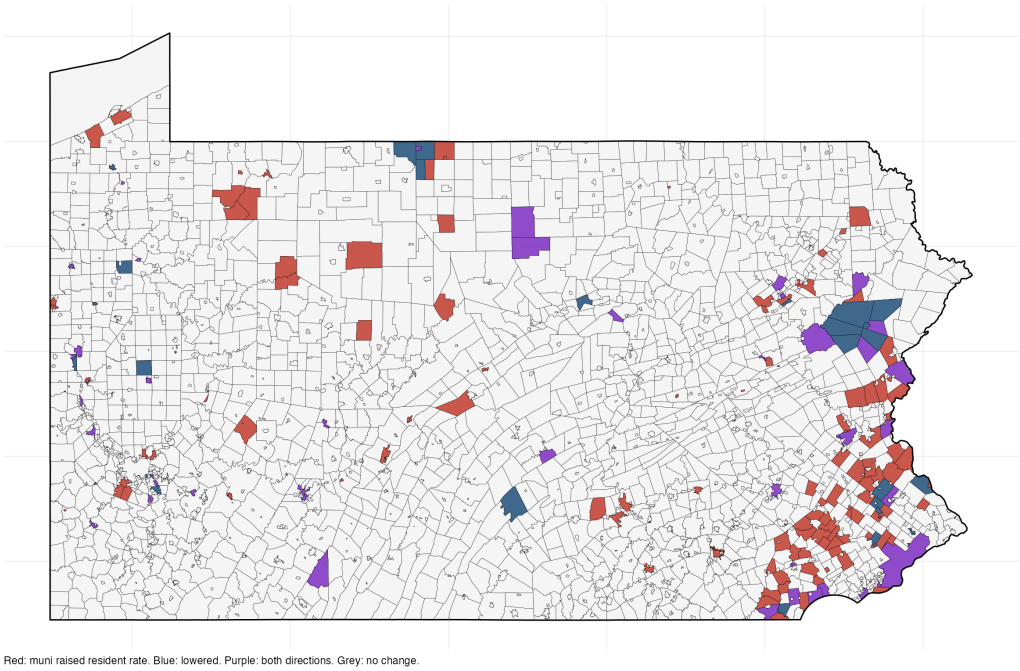


Figure 4: Pennsylvania EIT rate changes, 2000–2019. Red = increased; blue = decreased; purple = both; grey = no change.

4.2.1 Cell counts and stacked-design support

Before turning to the empirical strategy, this subsection documents the sample composition of the binding-state decomposition. The purpose is twofold: to verify that each cell in Section 5 has enough municipalities to identify a per-pp coefficient, and to characterize which types of Pennsylvania municipalities populate the treated cells—a within-state heterogeneity that matters because PA legal-organization classes (city, borough, township) have different revenue authorities and fiscal-capacity profiles.

Table 6 reports the number of unique municipalities in each binding-state cell on each rate-setter after the NC-purity filter described in Section 5.

Table 6: Cell composition by rate-setter (drop sample, post-NC-fix)

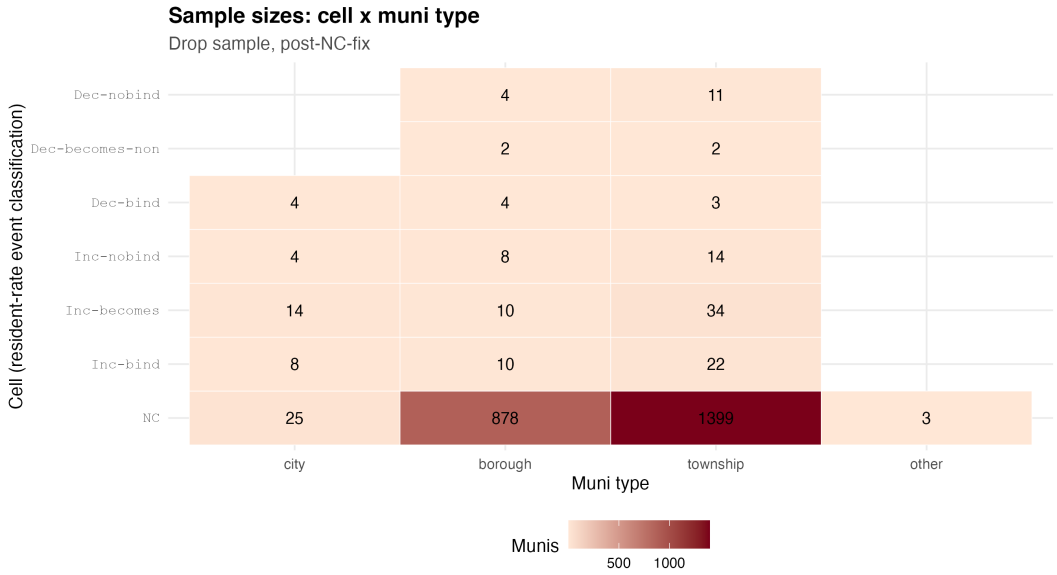
Side	Inc-bind	Inc-becomes	Inc-nobind	Dec-bind	Dec-becomes-non	Dec-nobind	NC
Muni resident (τ^{muni})	40	58	26	11	4	15	2,406
School district (τ^{SD})	75	197	14	4	10	—	2,253
Muni nonresident (τ^{nr})	0	0	171	0	5	109	2,276

Notes: Cells classified at each municipality's first observed rate event on the relevant rate-setter. Munis with rate changes whose first-event binding state cannot be determined (NA pre-event nonresident rate) are dropped (not pooled with NC). On the nonresident side, Inc-bind and Inc-becomes cells are empty in PA because the municipality's nonresident rate almost never exceeds the municipality's own home stack at the municipality-self-pair level.

The cell counts also vary substantially by Pennsylvania municipal classification. Under the Pennsylvania municipal code, every general local-government unit is incorporated as one of three types: a city (governed by the City Code; subdivided by population into First, Second, Second A, or Third class, with Pittsburgh as the sole Second-class city), a borough (governed by the Borough Code, 53 P.S. § 45101 et seq.), or a township (First-class under 53 P.S. § 55101 et seq. or Second-class under 53 P.S. § 65101 et seq.). The classes differ in revenue authority—cities have broader tax-imposition authority and are more likely to operate at or near the statutory caps—and in service responsibilities, which loads onto fiscal-

capacity heterogeneity within the treated cohort. Figure 5 disaggregates the cell composition by class. The 40 Inc-bind municipalities comprise 8 cities, 10 boroughs, and 22 townships; the municipality-type heterogeneity analysis in Appendix B.2 shows the headline magnitude is concentrated in the 8 post-industrial cities. Decrease cells are extremely thin: 11 Dec-bind, 4 Dec-becomes-non, 15 Dec-nobind units in total, with no cities in Dec-becomes-non or Dec-nobind. The binding-rule symmetry test between Inc and Dec is therefore not feasible under conventional standard-error inference.

Figure 5: Sample sizes by binding-state cell and municipality type (drop sample, post-NC-fix).

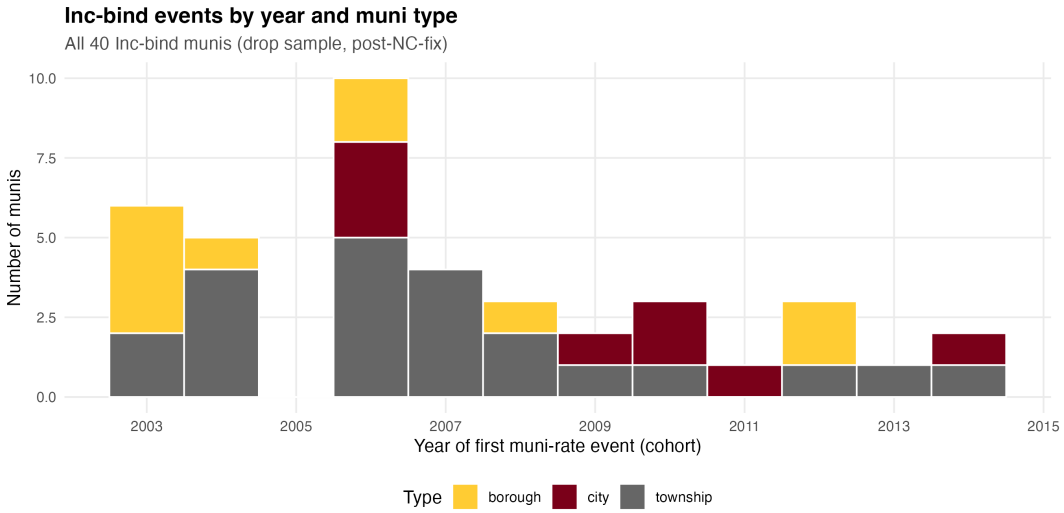


Notes: Each cell reports the count of unique municipalities in the (binding-state cell × municipality type) bin. NC dominates by sample size (878 borough + 1,399 township never-changers); treated cells are small per type.

The stacked design's identifying support comes from the joint distribution of cohort timing and event-time coverage. Figure 6 reports the distribution of municipality-rate first-event years for the 40 Inc-bind municipalities. The cohort distribution is bimodal: an early spike in 2003–2004 (boroughs and townships responding to baseline EIT setup) and a 2006 spike of 10 units (Act 1 of 2006 enabled SD-rate increases that mechanically flipped some municipalities into Inc-bind via the joint $\tau^{\text{muni}} + \tau^{\text{SD}}$ stack), with a long tail through 2014.

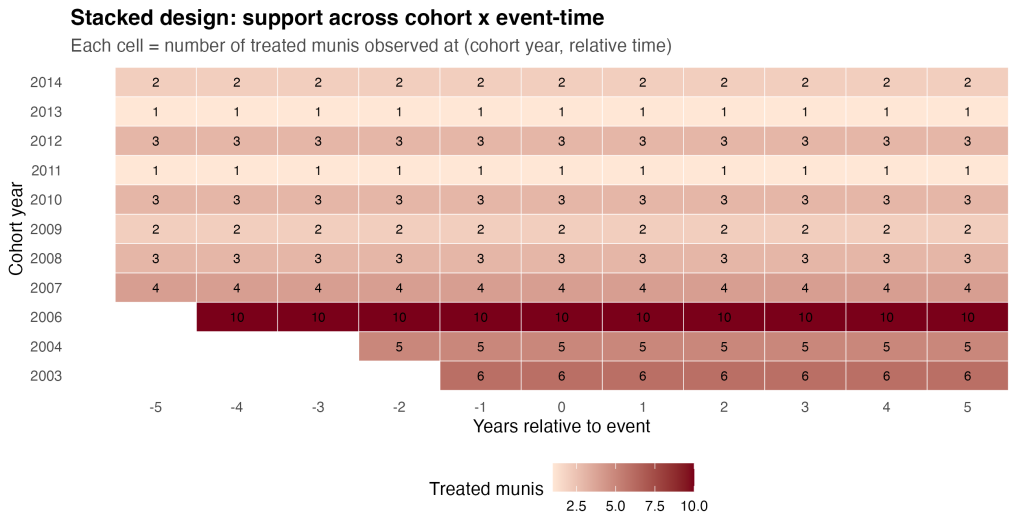
Figure 7 shows the resulting stacked-design support: the 2006 cohort dominates the post-event window with 10 treated municipalities, and most other cohorts contribute 1–4 treated municipalities to the stack. The thin tail cohorts limit the effective sample size for late-event-time coefficients in each cohort-specific stack.

Figure 6: Inc-bind municipality-rate first-event timing, by municipality type.



Notes: Counts of Inc-bind municipalities by year of first municipality-rate event. The 2006 spike reflects Act 1 of 2006 (school-district EIT authorization) flipping some municipalities into the binding cell via the joint home-stack mechanism.

Figure 7: Stacked design support: treated municipality count by cohort year × relative event-time.



Notes: Each cell reports the number of Inc-bind treated municipalities observed at the (cohort year, relative event-time k) combination. For example, the cell at row $c = 2006$ and column $k = +3$ shows the count of Inc-bind municipalities whose first-event year was 2006, observed in calendar year 2009 (three years after their rate change took effect). Pre-event cells with calendar year before 2002 are empty because LODES coverage begins in 2002, so cohorts at $c = 2003$ and $c = 2004$ have truncated pre-event windows.

5 Empirical Strategy

The empirical framework is a continuous stacked difference-in-differences design with binding-state cell decomposition, estimated on a panel spanning 2002–2019 at three units of observation: municipality, school district, and census block. Pre-2002 data are not consistently available in LODES; post-2019 years are excluded because COVID disrupted commuting patterns and ARPA fiscal transfers distorted municipal budgets. Each estimating unit plays a different role. The block-level estimate is the preferred causal specification because both rate-setting jurisdictions enter the regression at their native geography and the realized-dose pre-trend on the binding cell is clean. The municipality-level estimate is retained as policy-jurisdiction magnitude calibration, and the school-district aggregation is used as a

diagnostic. The first-difference dose-response specification estimates the post-event effect of a one percentage-point change in the relevant tax rate; a separate realized-dose event-study diagnostic tests for differential pre-event trajectories across units of varying eventual treatment intensity.

The resident-side specification includes two simultaneous rate-setter treatments: $\Delta\tau_{it}^{\text{muni}}$, the cumulative change in the municipality resident rate from baseline, and $\Delta\tau_{it}^{\text{SD}}$, the cumulative change in the school-district rate. Baseline for each is the rate at year $c_i^s - 1$, where c_i^s is unit i's first event year on rate-setter $s \in \{\text{municipality, SD}\}$, and cumulative dose accumulates with subsequent rate changes within i. The work-side nonresident rate τ^{nr} enters the binding-state classification of cells but is not itself a treatment in the headline regression: I focus on the resident-side mechanism and defer the work-side behavioral response to a separate analysis.

Each rate-change event is classified by direction $\in \{\text{Inc, Dec}\}$ and binding-state transition $\in \{\text{bind, becomes, nobind}\}$ to yield the six non-NC cells introduced in Section 3.1. For each unit i and each rate-setter s, I locate the first rate event year c_i^s and assign i to one of the six cells based on the direction of the change and the binding state of the home stack relative to $\tau^{\text{nr,muni}}$ in the pre- and post-event years. The NC reference category is restricted strictly to units whose rate τ^s never changes over 2002–2019. Units with rate changes whose first event has an undefined binding state—typically because the municipality's nonresident rate is missing in the pre-event year due to DCED reporting gaps before Act 32—are categorized as "Uncertain" and dropped: pooling them with NC would contaminate the control trend with treated-but-unclassified units, while reporting them in their own cell would mix true binding states with measurement gaps. Municipality-level regressions further impose a minimum employment threshold of 50 total jobs. The threshold protects the share-based outcomes (*share_live_in*, *share_work_outside*) from mechanical measurement noise: at N workers, a single worker moving changes the share by $1/N$, so at $N = 10$ a one-worker move changes the outcome by 10 percentage points, dominating the per-pp treatment magnitude. At $N = 50$

the one-worker noise floor is 2%, roughly an order of magnitude below the headline coefficient. The threshold also matches the LODES synthetic-noise-infusion design: the Census Bureau's noise calibration is validated at the tract level and above (Graham et al., n.d.), which in Pennsylvania typically contains at least 50 workers. Choosing exactly 50 rather than 20 or 100 is a judgment call; sensitivity analysis in Section 7 confirms the headline Inc-bind coefficient is stable across thresholds from 10 to 500. Imposing the threshold drops 185 of 2,559 PA municipalities (7.2%), almost all of which are sparsely populated rural townships in the NC pool; the treated cohort is essentially unaffected.

Outcomes are run in parallel across counts and shares. The municipality-level workplace and residence outcomes are the logs of `live_and_work` (residents employed in the municipality), `in_commuters` (workers in the municipality from elsewhere), `out_commuters`, `total_jobs`, and `total_residents`, plus the shares of workers living in their own municipality, of workers who are nonresidents, and of residents working outside the municipality. The school-district aggregation uses the same outcomes computed by summing block-level LODES counts up to the SD: each block has a unique cousub assignment (via the TIGER 2020 block-to-cousub crosswalk), each cousub has a primary SD assignment (the modal `sd_id` in the DCED rate panel; 97.6% of munis are single-SD), and the SD-level outcomes are block sums via this two-step block $\rightarrow$ cousub $\rightarrow$ SD path. At the block level, outcomes are the log count of workers living in the block, the log count of workers working in the block, the share of block residents working in their own municipality, and the share working outside; geometry and aggregation details are in Section 5.3.

Fiscal controls are share-based and time-varying, with the exact set keyed to the estimating unit. The shares are constructed as numerator-over-total ratios so that they absorb compositional changes in the municipality's or district's fiscal posture without conflating with the level of total spending.

Municipality-level regressions use six share controls from the DCED Annual Financial Report. Three are expenditure shares: `share_police` (police-services line / total expendi-

ture), share_streets (streets and highways), and share_debt_svc (debt service: principal and interest payments on outstanding general-obligation and revenue debt). Three are revenue shares: share_eit (Earned Income Tax revenue collected / total revenue), share_re (real-estate property tax revenue / total revenue), and share_intergov (intergovernmental revenue, which is grant and pass-through revenue from state and federal governments—e.g., Liquid Fuels Tax distributions, Community Development Block Grants, ARPA transfers—divided by total revenue).

School-district regressions use eight share controls from the Census Annual Survey of School System Finances (F-33) (U.S. Census Bureau, 2024a). Four are revenue shares: share_state_rev_sd (state aid: basic education subsidy, special-education subsidy, transportation reimbursement, etc.), share_local_rev_sd (locally raised revenue: property tax, EIT, occupation tax, public-utility taxes), share_federal_rev_sd (federal Title I, IDEA, school-lunch, and other direct or pass-through federal aid), and share_property_tax_sd (real-estate property-tax revenue as a share of total revenue, which is a subset of share_local_rev_sd and isolates the SD's primary local-tax lever). Four are expenditure shares: share_instructional_sd (teacher salaries, classroom supplies, instructional staff), share_support_sd (non-instructional operating costs: administration, transportation, maintenance, food service, student services), share_capital_sd (capital outlay—construction, renovation, equipment), and share_debt_int_sd (interest payments on the district's outstanding debt; principal appears as a capital outlay rather than as a current expense).

Block-level regressions use both sets of controls simultaneously, since each block is exposed to both rate-setters' fiscal posture. Figure 8 plots the empirical distribution of each share variable pooled across municipality-years (municipality shares) or SD-years (SD shares); cross-unit variation is broad on every dimension, so the share controls absorb meaningful fiscal heterogeneity without collinearity with the rate-treatment variable.

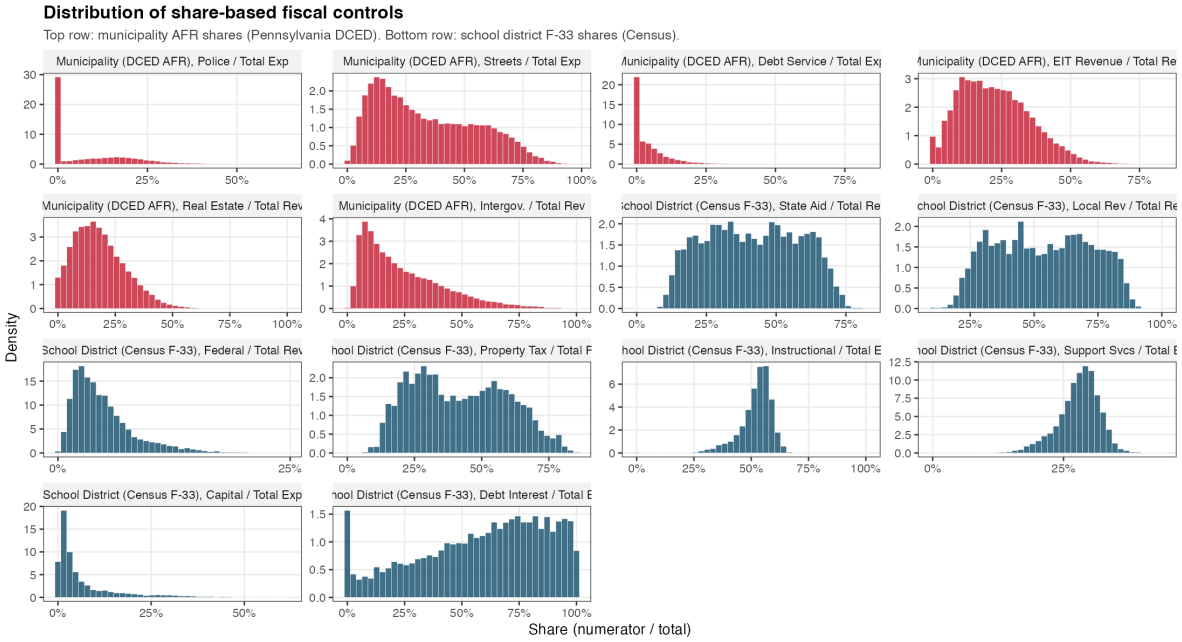


Figure 8: Distribution of share-based fiscal controls across the Pennsylvania panel, 2002–2019. Top six panels: municipality AFR shares (DCED). Lower eight panels: school-district F-33 shares (Census). Each panel pools all municipality-years (top) or SD-years (bottom) over 2002–2019.

5.1 Pooled Aggregate Specification

Before the binding-state decomposition I report an aggregate specification that pools all six binding-state cells per rate-setter into two direction tags (Inc for any rate increase, Dec for any rate decrease). The pooled regression is the literature-comparable baseline: it answers "what is the average per-pp effect of a municipality-rate hike on residents, ignoring binding state?" This is closest in form to the population-elasticity work of Yang and Heim (2017) on Indiana local income taxes. The binding-state decomposition that follows in Equation (4) then explains which sub-population of the treated municipality cohort drives the pooled coefficient.

For each unit-time observation (i, t) on the stacked panel, define the per-pp pooled spec

$$\begin{aligned} Y_{\tilde{i}\tilde{t}} \;=\; \alpha_{\tilde{i}} \;+\; \gamma_{\tilde{t}} &+ \beta^{M}_{\text{Inc}}\,\Delta\tau^{M}_{it}\,\mathbf{1}\{d^{M}_{i}=\text{Inc}\} \;+\; \beta^{M}_{\text{Dec}}\,\Delta\tau^{M}_{it}\,\mathbf{1}\{d^{M}_{i}=\text{Dec}\} \\ &+ \beta^{SD}_{\text{Inc}}\,\Delta\tau^{SD}_{it}\,\mathbf{1}\{d^{SD}_{i}=\text{Inc}\} \;+\; \beta^{SD}_{\text{Dec}}\,\Delta\tau^{SD}_{it}\,\mathbf{1}\{d^{SD}_{i}=\text{Dec}\} \\ &+ \mathbf{X}'_{it}\boldsymbol{\phi} \;+\; \varepsilon_{\tilde{i}\tilde{t}}, \end{aligned} \tag{3}$$

where $d^{s}_{i} \in \{\text{NC}, \text{Inc}, \text{Dec}\}$ is the direction tag on rate-setter $s \in \{M, SD\}$ (NC is the omitted reference), $\Delta\tau^{s}_{it}$ is the cumulative dose defined in Section 5.2, and the fixed effects and clustering follow the rule in Section 5.3: unit-by-stack and year-by-stack fixed effects, with clustering at the policy-jurisdiction level.

The pooled coefficient β^{M}_{Inc} is a weighted average of the six binding-state cell coefficients for the municipality rate-setter (Inc-bind, Inc-becomes, Inc-nobind; with equivalent decomposition for Dec). It is interpretable as the elasticity a prior-literature specification would recover. The decomposition in Section 6.2.1 then reveals that the pooled coefficient is dominated by the Inc-bind cell at the block level and attenuated toward zero in the Inc-nobind placebo cell—direct evidence that the institutional binding rule, not the statutory rate per se, drives the response.

I also report a binary analog of Equation (3) in which the four per-pp interactions are replaced by four post-treatment indicators $\{\text{everInc}^{M}_{it}, \text{everDec}^{M}_{it}, \text{everInc}^{SD}_{it}, \text{everDec}^{SD}_{it}\}$, each equal to 1 in the event year and after for the relevant direction on the relevant rate-setter. The binary coefficients are comparable to the event-study coefficients in Bagchi (2022); both versions appear together in Section 6.2.1.

5.2 Binding-State Decomposition: Stacked Continuous DiD

The binding-state decomposition extends the pooled spec by splitting each direction tag into its three binding-state cells, and is the main estimator of the paper. The construction is stacked continuous difference-in-differences. For each rate-setter $s \in \{\text{muni}, \text{SD}\}$ and

each cohort year c at which any municipality's first event on rate-setter s occurs, I build a sub-experiment $\mathcal{S}_c^s$ consisting of two types of units. The treated set is the municipalities with first-event year $c_i^s = c$ and a non-NC cell on rate-setter s; the control set is all NC^s municipalities (whose rate τ^s never changes over 2002–2019). The window covers calendar years $t \in [c-K,\, c+K]$ with $K = 5$. The choice $K = 5$ is conventional in the staggered-DiD literature (e.g., Cengiz et al., 2019) and is calibrated against the panel: $K = 5$ yields a 5-pre, 5-post window that fits inside 2002–2019 with full coverage for cohorts $c \in [2007, 2014]$ (i.e., the bulk of the treated mass) and gives the realized-dose ES four pre-event leads to test pre-trends. A shorter K would lose post-event coefficients before the cumulative dose has fully realized; a longer K would truncate the pre-period on early cohorts and the post-period on late cohorts. Sensitivity to $K \in \{3, 4, 6, 7\}$ leaves the headline coefficient stable to two decimal places (not reported separately).

A concrete example helps fix ideas. Wilkes-Barre is an Act 47 distressed city with a municipality resident-rate increase from 1.5% to 2.0% in 2009. Its work-side nonresident rate is 1.0%, so the home stack of $(1.5 + 0.5) = 2.0\%$ pre-event already exceeds the nonresident rate; the home stack continues to bind post-event at $(2.0 + 0.5) = 2.5\% > 1.0\%$. Wilkes-Barre is therefore an Inc-bind municipality with $c^M_{\text{WB}} = 2009$. The sub-experiment $\mathcal{S}^M_{2009}$ runs on the window $[2004, 2014]$. Its treated set includes Wilkes-Barre and any other Inc-bind / Inc-becomes / Inc-nobind / Dec-cell municipality whose first municipality event was 2009 ($\approx$ 10 munis); its control set is the $\sim$2,400 NC munis whose resident rate stayed flat across 2002–2019. Each NC municipality contributes 11 municipality-year observations ($[2004, 2014]$) to this sub-experiment. The dose variable $\Delta\tau^{\text{muni}}_{it} = \tau^{\text{muni}}_{it} - \tau^{\text{muni}}_{i,2008}$ is zero for $t \leq 2008$ and equals 0.005 for $t \geq 2009$ for Wilkes-Barre (and analogously for other 2009 cohort munis based on their own pre-event baseline). Relative event time is $\ell_{it} = t - 2009$, so $\ell = -5, -4, \ldots, 0, 1, \ldots, 5$.

Each municipality-year observation in $\mathcal{S}_c^s$ carries a stack indicator $s_i = c$, relative event time $\ell_{it} = t - c$, and stack-by-municipality and stack-by-year identifiers $\tilde{i} = (s_i, i)$ and $\tilde{t} =$

(s_i, t) that allow each unit to appear in the stacked panel multiple times under different stack identifiers (this is how NC controls re-appear across cohorts without contaminating cohort-specific identification). Stacking across all cohort years yields a single panel in which each NC municipality appears once per cohort and each treated municipality appears in exactly one sub-experiment—its first-event cohort. Within each sub-experiment, controls are clean (no treated municipality from another cohort serves as control), which is the property that fixes the negative-weighting problem in vanilla TWFE under staggered treatment (Goodman-Bacon, 2021).

Figure 9 sketches the construction. Each row is one example municipality. The colored bar is its $[c - K, c + K]$ window; the open circle marks its first-event year c; the cell label on the right is the binding-state classification. The six treated examples are one per binding-state cell. The bottom row is the never-changed (NC) control band: every NC municipality contributes observations to every treated cohort's sub-experiment, in calendar years matching that cohort's window. Early cohorts (e.g., the 2006 Inc-becomes example) have a slightly truncated pre-period because their $c - K$ year falls just before the panel's 2002 start; the stacked panel handles this by simply omitting the unobserved pre-event years.

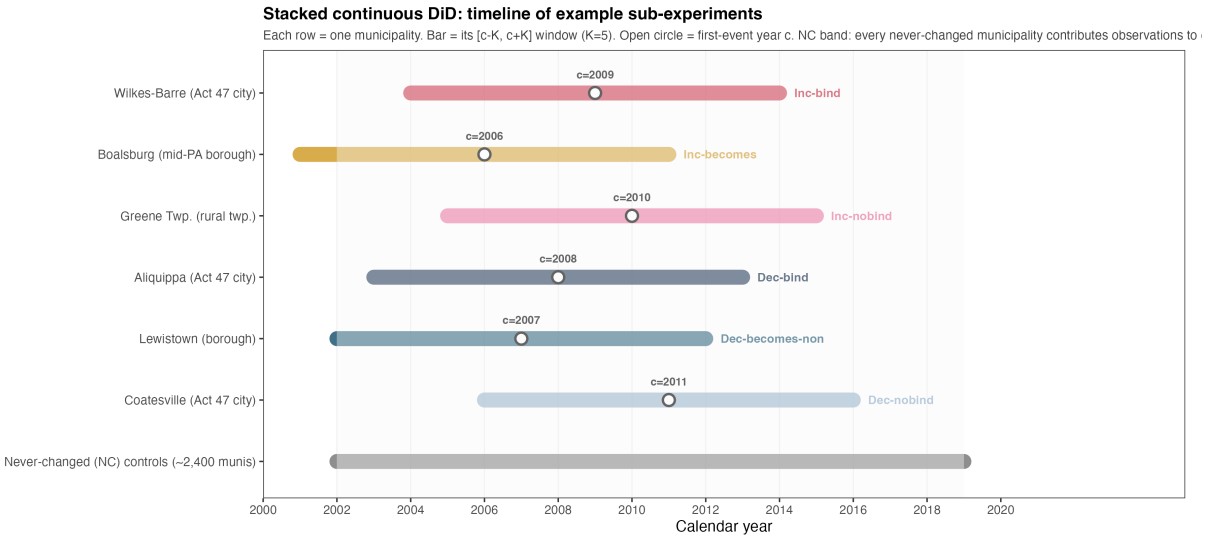


Figure 9: Stacked continuous DiD: timeline of example sub-experiments. Each row is a municipality; the colored bar is the $[c-K,\ c+K]$ window with $K = 5$; the open circle marks the first-event year c; the cell label gives the binding-state classification. The bottom (grey) row is the never-changed (NC) control pool; in the stacked panel, every NC municipality contributes observations to every cohort's window. Municipality names are illustrative.

The stacked-DiD construction is in the lineage of Cengiz et al. (2019) (138 minimum-wage events, stacked event-by-event with continuous treatment), Fadlon and Nielsen (2021) (severe health shocks, stacked with future-treated households as controls), and Deshpande and Li (2019) (SSA field-office closings, stacked with future-closing ZIP codes as controls). Three differences are worth flagging.

The three designs differ first in the choice of control group. **fadlonFamilyLaborSupply2021a** and Deshpande and Li (2019) use future-treated units (those that will be treated at $\tau + \Delta$) as controls, because never-treated units are scarce or selected in their settings (no never-dying households; selection into closing eligibility). Cengiz et al. (2019) uses clean controls (states with no minimum-wage events in the 8-year window around the focal event), which permits states with events outside the window to serve as controls. I use the strictest specification: never-changed (NC) municipalities, defined as those whose rate τ^s never changes across 2002–2019. Since 95% of PA municipalities are NC on at least one rate-setter, the never-changed pool is abundant and I do not need to relax to future-treated or clean-within-

window controls.

On the second dimension—continuous treatment—Cengiz et al. (2019) is the closest analogue: their treatment is continuous (log minimum wage) within each stacked sub-experiment. Their identification logic—"aligning events by event-time, not calendar time, equivalent to a setting where the events happen all at once" (Cengiz et al., 2019, Online Appendix C)—applies identically in my setup. I extend their continuous-treatment design by extracting cell-specific dose-response coefficients via the $i(\text{cell}, \Delta\tau)$ interaction, which their pooled event-by-event spec does not produce.

The third difference is the binding-state cell partitioning, which is new relative to all three antecedents. Cengiz et al. (2019) pools all minimum-wage events without a within-design taxonomy of treated units; Fadlon and Nielsen (2021) and Deshpande and Li (2019) have homogeneous treatments (death, closure). My decomposition exploits the piecewise PA EIT rule to partition treated municipalities into cells whose ex ante predicted response differs by binding state, generating a within-design placebo (Inc-nobind) that the standard stacked DiD cannot produce.

Some treated municipalities change rates more than once in 2002–2019, raising a coding question about which event defines the cohort. A treated municipality i may have more than one rate-change event over the panel. I classify each such municipality by the binding-state transition of its first event and assign i to the corresponding cell on rate-setter s. The municipality then appears in exactly one sub-experiment (its first-event cohort c_i^s), not in subsequent cohorts triggered by later events.

Subsequent rate changes within the same municipality accumulate into $\Delta\tau_{it}^s$, the cumulative dose from the municipality's pre-event baseline. Concretely: suppose municipality i has $\tau_i^s = 0.50$ in years 2005–2009, $\tau_i^s = 0.70$ in 2010 (first event, $c_i^s = 2010$, baseline $= 0.50$), and $\tau_i^s = 1.00$ from 2014 onward (second event). Then $\Delta\tau_{it}^s = 0$ for $t < 2010$, $\Delta\tau_{it}^s = 0.20$ for $t \in [2010, 2013]$, and $\Delta\tau_{it}^s = 0.50$ for $t \geq 2014$. The stacked-continuous specification absorbs this trajectory through the dose-stack interaction in a single regression on municipality i's

2005–2015 window (its first-event cohort 2010 with $K = 5$). Muni i is not duplicated as a 2014 cohort treated municipality; it stays in its 2010 cohort. The cumulative dose path encodes the second event by raising $\Delta\tau^s_{it}$ from 0.20 to 0.50 at year 2014.

This first-event classification has three deliberate consequences. First, no treated municipality serves as control to itself; each municipality appears in exactly one stack. Second, multi-event treated municipalities do not receive disproportionate weight from appearing in multiple stacks. Third, the per-pp coefficient β^s_k is interpretable as a dose-response elasticity over the cumulative rate change, which is the welfare-relevant magnitude for fiscal calibration.

For each treated municipality i in stack $\mathcal{S}^s_c$, define the cumulative dose-stack as $\Delta\tau^s_{it} = \tau^s_{it} - \tau^s_{i,c-1}$, with $\Delta\tau^s_{it} = 0$ for $t < c$ (pre-event) and for NC municipalities at all t. The stacked continuous DiD specification, for outcome Y, runs the regression on the stacked panel:

$$\begin{aligned} Y_{\tilde{i}\tilde{t}} \;=\; \alpha_{\tilde{i}} \;+\; \gamma_{\tilde{t}} + \sum_{k \in \mathcal{C}_{\setminus \mathrm{NC}}} \beta^{\mathrm{muni}}_k \, \Delta\tau^{\mathrm{muni}}_{it} \cdot \mathbf{1}\{\mathrm{cell}^{\mathrm{muni}}_i = k\} \\ + \sum_{k \in \mathcal{C}_{\setminus \mathrm{NC}}} \beta^{\mathrm{SD}}_k \, \Delta\tau^{\mathrm{SD}}_{it} \cdot \mathbf{1}\{\mathrm{cell}^{\mathrm{SD}}_i = k\} \\ + \mathbf{X}'_{it}\boldsymbol{\phi} \;+\; \varepsilon_{\tilde{i}\tilde{t}}, \end{aligned} \tag{4}$$

where $\mathcal{C}_{\setminus \mathrm{NC}} = \{$Inc-bind, Inc-becomes, Inc-nobind, Dec-bind, Dec-becomes-non, Dec-nobind$\}$ is the set of non-NC cells. NC is the omitted reference category on each rate-setter, so for each of the two rate-setters (municipality and SD) the regression estimates six β_k coefficients, each interpreted as the dose-response coefficient relative to the NC trend for that rate-setter, holding the other rate-setter's cell $\times$ dose interaction fixed. The unit of observation $\tilde{i}$ in the regression is a (municipality, stack) pair: each treated municipality contributes observations in exactly one stack (its first-event cohort), while each NC municipality contributes observations in every stack. The fixed effect $\alpha_{\tilde{i}}$ is therefore a separate intercept for each (municipality, stack) combination—so a single NC municipality has a different FE in

each cohort it appears in—and $\gamma_{\tilde{t}} = \gamma_{(c,t)}$ is a separate calendar-year shock per cohort. The control vector $\mathbf{X}_{it}$ holds the share-based fiscal controls described above. Standard errors are clustered at the underlying municipality, not at the (municipality, stack) level. Cluster-at-municipality is the right choice because the same municipality's outcomes are correlated across the multiple stacks in which it appears as an NC control; the municipality-by-stack FE absorbs the level but not the residual covariance.

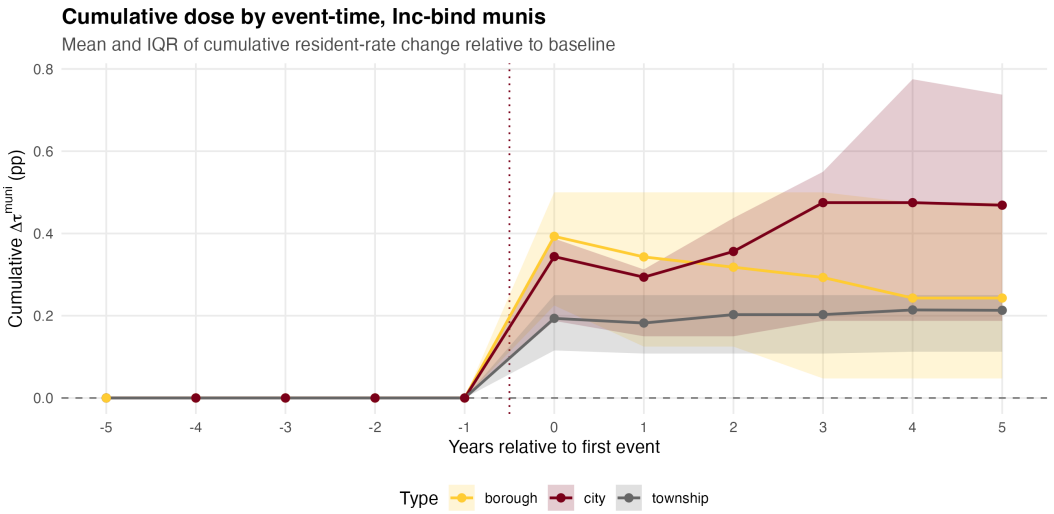
The cell coefficient β_k^s is the per-pp dose-response of outcome Y to a 1 percentage-point change in rate s, restricted to municipalities whose first event on rate-setter s falls in cell k, with NC municipalities on rate-setter s as controls. Specifically, $\beta^{\text{muni}}_{\text{Inc-bind}}$ is the dose-response of Y to municipality-rate hikes among municipalities whose home stack was binding both pre- and post-event, identified by comparing their post-event trajectory to NC municipalities' contemporaneous trajectory. The placebo within the design is $\beta^{\text{muni}}_{\text{Inc-nobind}}$: under the institutional rule, this should be zero (the rate hike does not move the bill).

Within each sub-experiment, the comparison is treated cohort vs NC controls in the same calendar-year window. The cell decomposition partitions the treated set by whether the rate change moved the bill. The reference for each cell coefficient is the NC trend for the same rate-setter. The triple decomposition (municipality × SD × nr) lets us read off, in a single regression: which rate moves which outcome, conditional on the municipality's binding state for that rate-setter.

The identifying variation in this design comes from the cross-unit distribution of cumulative rate changes after each unit's first event. Figure 10 traces the mean cumulative municipality-resident-rate change (with IQR) of Inc-bind units relative to their pre-event baseline. The pre-event window is mechanically flat at zero by construction (dose is measured relative to the unit's own $c-1$ rate); the post-event window shows the realized cumulative $\Delta\tau^{\text{muni}}$ rising to roughly 0.2 percentage points for townships, 0.3 for boroughs, and 0.4 for cities by $k = +5$. Figure 11 reports the cross-unit distribution of the cumulative dose at $k = +5$: most Inc-bind municipalities experience cumulative changes of 0.2–0.4 pp, with a

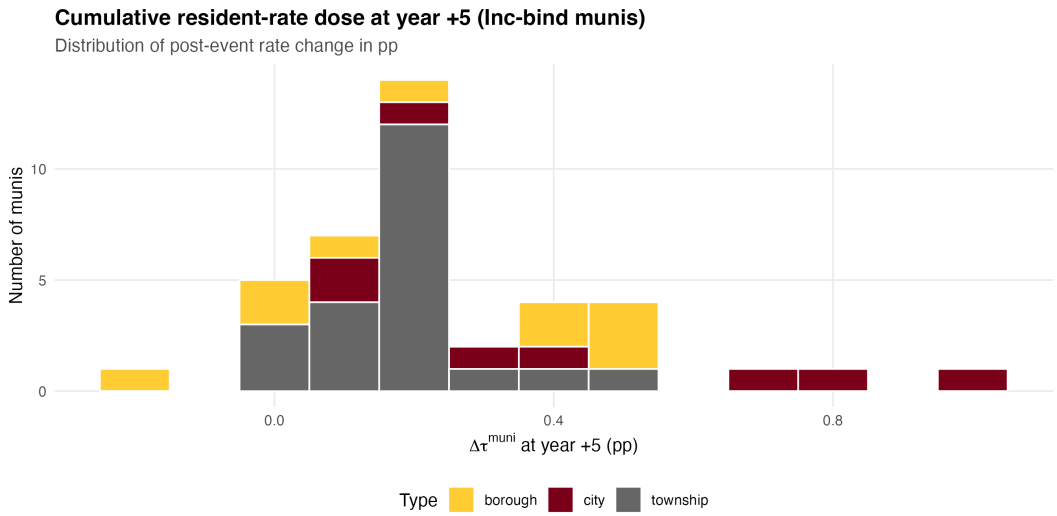
small right tail (Wilkes-Barre and Coatesville) at 0.8–0.9 pp. The realized-dose ES diagnostic exploits this across-unit variation to test pre-trends conditional on eventual treatment intensity.

Figure 10: Cumulative municipality-resident-rate dose by event-time, Inc-bind units, by municipality type.



Notes: Mean and inter-quartile range of cumulative $\Delta\tau_{it}^{\text{muni}}$ (cumulative resident-rate change from baseline $\tau_{i,c-1}^{\text{muni}}$) by relative event-time k, separately for cities, boroughs, and townships. Pre-event values are zero by construction. Post-event cumulative dose averages ≈ 0.4 pp for cities and ≈ 0.2 pp for townships by $k = +5$.

Figure 11: Distribution of cumulative municipality-resident-rate dose at $k = +5$, Inc-bind units, by municipality type.



Notes: Cross-unit distribution of cumulative resident-rate change at the end of the post-event window ($k = +5$). The mass at 0.2–0.4 pp reflects the standard half-point increment; the right tail at 0.8–0.9 pp is driven by a small set of Home Rule cities.

5.3 Three Estimating Units

The headline equation (4) is estimated at three units of observation. The choice of unit interacts directly with how the two rate-setting bodies (municipality and school district) enter the regression.

At the municipality level, the unit of observation is municipality × year. Workplace-panel and residence-panel outcomes are constructed by aggregating LODES origin–destination flows over blocks within the municipality. The municipal resident rate enters the regression at its native level. The school district resident rate is assigned to each municipality-year as $\max(\tau^{SD})$ across the school districts the municipality reports to in DCED records; for the 97.6% of municipalities that report to a single SD this is exact, and for the 2.4% that span multiple SDs the max-aggregation is conservative. Fiscal share controls are taken from the municipality AFR. Fixed effects: municipality-by-stack and year-by-stack. Standard errors clustered at the municipality level.

At the school district level, the unit is SD $\times$ year. Outcomes are constructed by aggregating LODES OD flows over blocks within the SD. The SD resident rate enters at its native level. The municipal resident rate is assigned to each SD-year as the within-SD aggregation across the SD's constituent municipalities. Fiscal share controls are taken from Census F-33 SD-level revenue and expenditure data, mirroring the municipality AFR shares. Fixed effects: SD-by-stack and year-by-stack. Standard errors clustered at the SD level.

At the block level, the unit is census block $\times$ year. Each block is assigned its own municipality and its own SD, so both rate-setters' policy variation enters the regression at its native level with no within-jurisdiction aggregation. Outcomes are block-defined: log of workers living in the block, log of workers working in the block, the share of block residents who work in the block's own municipality, and the share of block residents who work outside the municipality. Fiscal share controls include both municipality AFR shares and F-33 SD shares, since each block is exposed to both governments. Fixed effects: block-by-stack and year-by-stack. Standard errors two-way clustered at the municipality and SD levels.

The block level is the preferred causal estimate for two reasons. First, the block is the only level where both rate-setting jurisdictions enter at their native geography, with no within-jurisdiction aggregation step the reader has to defend. Second, the realized-dose pre-trend diagnostic is cleanest at the block level (Table 10): all four pre-event coefficients are individually $p > 0.10$, with a mean pre-period coefficient of -0.027. The block-level joint Wald F-test rejects parallel trends ($p < 0.001$) but the rejection is driven by the test's enormous statistical power at $N \approx 21$M block-years rather than by economically meaningful deviations; the Rambachan–Roth sensitivity bound at $\bar{M} = 0$ leaves the post-period CI bounded away from zero. The municipality-level analogue on the log count of residents working in their own municipality returns a joint Wald F-test that does not reject parallel trends ($p = 0.23$, stat $= 5.62$ on 4 df), with three of four pre-period coefficients individually $p > 0.10$. The municipality-level diagnostic is therefore the cleanly-non-rejecting one, while the block-level diagnostic is hyper-powered rather than a sharp failure, but it is enough

that I treat the block-level estimate as the causal headline and the municipality-level as policy-jurisdiction magnitude calibration: the typical PA municipality has ∼40 live-and-work workers, a meaningful per-jurisdiction benchmark, but the noisier pre-period in the diagnostic means I do not promote the municipality-level Inc-bind coefficient to a second causal headline. The school-district aggregation produces direction-mixed coefficients that do not survive the pre-trend diagnostic and is reported as an aggregation diagnostic in the Companion Appendix.

5.4 Pre-Trend Diagnostics

Because the first-difference dose $\Delta\tau_{it}$ is mechanically zero in all pre-event periods by construction, the headline specification cannot estimate conventional pre-treatment lead coefficients. I report two complementary event-study diagnostics designed for this setting. The first tests *strong* parallel trends across treatment-intensity magnitudes (recommended by the recent continuous-DiD literature for designs with a continuous dose). The second tests *average* parallel trends under a binary cohort treatment and complements the first by capturing pre-event mean differences regardless of dose.

5.4.1 Realized-Dose Sun–Abraham Event Study

To assess parallel trends in the cross-dose sense (Callaway & Sant'Anna, 2021; de Chaisemartin et al., 2024), I estimate a separate event-study specification using each unit's realized cumulative dose, $\widetilde{D}_i = \tau^M_{i,c+5} - \tau^M_{i,c-1}$, which is constant within unit and varies across units in proportion to eventual treatment intensity. For each municipality rate-setter cell c, restrict to the subset of units in $\{\text{NC}, c\}$ and run

$$Y_{it} = \alpha_{\tilde{i}} + \gamma_{\tilde{t}} + \sum_{k \neq -1} \beta^M_k \, \mathbf{1}\{\ell_{it} = k\} \, \widetilde{D}_i + \mathbf{X}'_{it}\boldsymbol{\phi} + \varepsilon_{it}, \tag{5}$$

where $\ell_{it} = t - c_i^M$ is event-time relative to the unit's first rate event. The β_k^M at pre-event leads ($k \leq -2$) test whether units with larger eventual rate changes were on differentially sloped pre-event trajectories; under parallel trends in the strong sense, these coefficients are zero. The post-event coefficients (β_k^M for $k \geq 0$) recover the per-pp dose-response at each event time and average to the corresponding cell coefficient in the headline specification when the unit's rate path is a single step at the event year.

5.4.2 Binary Sun–Abraham Interaction Weighted Event Study

The realized-dose diagnostic above tests whether *larger* eventual rate changes are associated with differential pre-trends. A complementary diagnostic asks the simpler question of whether ever-treated units differ in mean trajectory from never-treated units, ignoring dose magnitude. For this I use the Sun–Abraham interaction-weighted (SA-IW) event study (Sun & Abraham, 2021) implemented in `fixest::sunab()`:

$$Y_{it} = \alpha_i + \gamma_t + \sum_{e} \sum_{k \neq -1} \delta_{e,k} \, \mathbf{1}\{c_i = e\} \, \mathbf{1}\{\ell_{it} = k\} + \mathbf{X}_{it}'\boldsymbol{\phi} + \varepsilon_{it}, \tag{6}$$

estimating a separate dynamic ATT $\delta_{e,k}$ per cohort e and event-time k, then interaction-weighting the $\delta_{e,k}$ across cohorts to deliver a single cohort-robust path of pre- and post-event coefficients. Binary cohort assignment is constructed per binding-state cell: ever-Inc-bind municipalities are coded **1** in the year of their first Inc-bind event and thereafter, NC municipalities are coded **0** throughout, and other cells are excluded from the cell-specific subset.

The two diagnostics are complementary: failure of the realized-dose test indicates dose-magnitude selection (high-dose units were trending differently), while failure of the binary test indicates any mean differential trajectory between treated and control units. Passing both is required for clean identification under the strong parallel-trends assumption. Block-level binary SA-IW is unstable (thin cohort × event-time cells under stacking) and is reported

at the municipality and SD levels only.

I do not impose a uniform pre-trend standard across every outcome in the paper, because not every outcome carries the causal claim. The primary residential-sorting outcome is the block-level log count of workers living in the block, and the realized-dose pre-trend on the binding cell must be clean here for the headline to go through; this is the outcome the paper asks the reader to believe. Municipality-level resident counts—log residents working in their own municipality, log employed residents—are retained as policy-jurisdiction magnitude calibration. I acknowledge and caveat imperfect pre-trends at the municipality scale, but these outcomes do not carry the causal claim. Secondary equilibrium outcomes such as log in-commuters, log total jobs, and the industry decomposition are reported as suggestive evidence of a local-services demand channel; imperfect pre-trends are acceptable there because these outcomes are downstream of the residential-sorting margin and are not the identification object. Block-level share outcomes—the share of block residents working in their own municipality, the share of in-block jobs filled by out-commuters—are sensitive to sparse denominators and to endogenous denominator dynamics at block geography, so I report them in the Appendix and do not lean on them. The headline block-level log home count in the Inc-bind cell passes the strong standard; other outcomes are retained with the caveats appropriate to their role.

5.5 Sample Variants

The DCED nonresident-rate panel has reporting gaps—55–60% coverage pre-Act 32, 72–80% post—which affect the binding-state classification when the nonresident rate is missing. I report three sample variants. The strict variant drops municipality-years with no recorded nonresident rate. The filled variant imputes `muni_nonres_eit` as the contemporaneous resident rate when missing, which is statute-anchored: under Act 511 the nonresident rate is bounded above by the resident rate. The drop-incomplete variant keeps all municipality-years and drops NA only at regression time per outcome, maximizing N. The headline

tables use the drop-incomplete sample; the other two are reported as robustness columns, and coefficients are stable across them.

Section 6 reports the headline binding-state decomposition at each estimating unit, the pooled Inc/Dec literature benchmark, and the industry-mechanism decomposition. Auxiliary diagnostics (pre/post Act 32 split, dose-magnitude terciles, alternative control pools, sample variants, and Rambachan–Roth sensitivity) are in Section 7 and Appendix D.

6 Results

I report results in two layers. Section 6.1 presents the pooled Inc/Dec benchmark—what a conventional local-tax design would recover if it ignored the Pennsylvania max-rule. Section 6.2 reports the binding-state decomposition at three levels of aggregation. The block-level decomposition is the preferred causal estimate because both rate-setting jurisdictions enter at their native geography; the municipality-level decomposition is reported as policy-jurisdiction magnitude calibration; the school-district aggregation is a diagnostic non-result. The central empirical comparison throughout is between the Inc-bind cell (where the statutory rate change moves the worker's bill) and the Inc-nobind cell (where it does not), with Inc-nobind serving as the within-design placebo. Mechanism evidence from the industry decomposition appears in §6.3; the pre-/post-Act-32 enforcement-break split appears in the Robustness section.

6.1 Pooled Inc/Dec Benchmark

I first report the pooled Inc/Dec specification from Equation (3), which collapses each rate-setter to two direction tags. The pooled coefficients are useful because they approximate the estimand in a conventional local-tax design: the average per-pp dose-response of a municipality-rate change and an SD-rate change without distinguishing whether the rate change moves the worker's bill. In Pennsylvania, however, the pooled specification is not the

structural object of interest. It averages together rate changes that raise tax liability (Inc-bind) and rate changes that do not (Inc-nobind), and the binding-state decomposition in Section 6.2 shows which sub-population of the pooled treated cohort drives each coefficient. Cell counts and stacked-design support are reported in Section 4.2.

6.1.1 Block level

Table 7 reports the per-pp pooled coefficients at the block level. The table includes four coefficients of interest per outcome: the municipality resident-rate Inc effect β^{M}_{Inc}, the municipality resident-rate Dec effect β^{M}_{Dec}, and the analogous school-district-rate coefficients $\beta^{SD}_{\text{Inc}}, \beta^{SD}_{\text{Dec}}$.

Table 7: Pooled Inc/Dec headline — block level (per pp)

Outcome	Log(workers live)	Log(workers work)	Share live+work in muni
Municipality rate (τ^M)			
$\Delta\tau > 0$ (Inc)	-0.045	-0.066***	-0.041***
	(0.034)	(0.024)	(0.009)
$\Delta\tau < 0$ (Dec)	0.255***	0.097	-0.025
	(0.089)	(0.229)	(0.022)
School district rate (τ^{SD})			
$\Delta\tau > 0$ (Inc)	0.035	0.028	-0.000
	(0.031)	(0.025)	(0.005)
$\Delta\tau < 0$ (Dec)	0.001	-0.002	-0.008***
	(0.008)	(0.011)	(0.003)
N	23,567,655	8,751,581	23,567,655
R^2	0.858	0.865	0.563
Within R^2	0.000	0.000	0.000
Dep. mean	2.238	2.318	0.121
Dep. SD	1.186	1.655	0.184

Notes: Cluster-robust standard errors in parentheses. Each regression includes both rate-setters' direction-pooled treatments jointly; the four reported coefficients are β^M_{Inc}, β^M_{Dec}, β^{SD}_{Inc}, β^{SD}_{Dec}. $^*p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

6.1.2 Municipality level

Table 8 reports the analogous per-pp pooled coefficients at the municipality level, with the same four coefficients of interest per outcome.

Table 8: Pooled Inc/Dec headline — municipality level (per pp)

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Municipality rate (τ^M)				
$\Delta\tau > 0$ (Inc)	-0.228***	-0.209***	-0.035	-0.019**
	(0.058)	(0.056)	(0.033)	(0.008)
$\Delta\tau < 0$ (Dec)	0.260**	0.014	0.058	0.022
	(0.107)	(0.143)	(0.138)	(0.019)
School district rate (τ^{SD})				
$\Delta\tau > 0$ (Inc)	0.128*	0.105	0.051	0.007
	(0.070)	(0.067)	(0.038)	(0.005)
$\Delta\tau < 0$ (Dec)	-0.310***	-0.170**	0.020	-0.003
	(0.110)	(0.076)	(0.051)	(0.007)
N	218,677	239,023	242,722	239,033
R^2	0.964	0.959	0.981	0.823
Within R^2	0.001	0.001	0.001	0.001
Dep. mean	3.677	5.847	6.697	0.093
Dep. SD	1.934	1.887	1.299	0.076

Notes: Cluster-robust standard errors in parentheses. Each regression includes both rate-setters' direction-pooled treatments jointly; the four reported coefficients are β^M_{Inc}, β^M_{Dec}, β^{SD}_{Inc}, β^{SD}_{Dec}. $^*p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

6.1.3 School-district level (brief)

The school-district-level analogue (pooled across the SD's constituent municipality rates) is reported in Table 9; SD-level identification is weaker than at the municipality or block level and the table is included for completeness.

Table 9: Pooled Inc/Dec headline — school district level (per pp)

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Municipality rate (τ^M)				
$\Delta\tau > 0$ (Inc)	0.017	0.013	0.014	0.000
	(0.022)	(0.017)	(0.014)	(0.001)
$\Delta\tau < 0$ (Dec)	-0.017***	-0.010*	-0.005***	-0.000
	(0.003)	(0.005)	(0.002)	(0.000)
School district rate (τ^{SD})				
$\Delta\tau > 0$ (Inc)	0.150**	0.154**	0.196***	0.001
	(0.070)	(0.069)	(0.054)	(0.004)
$\Delta\tau < 0$ (Dec)	-0.008	0.089	0.139***	-0.012
	(0.069)	(0.078)	(0.040)	(0.010)
N	51,440	51,474	51,474	51,474
R^2	0.986	0.987	0.987	0.953
Within R^2	0.003	0.002	0.004	0.004
Dep. mean	6.260	8.335	8.768	0.125
Dep. SD	1.298	1.202	0.909	0.068

Notes: SD-level regression: τ^{SD} is the school-district's own rate; the muni-rate row uses the within-SD mean of constituent muni resident rates as an aggregated treatment. Cluster-robust standard errors in parentheses. Each regression includes both rate-setters' direction-pooled treatments jointly; the four reported coefficients are β^M_{Inc}, β^M_{Dec}, β^{SD}_{Inc}, β^{SD}_{Dec}. $^*p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

6.1.4 Realized-dose ES benchmark diagnostics

Tables 10 and 11 report the ATT-summary statistics for the block-level pooled event study, separately on the municipality-cohort and SD-cohort stacks, with cluster-robust Wald pre-trend p-values. Figures 12 and 13 plot the event-time paths for each stack, side by side, with color encoding direction (red Inc, blue Dec). Tables 12 and 13 report the analogous ATT summaries for the municipality- and SD-level pooled specifications; Figure 14 overlays those event studies on the same axes for visual comparison. The Wald p-values are the joint pre-trend test $H_0\colon \beta_k = 0$ for all $k \leq -2$ using the model's cluster-robust variance matrix.

Table 10: Pooled realized-dose event study — block level (muni-cohort).

Direction	Log(workers live)	Log(workers work)	Share live+work in muni
Inc (post)	-0.032	-0.053	-0.035
	(0.021)	(0.014)	(0.005)
Inc (pre)	0.079	0.028	0.008
	(0.020)	(0.011)	(0.005)
Inc Wald p	0.0019	0.2118	0.0002
Dec (post)	0.217	0.123	-0.045
	(0.046)	(0.100)	(0.007)
Dec (pre)	0.171	0.160	0.018
	(0.052)	(0.118)	(0.010)
Dec Wald p	<0.0001	<0.0001	0.0021

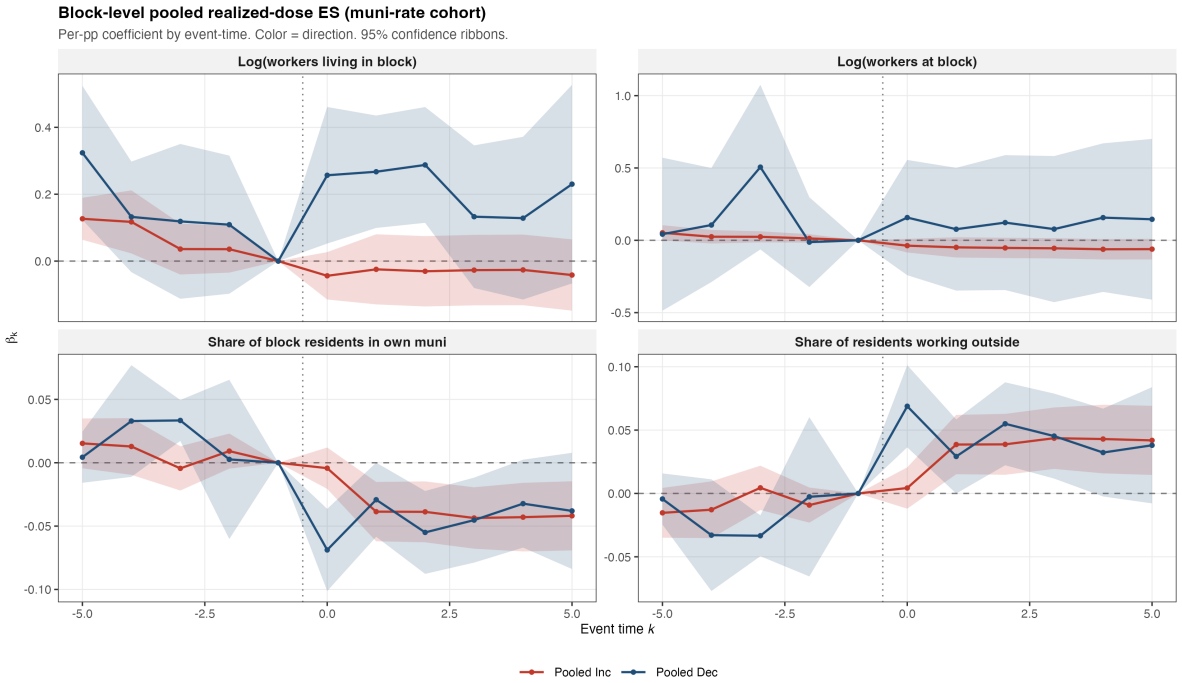
Notes: Cluster-robust standard errors in parentheses. Post = mean of β_k for $k = 0, \ldots, 5$; pre = mean for $k \leq -2$. Wald p is the cluster-robust joint test of $H_0 \colon \beta_k = 0$ across pre-event k. Stack built on muni-cohort years; event time defined against each block's muni-rate event. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 11: Pooled realized-dose event study — block level (SD-cohort).

Direction	Log(workers live)	Log(workers work)	Share live+work in muni
Inc (post)	0.022	-0.012	—
	(0.017)	(0.014)	—
Inc (pre)	-0.126	-0.003	—
	(0.040)	(0.047)	—
Inc Wald p	0.0573	0.0151	—
Dec (post)	—	0.322	-0.117
	—	(0.063)	(0.011)
Dec (pre)	—	0.540	-0.012
	—	(0.131)	(0.008)
Dec Wald p	—	$<$0.0001	0.0067

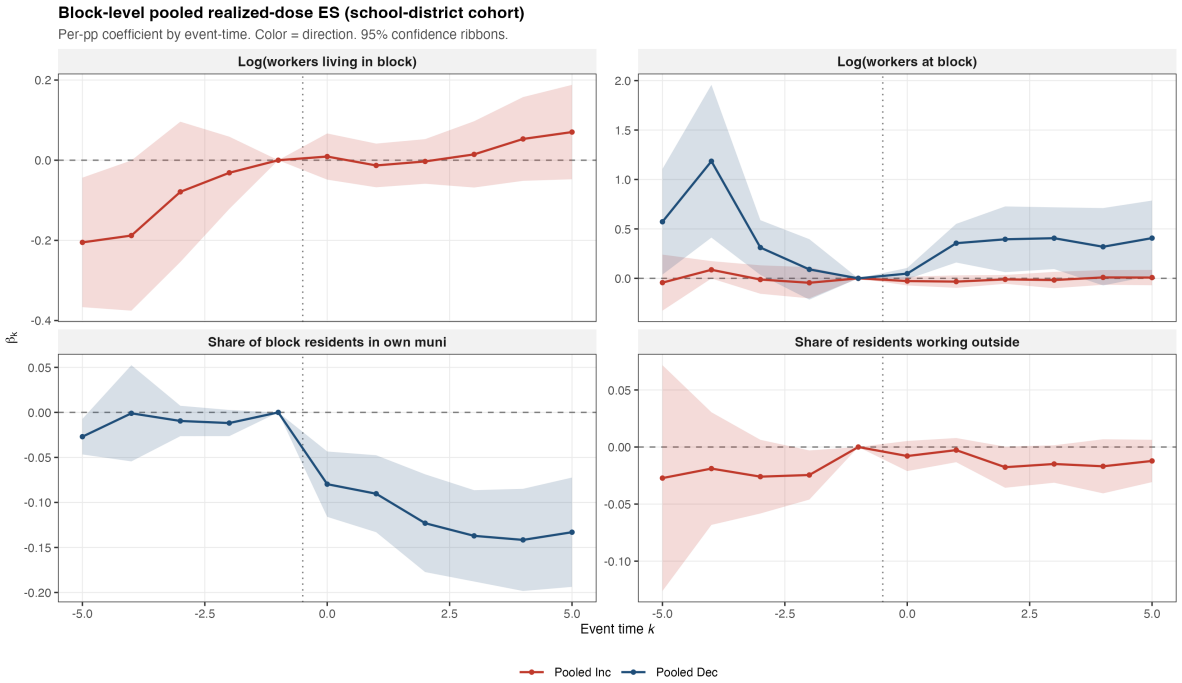
Notes: Cluster-robust standard errors in parentheses. Post = mean of β_k for $k = 0, \ldots, 5$; pre = mean for $k \leq -2$. Wald p is the cluster-robust joint test of H_0: $\beta_k = 0$ across pre-event k. Stack built on SD-cohort years; event time defined against each block's SD-rate event. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Figure 12: Block-level pooled realized-dose event study, municipality-cohort stack. Inc (red) and Dec (blue) coefficients on the four block outcomes corresponding to the four pooled coefficients in Table 7.



Notes: Per-pp realized-dose coefficient at each event-time. Block-level stacked-continuous DiD with block-by-stack and year-by-stack fixed effects and two-way clustering at the municipality and school-district levels. Relative event-time k is defined against each block's first municipality-rate event year c_b^M. Ribbons are 95% confidence bands; reference period $k = -1$ omitted.

Figure 13: Block-level pooled realized-dose event study, SD-cohort stack. Inc (red) and Dec (blue) coefficients on the four block outcomes corresponding to the four pooled coefficients in Table 7.



Notes: As in Figure 12, but with relative event-time k defined against each block's first school-district-rate event year c_b^{SD}. The combined municipality-cohort and SD-cohort overlay is reported in Appendix B.4 for visual comparison.

Table 12: Pooled realized-dose event study — municipality level.

Direction	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc (post)	-0.197	-0.218	-0.039	-0.015
	(0.026)	(0.028)	(0.014)	(0.003)
Inc (pre)	0.094	0.010	0.054	0.011
	(0.034)	(0.058)	(0.032)	(0.006)
Inc Wald p	<0.0001	0.0513	0.4834	0.5071
Dec (post)	0.674	0.264	0.175	0.020
	(0.178)	(0.091)	(0.041)	(0.008)
Dec (pre)	0.190	-0.139	0.169	0.004
	(0.217)	(0.100)	(0.044)	(0.007)
Dec Wald p	<0.0001	<0.0001	0.0003	0.0001

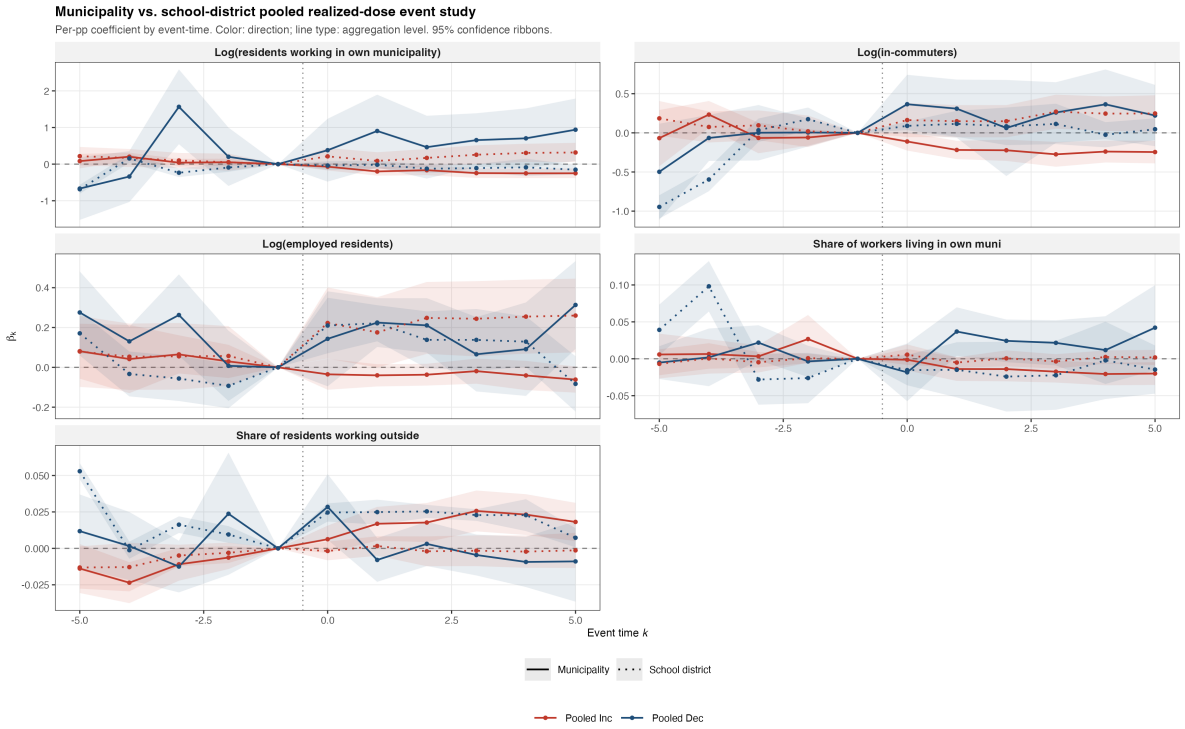
Notes: Cluster-robust standard errors in parentheses. Post = mean of β_k for $k = 0, \ldots, 5$; pre = mean for $k \leq -2$. Wald p is the cluster-robust joint test of H_0: $\beta_k = 0$ across pre-event k. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 13: Pooled realized-dose event study — school district level.

Direction	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc (post)	0.225	0.204	0.234	0.000
	(0.051)	(0.045)	(0.038)	(0.002)
Inc (pre)	0.134	0.094	0.061	-0.003
	(0.060)	(0.052)	(0.042)	(0.004)
Inc Wald p	0.4556	0.2738	0.8816	0.8688
Dec (post)	-0.088	0.071	0.126	-0.016
	(0.030)	(0.038)	(0.028)	(0.009)
Dec (pre)	-0.214	-0.334	-0.003	0.021
	(0.030)	(0.038)	(0.029)	(0.009)
Dec Wald p	<0.0001	<0.0001	<0.0001	<0.0001

Notes: Cluster-robust standard errors in parentheses. Post = mean of β_k for $k = 0, \ldots, 5$; pre = mean for $k \leq -2$. Wald p is the cluster-robust joint test of H_0: $\beta_k = 0$ across pre-event k. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Figure 14: Municipality vs. school-district pooled realized-dose event study: direction (color) and aggregation level (line type).



Notes: Per-pp realized-dose coefficient at each event-time. Color: direction (red = Inc, blue = Dec). Line type: aggregation level (solid = municipality stacked-continuous DiD; dotted = school-district stacked-continuous DiD on the aggregated outcomes). Ribbons are 95% confidence bands.

The pooled coefficients in Tables 7–9 mask substantial within-direction heterogeneity relative to the cell-decomposed magnitudes. At the block level, the pooled-Inc post-event coefficient on the log count of workers living in the block is approximately −0.08, while the realized-dose ES post-event mean restricted to the Inc-bind cell alone is −0.305 (a fourfold attenuation when Inc-bind is averaged with Inc-becomes and the Inc-nobind placebo cells). At the municipality level the pooled-Inc coefficient on the log count of residents working in their own municipality is approximately −0.23, compared to an Inc-bind cell coefficient of −0.401. The pooled coefficient is a sample-share weighted average of the binding-state cell coefficients; the decomposition in Sections 6.2.1– 6.2.2 shows that this aggregate is driven almost entirely by Inc-bind, while Inc-nobind delivers the predicted null behavioral response. This is the central point of the pooled-then-decompose structure: the institutional binding

rule, not the statutory rate change itself, identifies the response.

Pooled-Dec coefficients deserve a separate caveat on pre-trends. The block-level pooled-Dec ES coefficient on the log count of workers living in the block has a post-event mean of +0.22 and a pre-event lead of +0.17; the post/pre difference is small and the pre-trend fails parallel trends. The Dec direction reflects selection of a small set of declining or politically distinct municipalities into rate cuts, not a clean binding-rule mirror of the Inc effect, and is reported here for completeness. The Inc direction, with its clean pre-trend on the count outcomes, carries the causal weight in the analysis.

6.1.5 Binary ES benchmark diagnostics (appendix)

A binary pooled-spec analog (Companion Appendix, "Pooled Inc/Dec Binary-Event Analog") replaces the per-pp dose interactions with four post-treatment indicators, comparable to a Bagchi (2022) binary DiD spec. Coefficients agree in sign with the per-pp pooled estimates and are deferred to the Companion Appendix in long form; I report them here only as a robustness column at the headline outcome.

6.2 Binding-State Decomposition

6.2.1 Block level (main causal evidence)

The preferred causal estimate is the block-level municipal Inc-bind coefficient on the log count of workers living in the block. This is the cleanest test of the EIT mechanism for two reasons. First, both rate-setting jurisdictions enter the regression at their native geography—each block is exposed to a single municipality and a single school district, so no within-jurisdiction aggregation is required to assign policy variation. Second, the realized-dose pre-trend on the binding cell is near zero (Table 10; $\beta_{\text{pre}} = -0.027$, $|t| < 1$). Together these features make the block-level decomposition the level at which the paper asks the reader to believe a causal claim.

Table 14 reports the block-level dose-response on the municipality-rate cells, and Table 15

reports the SD-rate cells from the same unified regression. The main residential-sorting outcome is the log count of workers living in the block; the log count of workers working in the block is reported alongside as a secondary count outcome. Block-level share outcomes—the share of block residents working in their own municipality, the share of in-block jobs filled by out-commuters—are mechanically fragile at block geography because both the numerator and denominator can move and many blocks contain small worker counts. I report them in Appendix B.1 as supporting estimates and do not use them to identify the headline causal effect.

Table 14: Block-level dose-response: muni resident-rate cells (count outcomes)

Outcome	Log(workers live)	Log(workers work)
Inc-bind	-0.208***	-0.097*
	(0.054)	(0.052)
Inc-becomes	-0.013	-0.053**
	(0.032)	(0.022)
Inc-nobind	-0.056	-0.113
	(0.098)	(0.075)
Dec-bind	0.269*	0.167
	(0.153)	(0.353)
Dec-becomes-non	0.193	0.413***
	(0.299)	(0.137)
Dec-nobind	0.242***	-0.053
	(0.086)	(0.254)
N	23,567,655	8,751,581
R^2	0.858	0.865
Within R^2	0.000	0.000
Dep. mean	2.238	2.318
Dep. SD	1.186	1.655

Notes: Cluster-robust standard errors in parentheses. Continuous stacked-DiD with block-by-stack and year-by-stack fixed effects; standard errors two-way clustered at the muni and SD levels. Each block is exposed to one municipality and one school district; both rate-setters' policy variation enters at its native level. Share outcomes are reported in the appendix; sparse block-level denominators make share estimates mechanically fragile. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 15: Block-level dose-response: SD resident-rate cells (count outcomes)

Outcome	Log(workers live)	Log(workers work)
Inc-bind	0.097	0.099**
	(0.062)	(0.047)
Inc-becomes	0.016	0.009
	(0.032)	(0.025)
Dec-bind	0.207***	-0.052
	(0.021)	(0.144)
Dec-becomes-non	-0.013*	0.001
	(0.008)	(0.007)
Dec-nobind	-0.486***	0.502***
	(0.099)	(0.168)
N	23,567,655	8,751,581
R^2	0.858	0.865
Within R^2	0.000	0.000
Dep. mean	2.238	2.318
Dep. SD	1.186	1.655

Notes: Cluster-robust standard errors in parentheses. Continuous stacked-DiD with block-by-stack and year-by-stack fixed effects; standard errors two-way clustered at the muni and SD levels. Each block is exposed to one municipality and one school district; both rate-setters' policy variation enters at its native level. Share outcomes are reported in the appendix. Rows omitted (no estimable coefficients in any outcome column): Inc-nobind. These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

The Inc-nobind cell serves as the within-design placebo: it contains municipalities whose home stack remained below the municipality's nonresident rate after the rate hike, so the bill never moved for in-municipality residents. The institutional rule predicts no behavioral response, and the Inc-nobind coefficient on the log count of residents working in their own municipality is statistically null (-0.062, SE 0.169); on log total jobs is -0.201 (SE 0.203);

and on the share of workers living in their own municipality is −0.012 (SE 0.008). The placebo passes on the headline outcomes. Two outcomes show small marginally significant deviations (the log count of out-commuters $+0.217^{**}$ and log employed residents $+0.194^{**}$); with $n = 26$ in this cell these are imprecise and I disclose rather than claim a perfectly clean placebo.

The joint regression also separates the two rate-setting channels, since it includes both municipality-resident-rate cells and SD-resident-rate cells. The municipality resident rate moves residents (Inc-bind coefficient on log residents working in their own municipality: -0.401^{***}; on log employed residents: -0.100^{***}) and in-commuters (-0.212^{**}), consistent with the local-services demand spillover in Section 6.3. The SD resident-rate Inc-bind coefficient on the same outcomes is $+0.177^{*}$ on log residents working in their own municipality and +0.063 on log employed residents: SD rate changes do not produce a clean sorting response through the binding channel even when SD changes flip the binding state. An earlier draft using tract-centroid LODES recovered a positive SD-rate "attraction" effect; under block-level LODES with the unified two-rate-setter design that artifact disappears. The work-side nonresident rate τ^{nr} enters the binding-state classification but is not itself a treatment in the headline regression; its cell coefficients are imprecise because the municipality nonresident rate almost never exceeds the municipality's own home stack at the municipality-self-pair level in PA.

6.2.2 Municipality level (calibration)

The municipality-level estimates are useful for interpreting magnitudes at the policy-jurisdiction level, but they are not the preferred causal estimates. Both rate-setter cells enter the same unified regression as in the block-level model, with one within-municipality aggregation step: each municipality-year is assigned $\max(\tau^{SD})$ across the school districts the municipality reports to (exact for the 97.6% of municipalities that report to a single SD, conservative for the rest). Table 16 reports the municipality-level municipality-rate cells with full diagnostics;

Table 17 reports the SD-rate cells from the same regression.

Table 16: Municipality-level dose-response: muni resident-rate cells

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind	-0.401***	-0.212**	-0.100***	-0.045***
	(0.096)	(0.083)	(0.033)	(0.009)
Inc-becomes	-0.185***	-0.203***	-0.068*	-0.012
	(0.068)	(0.067)	(0.039)	(0.009)
Inc-nobind	-0.057	-0.250	0.074	-0.014
	(0.178)	(0.251)	(0.065)	(0.011)
Dec-bind	0.220**	0.049	-0.054	0.052*
	(0.102)	(0.108)	(0.191)	(0.030)
Dec-becomes-non	-0.291	0.422*	-0.682	-0.073
	(0.854)	(0.241)	(0.721)	(0.066)
Dec-nobind	0.308**	-0.014	0.131	0.013
	(0.126)	(0.194)	(0.171)	(0.026)
N	218,685	239,031	242,730	239,041
R^2	0.964	0.959	0.981	0.823
Within R^2	0.001	0.001	0.001	0.001
Dep. mean	3.677	5.847	6.697	0.093
Dep. SD	1.934	1.887	1.299	0.076

Notes: Cluster-robust standard errors in parentheses. Continuous stacked-DiD with municipality-by-stack and year-by-stack fixed effects; standard errors clustered at the municipality level. Each column is a separate regression on the indicated outcome. Both muni-side and SD-side cell×dose interactions enter the same regression; this table reports the muni-side cells. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 17: Municipality-level dose-response: SD resident-rate cells (partial)

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind	0.177*	0.160*	0.063	0.005
	(0.099)	(0.092)	(0.046)	(0.007)
Inc-becomes	0.103	0.081	0.045	0.008
	(0.090)	(0.086)	(0.052)	(0.007)
Dec-bind	-0.721***	-0.293***	0.212**	-0.034***
	(0.191)	(0.060)	(0.090)	(0.013)
Dec-becomes-non	-0.223**	-0.140	-0.006	0.005
	(0.109)	(0.090)	(0.031)	(0.006)
Dec-nobind	0.048	-0.241	-0.413***	-0.007
	(0.179)	(0.423)	(0.110)	(0.010)
N	218,685	239,031	242,730	239,041
R^2	0.964	0.959	0.981	0.823
Within R^2	0.001	0.001	0.001	0.001
Dep. mean	3.677	5.847	6.697	0.093
Dep. SD	1.934	1.887	1.299	0.076

Notes: Cluster-robust standard errors in parentheses. Continuous stacked-DiD with municipality-by-stack and year-by-stack fixed effects; standard errors clustered at the municipality level. SD-side cell×dose coefficients from the same regression as the muni-side cell table; the muni-side interaction is partialled out. Rows omitted (no estimable coefficients in any outcome column): Inc-nobind. These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

The municipality-level Inc-bind coefficient on the log count of residents working in their own municipality is -0.401^{***}, larger in absolute value than the block-level coefficient on the log count of workers living in the block (-0.208^{***}). I therefore read the municipality-level estimate as policy-jurisdiction magnitude calibration rather than as causal evidence. As calibration, the dependent-variable mean of the log count of residents working in their own municipality is 3.68, so the typical municipality has $e^{3.68} \approx 40$ live-and-work workers. A 1 pp Inc-bind hike implies $40 \cdot (e^{-0.401} - 1) \approx -13$ workers at the municipality-level coefficient and

$40 \cdot (e^{-0.208} - 1) \approx -7.5$ workers at the cleaner block-level coefficient. Actual Inc-bind hikes average 0.27 pp. The municipality-level estimate is a coarser calibration of the magnitude implied by the block-level estimate, and the two directionally agree.

6.2.3 School-district level (diagnostic)

Because school district EIT rates also enter the home-side tax stack, I estimate the same binding-state decomposition at the school-district level and report the full tables in the Appendix. The SD aggregation does not produce robust evidence of residential sorting: the SD-rate Inc-bind coefficient on the log count of residents working in their own municipality is $+0.066^{*}$ (sign-reversed relative to both the municipality- and block-level findings), other Inc-bind outcomes are direction-mixed, and the realized-dose pre-trend diagnostic is weaker than in the block-level municipal specification. I therefore treat the SD-level estimates as an aggregation diagnostic rather than as a second causal treatment channel. The substantive conclusion is negative: the sorting response is concentrated in municipal resident-rate increases that bind, not in school-district rate changes. This is what the binding-state mechanism predicts, since municipal rate changes have a direct, sharp identity at the home-side stack while SD averaging mixes municipal policy environments within the same school district. Full SD tables are deferred to the Appendix.

6.2.4 Realized-dose ES decomposition diagnostics

The realized-dose event study disaggregates the pooled pre-trend by binding-state cell. Table 18 reports the post/pre ATT summary for the per-cell ES at the block level; Figure 15 plots the same coefficients by event-time, faceted by cell, and Figure 17 reports the corresponding municipality-level picture. Two findings stand out, and they tell consistent substantive stories even though the joint Wald test reports opposite verdicts at the two estimating units.

Block level. On the log count of workers living in the block, all four pre-event coefficients

are individually indistinguishable from zero ($p > 0.10$ at $k = -5, -4, -3, -2$) and the mean pre-period coefficient is -0.027. The joint Wald F-test nonetheless rejects parallel trends (stat $= 26.9$ on 4 df, $p < 0.001$) because at $N \approx 21$M block-years the test has extraordinary statistical power to detect economically tiny correlated deviations across leads. Following **rambachanMoreCredibleApproach2023**, I treat the joint Wald as a flag rather than a verdict and rely on the Rambachan–Roth sensitivity bounds in Section 7.5: the post-period CI at strict parallel trends ($\bar{M} = 0$) is $[-0.49, -0.12]$, bounded comfortably away from zero. The Inc-becomes and Inc-nobind cells fail the joint pre-trend test sharply (both $p < 0.001$), which is what the institutional rule predicts: the cohorts that should not respond are selected on baseline trends.

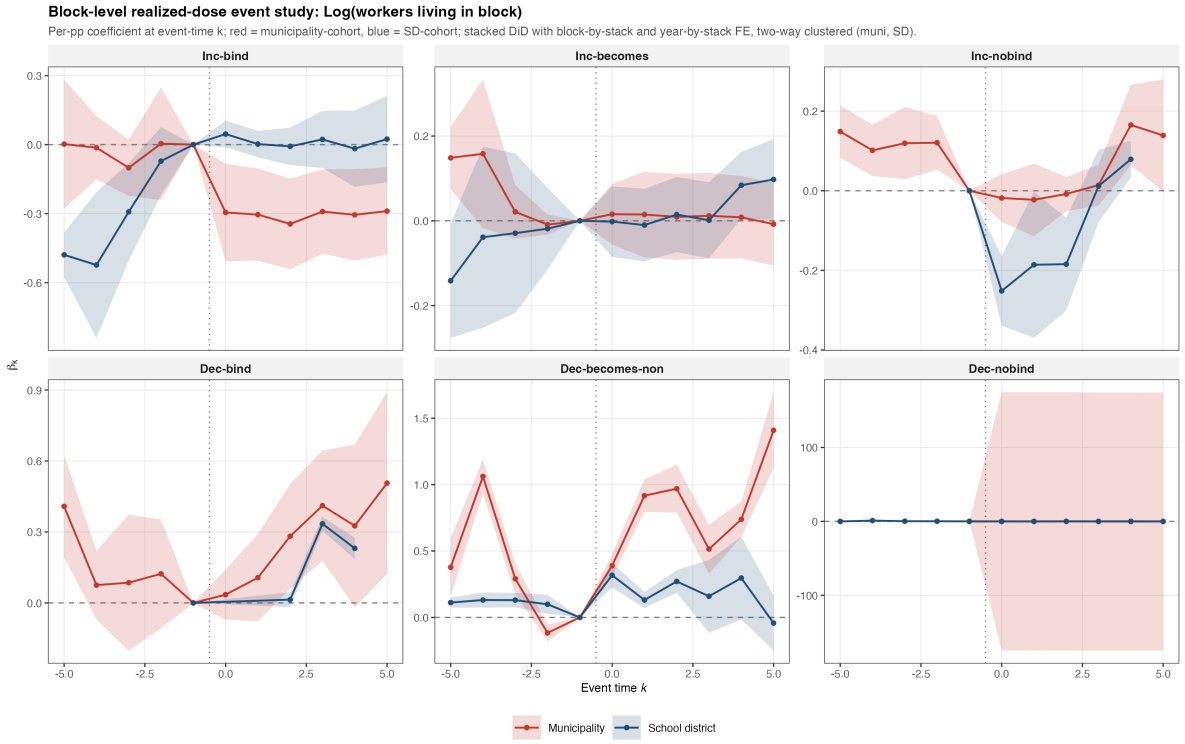
Municipality level. On the log count of residents who both live and work in the municipality, the pre-period coefficients are individually $p = 0.13, 0.07, 0.04, 0.58$ at $k = -5$ through $k = -2$, and the joint Wald F-test does not reject parallel trends (stat $= 5.62$ on 4 df, $p = 0.23$). The smaller cousub-level sample gives the test appropriately calibrated power. Inc-becomes and Inc-nobind again fail the joint test ($p < 0.001$ on both), consistent with the cohort-selection reading.

At both levels the substantive conclusion is the same: the Inc-bind cell is consistent with parallel trends under sensitivity bounds, the Inc-becomes and Inc-nobind placebo cells reflect cohort selection on pre-event trajectories rather than treatment-induced trends, and the post-event Inc-bind response is the cell the binding rule predicts.

Table 18: Realized-dose event study by binding-state cell — block level (muni cohort).

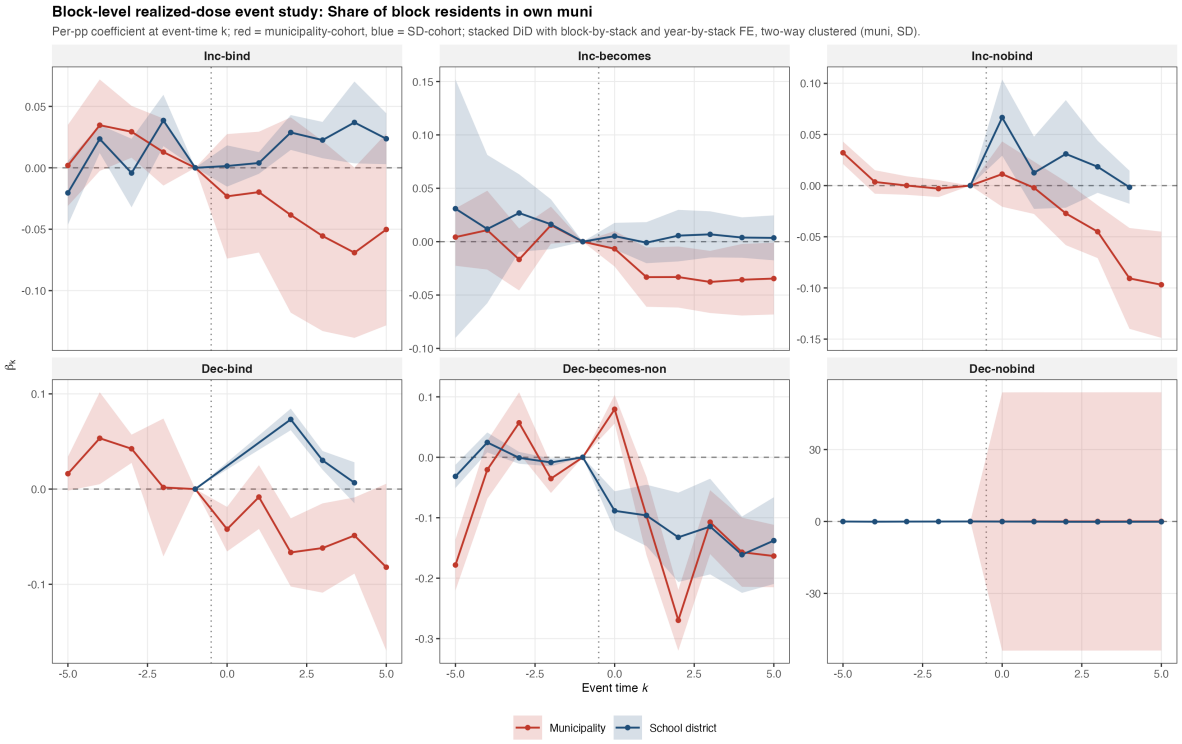
Cell	Log(workers live)	Log(workers work)	Share live+work in muni
Inc-bind (post)	-0.305	-0.079	-0.043
	(0.041)	(0.048)	(0.014)
Inc-bind (pre)	-0.027	0.017	0.020
	(0.053)	(0.033)	(0.008)
Inc-bind Wald p	<0.0001	0.0031	0.0086
Inc-becomes (post)	0.009	-0.046	-0.030
	(0.020)	(0.015)	(0.006)
Inc-becomes (pre)	0.080	0.032	0.003
	(0.026)	(0.014)	(0.007)
Inc-becomes Wald p	0.0002	0.2120	<0.0001
Inc-nobind (post)	0.045	-0.050	-0.042
	(0.018)	(0.019)	(0.008)
Inc-nobind (pre)	0.123	0.025	0.008
	(0.019)	(0.016)	(0.003)
Inc-nobind Wald p	<0.0001	<0.0001	<0.0001
Dec-bind (post)	0.278	0.294	-0.052
	(0.055)	(0.110)	(0.010)
Dec-bind (pre)	0.173	0.457	0.028
	(0.057)	(0.099)	(0.011)
Dec-bind Wald p	<0.0001	<0.0001	<0.0001
Dec-becomes-non (post)	0.823	-0.815	-0.119
	(0.037)	(0.085)	(0.011)
Dec-becomes-non (pre)	0.403	-1.387	-0.044
	(0.036)	(0.036)	(0.011)
Dec-becomes-non Wald p	<0.0001	<0.0001	<0.0001
Dec-nobind (post)	-0.095	0.060	0.015
	(36.376)	(0.042)	(11.217)
Dec-nobind (pre)	—	—	—
Dec-nobind Wald p	—	—	—

Figure 15: Block-level realized-dose event study on the log count of workers living in the block: municipality-cohort (red) and school-district-cohort (blue), faceted by binding-state cell.



Notes: Per-pp realized-dose coefficient at each event-time, block-level stacked-continuous DiD with block-by-stack and year-by-stack fixed effects, two-way clustered at the municipality and school-district levels. Red: the municipality-cohort stack (event time relative to first municipality-rate change). Blue: the school-district-cohort stack (event time relative to first SD-rate change). Ribbons are 95% confidence bands; reference period $k = -1$ omitted.

Figure 16: Block-level realized-dose event study on the share of block residents working in their own municipality: municipality-cohort (red) and school-district-cohort (blue), faceted by binding-state cell.



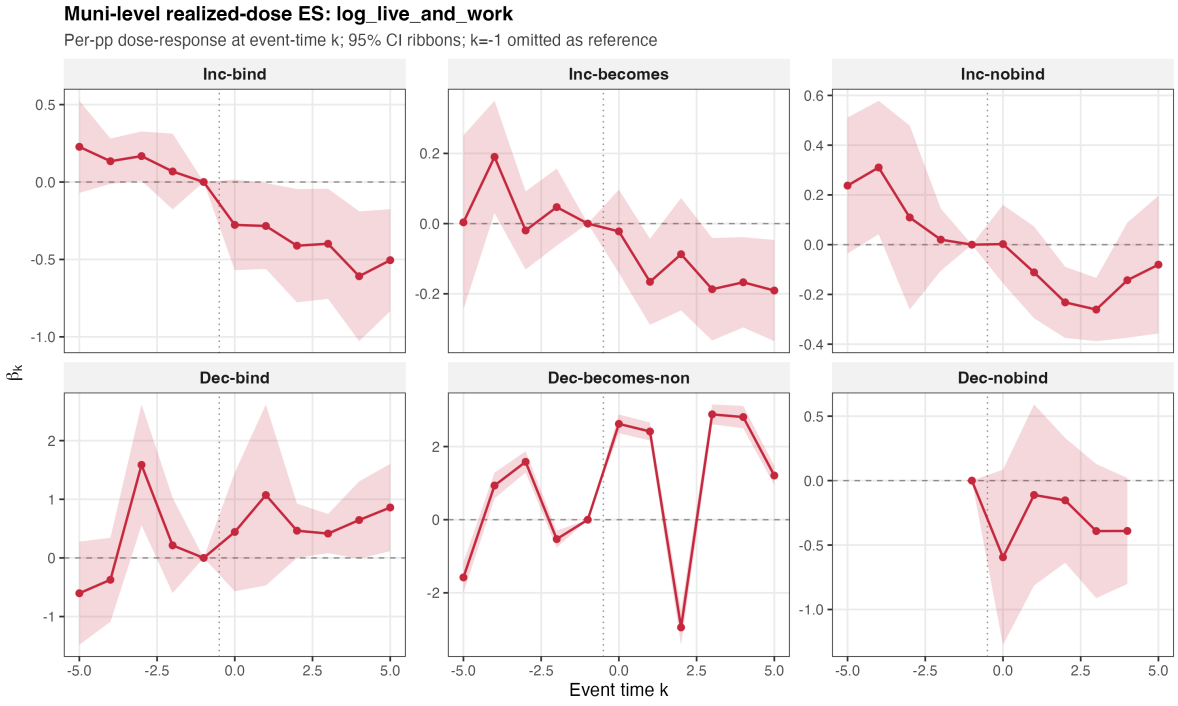
Notes: See Figure 15.

The Inc-becomes cell warrants a separate note. The headline stacked-DiD coefficient is essentially zero at the block level (-0.013), which initially looks like a failure of the predicted intermediate response. The realized-dose ES tells a different story: the pre-event lead is $+0.080$ ($|t| \approx 3$) and the post-event mean is $+0.009$, so the post-minus-pre difference is -0.071 per pp—about a third of the Inc-bind magnitude, in the predicted direction, but masked in the single-coefficient specification by selection on positive pre-event growth in Inc-becomes municipalities. The block-level Inc-becomes coefficient is therefore consistent with the partial-bill-change prediction of the institutional rule once the pre-trend is netted out, but I do not claim it as a clean causal estimate because the strong-parallel-trends assumption fails for this cell. At the municipality level the corresponding Inc-becomes coefficient is -0.181^{**} in the headline specification, and its realized-dose pre-trend is smaller ($+0.055$, $|t| \approx 1.3$);

the municipality-level signal is detectable in the post coefficient because the municipality-by-stack fixed effects absorb less of the within-municipality heterogeneity that attenuates the block-level coefficient. The block-level Inc-bind cell is the only combination of cell and estimating unit where the strong parallel-trends diagnostic passes cleanly, and is therefore the specification on which the paper rests its primary causal claim. The full event-time coefficient detail is in the Companion Appendix.

For completeness, the municipality-level post-event trajectory on Inc-bind log live-and-work is uniformly negative and individually significant from $k = +1$ onward (coefficients of -0.28, -0.41, -0.40, -0.61, -0.50 at $k = 1$ through $k = 5$, all $p < 0.05$), walking down monotonically across the post-window.

Figure 17: Municipality-level realized-dose event study on log(residents who live and work in the municipality), faceted by binding-state cell.



Notes: Per-pp realized-dose coefficient at each event-time, municipality-level stacked-continuous DiD with municipality-by-stack and year-by-stack fixed effects, cluster-robust by municipality. The Inc-bind facet shows pre-event coefficients mostly null pointwise (joint Wald F-test $p = 0.23$), followed by a monotonically negative post-event trajectory.

Appendix B collects the municipality-level realized-dose event-study figures for the four remaining residence-side outcomes (log total residents, log in-commuters, share of workers living in their own municipality, share of residents working outside) as well as the school-district-level and binary Sun–Abraham interaction-weighted analogues at all three estimating units. The patterns mirror the live-and-work panel: Inc-bind has visible pre-event drift consistent with high-dose post-industrial selection, while the within-design Inc-nobind placebo is approximately flat.

6.2.5 Binary ES decomposition diagnostics (appendix)

Binary Sun–Abraham interaction-weighted event studies for the binding-state cells at the municipality and school-district levels are reported in Appendix B; the block-level binary IW analogue is unstable because of thin cohort × event-time cells and is omitted. The binary IW patterns agree in sign with the per-pp realized-dose ES patterns reported above.

6.3 Mechanism: Industry Decomposition

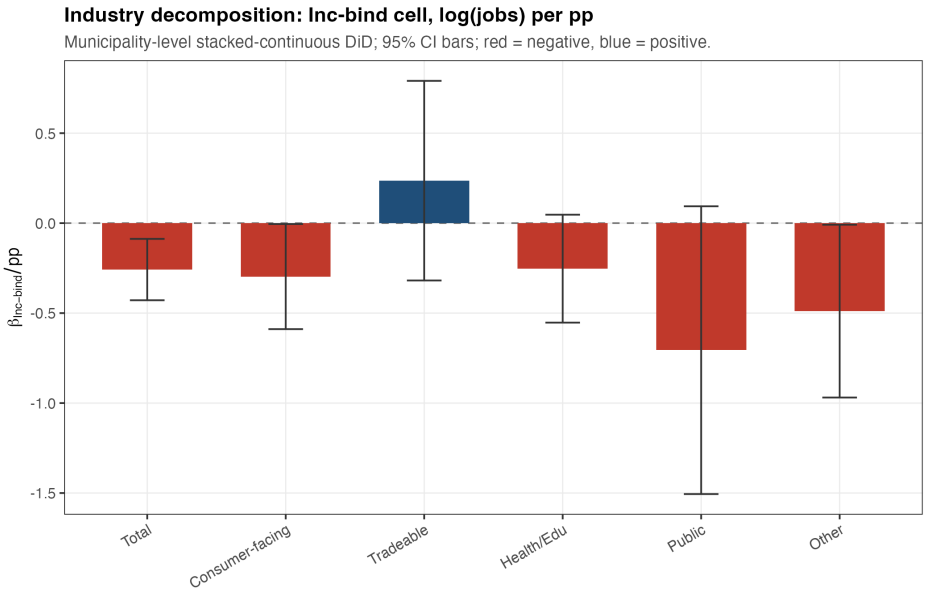
The employment and industry results are mechanism evidence rather than the primary identification object. Headline residential sorting operates on the residence side; the question here is whether the downstream employment response is consistent with a local-services demand channel rather than direct nonresident-tax incidence on jobs. The prediction is that consumer-facing jobs (retail, food, arts, other services—sectors whose demand depends on local population) fall when the municipality resident rate hikes, while tradeable jobs (manufacturing, wholesale, information, finance, professional services—sectors selling to broader markets) are insulated. Using LODES Workplace Area Characteristics aggregated to the municipality by NAICS supersector, Table 19 reports the municipality-resident-rate cell coefficients by industry bucket.

Table 19: Industry decomposition — municipality level.

Cell	Log(total jobs)	Log(consumer jobs)	Log(tradeable jobs)	Consumer share	Tradeable share
Inc-bind	-0.258***	-0.297**	0.236	-0.001	0.106*
Inc-becomes	-0.235***	-0.146**	-0.150	0.008	0.028
Inc-nobind	-0.074	-0.029	-0.265*	0.001	-0.026
Dec-bind	0.018	-0.012	-0.768	-0.014	-0.104
Dec-becomes-non	0.349	-0.742	1.967***	-0.263*	0.551***
Dec-nobind	-0.019	0.195**	0.347**	0.031	0.031*
N	239,001	225,691	219,237	239,001	239,001
R^2	0.961	0.952	0.935	0.760	0.779
Within R^2	0.001	0.001	0.001	0.001	0.001
Dep. mean	5.995	4.553	4.639	0.249	0.250
Dep. SD	1.899	1.935	2.026	0.189	0.202

Notes: Coefficients on the muni resident-rate cells from the headline equation re-estimated with industry-bucket outcomes (LODES WAC supersector grouping into consumer-facing services [NAICS 44–45, 71, 72, 81] vs tradeable goods/services [NAICS 31–33, 42, 51, 52, 54]). Cluster-robust standard errors in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Figure 18: Industry decomposition: municipality-Inc-bind cell coefficient on log(jobs) per pp, municipality level.



Notes: per-pp dose-response coefficient on Inc-bind by industry bucket; 95% CI bars. Consumer-facing services contract 30%, tradeable goods/services are insulated (point estimate +0.236, statistically indistinguishable from zero). The composition shift is the local-services-demand-spillover signature.

Consumer-facing jobs fall by 30 percent in the Inc-bind cell; tradeable jobs do not (point estimate +0.236, statistically indistinguishable from zero). The composition shift is captured by the share of tradeable jobs, which rises 11 percentage points. The pattern is consistent with a local-services demand channel: resident loss reduces local demand and shrinks the sectors most exposed to that demand, while tradeable sectors selling to broader markets are insulated. The pre-trend on Inc-bind employment outcomes is not as clean as on the headline residence-side count outcome, and I do not claim the employment estimates are causal in the same sense; they are reported as suggestive mechanism evidence.

The same decomposition at the SD level (Table 20) is qualitatively different: aggregating municipality-level industry counts to school district via the municipality–SD assignment produces a positive Inc-bind coefficient on log consumer-facing jobs ($+0.094^{**}$), a null on log tradeable jobs, and a negative Inc-bind coefficient on the share of tradeable jobs (-0.031^{*}).

The SD-level pattern indicates that the consumer-services contraction documented at the municipality scale is a within-municipality, municipality-local response that averages out (or even partially reverses) when aggregated to the school- district level. This is consistent with workers leaving the municipality that hiked but staying within the SD, so SD-level consumer jobs do not contract.

Table 20: Industry decomposition — school district level

Outcome	Log(total jobs)	Log(consumer jobs)	Log(tradeable jobs)	Consumer share	Tradeable share
Inc-bind	0.033	0.124*	-0.094	0.023*	-0.042*
	(0.043)	(0.071)	(0.110)	(0.012)	(0.024)
Inc-becomes	0.047	0.045	0.010	-0.000	-0.013
	(0.030)	(0.031)	(0.089)	(0.009)	(0.020)
Inc-nobind	0.519*	0.337***	0.541***	-0.033	0.028
	(0.277)	(0.077)	(0.140)	(0.059)	(0.031)
Dec-bind	0.027	0.003	0.068	-0.007	0.005
	(0.107)	(0.081)	(0.201)	(0.007)	(0.018)
Dec-nobind	0.094***	0.184***	-0.116***	0.015***	-0.050***
	(0.019)	(0.032)	(0.028)	(0.005)	(0.010)
N	51,454	51,454	51,454	51,454	51,454
R^2	0.992	0.989	0.979	0.911	0.903
Within R^2	0.002	0.006	0.002	0.005	0.003
Dep. mean	8.531	7.046	7.130	0.245	0.276
Dep. SD	1.179	1.285	1.383	0.097	0.124

Notes: Cluster-robust standard errors in parentheses. Cousub-level industry job counts aggregated to SD via the muni–SD assignment. Rows omitted (no estimable coefficients in any outcome column): Dec-becomes-non. These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Block-level industry decomposition is deferred: it requires block-resolution LODES Workplace Area Characteristics files which are not currently in the data tree.

The pre-/post-Act-32 institutional break is itself an important heterogeneity exercise, but because it involves a full re-estimation of the cell × dose regression on pre-2012 and post-2012

subsamples at each estimating unit, all Act-32 results are consolidated in Section 7.2.

6.4 Summary

The results follow a benchmark–decomposition–mechanism narrative. Section 6.1 reports the pooled Inc/Dec benchmark, which approximates a conventional local-tax design but mixes binding and non-binding events. Section 6.2.1 reports the preferred causal estimate: at the block level, a 1 pp municipal resident-rate increase that binds reduces the log count of workers living in the block by -0.208^{***} ($\approx 19\%$), with a clean realized-dose pre-trend and a null Inc-nobind within-design placebo on the same count outcome. Section 6.2.2 reports the municipality-level estimates as policy-jurisdiction calibration; the municipality Inc-bind coefficient is larger but pre-trended, so it is not the causal headline. Section 6.3 documents the binding-cell employment composition shift—consumer-facing services contract while tradeables are insulated—as suggestive evidence of a local-services demand channel. Section 6.2.3 reports the school-district aggregation as a diagnostic non-result. Robustness and pre-trend diagnostics are in Section 7; full SD tables, block share outcomes, and the binary Sun–Abraham IW event study are in the Appendix.

7 Robustness

This section summarizes the battery of robustness checks behind the headline dose-response; full tables, figures, and detailed prose are collected in Appendix D. The realized-dose event-study pre-trend diagnostic (Table 10) was discussed alongside the headline in Section 6.2.1.

7.1 Sample Variant Robustness

Three sample variants for the DCED nonresident-rate panel (strict/filled/drop-incomplete) leave the headline Inc-bind coefficient on log residents working in own municipality essentially unchanged at -0.401^{***} (see Appendix D.3).

7.2 Pre/Post Act 32 Decomposition

Splitting the headline cell-×-dose specification at the Act-32 enforcement break (cohort first event year < 2012 vs ≥ 2012) shows the municipality-level binding-cell effect strengthens after consolidated employer withholding takes effect (-0.339^{**} pre vs -0.534^{***} post on log residents working in own municipality); the block-level effect is roughly constant pre and post, so the finer-geography sorting response predates the Act-32 enforcement break (see Appendix D.1).

7.3 Dose-Magnitude Linearity

Splitting Inc-bind units into terciles of total post-event dose and re-estimating the headline per tercile delivers essentially identical per-pp coefficients in the small and large terciles (-0.440 vs -0.417); the linearity assumption embedded in $\beta_k^s \Delta\tau_{it}^s$ is supported (see Appendix D.2).

7.4 Alternative Control Group: Clean-in-Window

Relaxing the NC-only control restriction to the Cengiz et al. (2019) clean-in-window definition (units with no events in the focal cohort's $[c-K, c+K]$ window) leaves all headline coefficients essentially unchanged on Inc-bind, Inc-becomes, and Inc-nobind cells across six outcomes (see Appendix D.4).

7.5 Rambachan–Roth Sensitivity to Pre-Trend Violations

I implement the **rambachanMoreCredibleApproach2023** relative-magnitudes sensitivity $\Delta^{RM}(\bar{M})$ on the Inc-bind realized-dose ES at each estimating unit. The block-level log home count headline survives strict parallel trends with $[-0.49, -0.12]$ at $\bar{M} = 0$, as do the cousub-level log live-and-work and share-live-in coefficients; all three lose significance between $\bar{M} = 0$ and $\bar{M} = 0.5$, while none of the SD-level Inc-bind residence-side coefficients survives even

$\bar{M} = 0$ (see Appendix D.6).

8 Discussion

The continuous stacked DiD design produces four substantive findings. The block-level model is the preferred causal estimate; the municipality-level model is the policy-jurisdiction magnitude calibration; the school-district-level aggregation is an aggregation diagnostic and is reported in the Companion Appendix.

The central finding is that sorting is identified through the binding channel. The Inc-bind cell—rate hikes that move the bill—reduces the log count of workers living in the block at the block level by approximately 19 percent per pp (-0.21 log points, SE = 0.054). The municipality-level analogue is larger (-0.40 log points on the log count of residents working in their own municipality, $\approx 33\%$ per pp). The Inc-nobind placebo—rate hikes whose post-event home stack remains below the nonresident rate—is statistically indistinguishable from zero on the block-level count outcomes (-0.056 on the log count of workers living in the block, -0.113 on the log count of workers at the block) and on the municipality-level the log count of residents working in their own municipality (-0.057). Inc-becomes lies between the two. The pattern matches the institutional rule's prediction: the same statutory rate change can produce a strong, moderate, or null behavioral response depending solely on whether the rate moves the bill. This is direct evidence that workers respond to taxes they actually pay rather than to statutory rates per se.

Block-level identification is preferred not because the joint Wald F-test is cleaner—it is not—but because the block observes both rate-setting jurisdictions at their native geography without a within-jurisdiction aggregation step. The two estimating units give opposite verdicts on the joint Wald: at the block level (N $\approx$ 21M) the joint test rejects parallel trends despite pointwise-null pre-event coefficients, reflecting hyper-precision rather than an economically meaningful violation; at the municipality level (N $\approx$ 198K) the joint test does not

reject ($p = 0.23$). The **rambachanMoreCredibleApproach2023** sensitivity bounds confirm the substantive conclusion under either reading: the block-level Inc-bind post-period CI at strict parallel trends is $[-0.49, -0.12]$, bounded comfortably away from zero. The block-level result is the defensible per-pp causal quantity; the municipality-level estimate is useful for back-of-envelope welfare calibration.

The joint regression also separates the two rate-setting channels. The municipality resident-rate Inc-bind coefficient on resident outcomes is large and consistent at both the block and municipality levels. The SD-resident-rate cells, estimated jointly in the same regression, produce smaller and direction-mixed effects: $+0.177^{*}$ at the municipality level on the log count of residents working in their own municipality, $+0.097$ on block-level log count of workers living in the block ($p = 0.12$), and $+0.066^{*}$ at the SD level (sign-reversed). The SD rate is not the dominant driver of residential sorting through the binding channel. This is a substantive change from an earlier draft that found a positive SD-rate "attraction" effect, which was an artifact of tract-centroid LODES aggregation; under block-level LODES with the unified two-rate-setter design, SD-rate effects on residents are small or wrong-signed.

The mechanism behind the workplace response appears to be local-services demand spillover rather than direct tax incidence on jobs. Industry-decomposed outcomes at the municipality level show consumer-facing jobs (retail, food, arts, other services) fall by approximately 26 percent in the Inc-bind cell (-0.30 log points, SE $= 0.15$) while tradeable jobs (manufacturing, wholesale, finance, professional services) are insulated (point estimate $+0.24$, statistically indistinguishable from zero). The composition shift is captured by the share of tradeable jobs rising approximately 11 percentage points in the binding cell. The decline in total jobs operates via local-services demand spillover from departing residents, not via direct nonresident-tax incidence.

The block-level Inc-bind estimate also has a direct revenue implication. Apply the envelope-theorem decomposition of Saez (2001): for a flat local rate τ on tax base WN

with residential semi-elasticity $\beta = \partial \ln N / \partial \tau$, the first-order effect of a small hike is

$$\frac{dR}{d\tau} \;=\; W \cdot N \cdot (1 + \tau\beta),$$

a mechanical gain WN minus the behavioral offset $\tau\beta$ from residents relocating. At the PA municipality baseline $\tau \approx 1\%$ and the block-level Inc-bind estimate $\beta \approx -21$ per unit rate (-0.21 log points per pp), $\tau\beta \approx -0.21$: the residential response offsets roughly 21% of the mechanical revenue gain. The hike remains revenue-positive, paralleling the Spanish regional finding in Agrawal and Foremny (2019).

The flat-tax revenue-maximizing rate, obtained by setting $dR/d\tau = 0$, is $\tau^* = 1/|\beta| \approx 4.8\%$. This is the analog, for a flat local tax with only an extensive-margin mobility response, of the top-marginal-rate formula $\tau^* = 1/(1 + a\varepsilon + \eta)$ in the optimal-taxation-with-mobility literature reviewed by Agrawal (2026), where a is the Pareto coefficient of top earners, ε is the intensive-margin elasticity of taxable income, and η is the extensive-margin mobility elasticity. In my flat-tax setting neither a nor ε enters: there is no top tail and no intensive labor-supply channel, only the mobility margin captured by β. PA's statutory EIT range (roughly 0.5–3.4%) sits entirely in the revenue-positive interior; the headline sorting estimate is consistent with a revenue gain to the raising municipality at the municipality boundary, with the welfare cost accruing elsewhere through relocation and the local-services demand spillover documented above.

Eight limitations merit acknowledgment. First, while the block-level stacked headline on the count outcomes is large, negative, and well-identified, the block-level share outcome (the share of block residents working in their own municipality) is null in the binding cell (-0.035, $p = 0.26$) and marginally significant in the Inc-nobind placebo (-0.037, $p = 0.015$). The share outcome at block scale appears to reflect mechanical denominator dynamics on very small populations rather than binding-state-specific behavior, so the share-based block evidence should be read as suggestive rather than identified. The count-based block evidence re-

mains the cleanest test. Second, the school-district aggregation produces direction-mixed coefficients and a marginally non-zero realized-dose pre-event lead; the SD-level results should be read as confirming the directional sign of the municipality-rate effect under a coarser aggregation rather than as standalone identification of an SD-channel response. Third, the LODES OD data covers PA municipalities but not cross-state inflows in the intra-state main file used here; the auxiliary file (cross-state inflows) is downloaded but not yet incorporated, which would tighten the estimates of in-commuters for border municipalities.

Fourth, the paper estimates the residential-sorting response to the home-side resident rate but does not separately estimate the work-side response to the municipality nonresident rate (τ^{nr}). The nonresident rate enters the design only through the binding-state classification of cells. A complementary work-side specification—estimating in-commuter and total-jobs responses to τ^{nr} events at the municipality level, with a parallel binding-state decomposition—would test whether the local-services demand spillover I document on the residence side (§6.3) is symmetric with a workplace-side response. I defer this to a separate paper.

Fifth, the design treats each treated municipality's stack as if its identifying variation were independent of the rate trajectories of its geographic neighbors. Pennsylvania municipalities are small and dense; a worker who moves out of a rate-raising municipality often relocates to an adjacent municipality whose own rate is also moving in some year of the panel. The headline specification absorbs this through never-changed-only controls and municipality-by-stack fixed effects, but it does not formally model spatial spillovers, so the per-pp coefficient is best read as a local displacement elasticity rather than a net-migration one. The **butts2023DiDSpatialSpillovers** estimator generalizes stacked DiD to allow for SUTVA violations from spatial spillovers across treated units, using an explicit spatial weighting matrix on the adjacency graph. Applying it to the PA cousub adjacency graph is left for future work.

Sixth, I do not currently estimate the headline using the heterogeneity-robust continuous-

treatment difference-in-differences estimator of **deChaisemartin2022Continuous**, which generalizes the de Chaisemartin–D'Haultfœuille framework (**deChaisemartin2020TWFE**; **deChaisemartin2023ImperfectPT**) to continuously distributed treatments at every period and leverages switchers (units that move both up and down in dose) for identification. An earlier iteration of this paper used a prototype CDH-continuous specification at the cousub level; the block-level analog is computationally prohibitive on the current stacked panel ($\sim$ 21M block-years $\times$ multiple cohorts) and is left for future work. Re-estimating the headline with `did_multiplegt_dyn` (**deChaisemartin2025MultipleGtDyn**)—once the package documentation and switcher-based identification at this scale are stable—would also let me exploit the rate-decrease (Dec) cohorts, currently too thin under conventional cluster-robust inference, as identification rather than just as a directional mirror.

Seventh, LODES is a stock dataset rather than a person-linked flow dataset. For each (residence block, workplace block, year), I observe a count of workers; I do not observe individual migration histories. The headline coefficient on the log count of workers living in the block is therefore a net elasticity—the combined effect of move-outs and move-ins—rather than separate estimates of the gross outflow and gross inflow margins. The paper's evidence that municipal rate hikes reduce the residential count is consistent both with workers leaving and with workers not arriving to replace those who left; LODES alone cannot decompose the two. Decomposing into gross outflow and gross inflow would require person-linked administrative data: the IRS Statistics of Income individual master file, the American Community Survey microdata with full residence-history linkage, or—most directly—the LEHD Job-to-Job Flows file, which tracks the same workers across employer transitions at quarterly frequency. All three are restricted-access products requiring approved-project status under the Census Bureau's Federal Statistical Research Data Centers (FSRDCs) or IRS Statistics of Income agreements; the microdata cannot be exported and must be analyzed at a physical FSRDC location. Replicating the present analysis on person-linked data is a clear extension but a substantially larger undertaking with its own access timeline.

Eighth, the paper does not estimate a comparable elasticity for the property tax, the other major local-revenue lever in Pennsylvania. The Oates (1969) capitalization theory predicts a substantially smaller residential-sorting elasticity for property tax than for the EIT, because property tax falls on owners and capitalizes into house prices rather than on workers directly. A Rambachan–Roth-style sensitivity bound on the property-tax elasticity would test the falsification implication "is the EIT effect specifically about the EIT, or any local tax raising?" but requires a 2002–2019 municipality-level millage panel that DCED does not publish (the DCED Annual Financial Report schema dropped per-municipality millage in 2006). Constructing this panel from county assessor records is left for future work.

9 Conclusion

When do local income taxes cause residential sorting? The institutional answer in Pennsylvania is that they cause sorting when the piecewise max-rule makes the statutory rate increase change the worker's actual bill. At the census block, the level at which both rate-setting jurisdictions enter the regression at their native geography, a 1 percentage-point municipal resident-rate increase in the Inc-bind cell reduces the log count of workers living in the block by roughly -0.21 log points, approximately 19%, with a clean realized-dose pre-trend ($\beta_{\text{pre}} = -0.027$, $|t| < 1$). The within-design Inc-nobind placebo, where the same statutory increase does not move the bill, is statistically null on the same count outcome. The response is therefore tied to tax liability rather than to the statutory rate change itself, and the binding-state decomposition is what distinguishes the two.

Municipality-level estimates are larger (-0.40 log points on the log count of residents working in their own municipality) but pre-trended, and they enter the paper as policy-jurisdiction magnitude calibration rather than as a second causal headline. School-district aggregations do not robustly detect a sorting response and are reported as a diagnostic non-

result. Industry decomposition shows that consumer-facing jobs fall in the binding cell while tradeable jobs are insulated, suggestive of a local-services demand spillover from departing residents rather than direct nonresident-tax incidence on jobs.

The narrow, credible claim is that local income taxes cause residential sorting when they bind. The same design framework—a continuous stacked DiD with binding-state cell decomposition, a realized-dose event study as a strong-parallel-trends diagnostic, and a within-design placebo cell whose theoretical prediction is zero—is portable to other settings where a piecewise tax rule generates predictable cross-sectional variation in whether a statutory rate change moves the worker's bill. Local sales taxes with use-tax-credit rules, county-level income-tax credits across state borders, and special tax districts with overlapping jurisdictions are candidate extensions.

References

Abowd, J. M., Ahearn, M. C., Andersson, F., Baldwin, J., Bartelsman, E., Bernard, A. B., Clayton, R. L., Davis, S. J., Dunne, T., Dunne, T., Jensen, J. B., & Roberts, M. J. (Eds.). (2009). *Producer Dynamics: New Evidence from Micro Data*. University of Chicago Press.

Agrawal, D. R. (2015). The tax gradient: Spatial aspects of fiscal competition. *American Economic Journal: Economic Policy*, *7*(2), 1–29. https://doi.org/10.1257/pol.20120360

Agrawal, D. R. (2026). Lessons on State and Local Income Taxes from the Twenty-first Century and Challenges for the Future. *CESifo Working Papers*.

Agrawal, D. R., & Foremny, D. (2019). Relocation of the Rich: Migration in Response to Top Tax Rate Changes from Spanish Reforms. *The Review of Economics and Statistics*, *101*(2), 214–232. https://doi.org/10.1162/rest_a_00764

Agrawal, D. R., Hoyt, W. H., & Wilson, J. D. (2022). Local Policy Choice: Theory and Empirics. *Journal of Economic Literature*, *60*(4), 1378–1455. https://doi.org/10.1257/jel.20201490

Bagchi, S. (2022). The Effects of Introducing Withholding on Tax Compliance: Evidence from Pennsylvania's Local Earned Income Tax. *SSRN Electronic Journal*. https://doi.org/10.2139/ssrn.4291598

Callaway, B., & Sant'Anna, P. H. C. (2021). Difference-in-differences with multiple time periods. *Journal of Econometrics*, *225*(2), 200–230. https://doi.org/10.1016/j.jeconom.2020.12.001

Cengiz, D., Dube, A., Lindner, A., & Zipperer, B. (2019). The Effect of Minimum Wages on Low-Wage Jobs*. *The Quarterly Journal of Economics*, *134*(3), 1405–1454. https://doi.org/10.1093/qje/qjz014

Coomes, P. A., & Hoyt, W. H. (2008). Income taxes and the destination of movers to multi-state MSAs. *Journal of Urban Economics*, *63*(3), 920–937. https://doi.org/10.1016/j.jue.2007.07.004

de Chaisemartin, C., D'Haultfœuille, X., & Vazquez-Bare, G. (2024). Difference-in-difference estimators with continuous treatments and no stayers. *AEA Papers and Proceedings*, *114*, 393–397. https://doi.org/10.1257/pandp.20241049

Deshpande, M., & Li, Y. (2019). Who Is Screened Out? Application Costs and the Targeting of Disability Programs. *American Economic Journal: Economic Policy*, *11*(4), 213–248. https://doi.org/10.1257/pol.20180076

Fadlon, I., & Nielsen, T. H. (2021). Family Labor Supply Responses to Severe Health Shocks: Evidence from Danish Administrative Records. *American Economic Journal: Applied Economics*, *13*(3), 1–30. https://doi.org/10.1257/app.20170604

Goodman-Bacon, A. (2021). Difference-in-differences with variation in treatment timing. *Journal of Econometrics*, *225*(2), 254–277. https://doi.org/10.1016/j.jeconom.2021.03.014

Graham, M. R., Kutzbach, M. J., & McKenzie, B. (n.d.). DESIGN COMPARISON OF LODES AND ACS COMMUTING DATA PRODUCTS.

Kaplan, J. (2024). Jacob Kaplan's Concatenated Files: Uniform Crime Reporting (UCR) Program Data: Offenses Known and Clearances by Arrest, 1960–Present. https://doi.org/10.3886/E100707V18

Kleven, H., Landais, C., Muñoz, M., & Stantcheva, S. (2020). Taxation and migration: Evidence and policy implications. *Journal of Economic Perspectives*, *34*(2), 119–142. https://doi.org/10.1257/jep.34.2.119

Moretti, E., & Wilson, D. J. (2017). The Effect of State Taxes on the Geographical Location of Top Earners: Evidence from Star Scientists. *American Economic Review*, *107*(7), 1858–1903. https://doi.org/10.1257/aer.20150508

Pennsylvania Department of Community and Economic Development. (2024a). MunStats portal: Municipal Statistics and Local Government Data.

Pennsylvania Department of Community and Economic Development. (2024b). Statewide Municipal Annual Financial Report (AFR) Database.

Pennsylvania General Assembly. (1932). Sterling Act, Act of August 5, 1932, P.L. 45.

Pennsylvania General Assembly. (1965). Local Tax Enabling Act, Act 511 of 1965.

Pennsylvania General Assembly. (1972). Home Rule Charter and Optional Plans Law.

Pennsylvania General Assembly. (1987). Municipalities Financial Recovery Act, Act 47 of 1987.

Pennsylvania General Assembly. (2006). Taxpayer Relief Act, Act 1 of 2006 (Special Session).

Pennsylvania General Assembly. (2008). Act 32 of 2008.

Rohlin, S., Rosenthal, S. S., & Ross, A. (2014). Tax avoidance and business location in a state border model. *Journal of Urban Economics*, *83*, 34–49. https://doi.org/10.1016/j.jue.2014.06.003

Saez, E. (2001). Using Elasticities to Derive Optimal Income Tax Rates. *The Review of Economic Studies*, *68*(1), 205–229. https://doi.org/10.1111/1467-937X.00166

Sun, L., & Abraham, S. (2021). Estimating dynamic treatment effects in event studies with heterogeneous treatment effects. *Journal of Econometrics*, *225*(2), 175–199. https://doi.org/10.1016/j.jeconom.2020.09.006

Tiebout, C. M. (1956). A pure theory of local expenditures. *Journal of Political Economy*, *64*(5), 416–424. https://doi.org/10.1086/257839

U.S. Census Bureau. (2020). TIGER/Line Shapefiles, 2020.

U.S. Census Bureau. (2024a). Annual Survey of School System Finances (F-33).

U.S. Census Bureau. (2024b). Building Permits Survey (BPS).

U.S. Census Bureau. (2024c). LEHD Origin-Destination Employment Statistics (LODES), Version 8.

Yang, L. (, & Heim, B. T. (2017). RESPONSIVENESS OF INCOME TO LOCAL INCOME TAXES: EVIDENCE FROM INDIANA. *National Tax Journal*, *70*(2), 367–391. https://doi.org/10.17310/ntj.2017.2.05

Young, C., Varner, C., Lurie, I. Z., & Prisinzano, R. (2016). Millionaire migration and taxation of the elite: Evidence from administrative data. *American Sociological Review*, *81*(3), 421–446. https://doi.org/10.1177/0003122416639625

A Pooled Inc/Dec Benchmark

This appendix collects the full pooled-aggregate (Inc/Dec) benchmark estimates and event-study diagnostics by estimating unit. The pooled specification is the symmetric Inc/Dec aggregate that precedes the binding-state decomposition reported in Appendix B.

A.1 Block-level pooled benchmark

The block-level per-pp pooled Inc/Dec specification (Equation 3) is reported as Table 7 in Section 6.1. Reading the table by sign and rate-setter: the municipality resident-rate increase reduces the log count of workers living in the block; the municipality resident-rate decrease

moves the same outcome in the opposite direction with a smaller magnitude; the SD rate is essentially null on both rate-setters. The pooled specification mixes binding and non-binding events on each rate-setter and so the magnitudes are dose-attenuated relative to the binding-state-decomposed estimate (Appendix B.1); the pooled benchmark is reported as the conventional local-tax literature comparator and as motivation for the cell decomposition. The binary-event analog of this specification is in Appendix A.5 (Table 21); coefficients agree in sign with the per-pp benchmark at every level.

A.2 Municipality-level pooled benchmark

The municipality-level per-pp pooled Inc/Dec specification is reported as Table 8 in Section 6.1. At the municipality level the pooled Inc coefficient on `log(live_and_work)` is sizable and negative; sign-symmetry with the Dec coefficient is partial (asymmetric magnitudes), and the residence-side outcomes (`share_live_in`, `log_total_residents`) move in the same direction. The pooled municipality Inc coefficient understates the binding-cell effect because it mixes Inc-bind (binding incidence) with Inc-becomes and Inc-nobind (partial / no incidence); the decomposed cells in Appendix B.2 show this directly. The binary analog is Table 22.

A.3 School-district-level pooled benchmark

The school-district-level per-pp pooled Inc/Dec specification is reported as Table 9 in Section 6.1. SD-level identification is weaker than at the municipality and block levels: the pooled Inc coefficient on residence-side outcomes is small and sign-reversed relative to the municipality and block estimates, the SD pooled Dec coefficient is also direction-mixed, and SD-level realized-dose pre-trends fail parallel trends on `log_live_and_work` ($\beta_{\text{pre}} = +0.061$, $|t| = 1.74$; see Appendix B.4). The SD-aggregation is reported as an aggregation diagnostic non-result and is not the load-bearing causal estimate. The binary analog is Table 23.

A.4 Realized-dose ES benchmark diagnostics

The realized-dose Sun–Abraham specification tests strong parallel trends across treatment-intensity magnitudes by interacting relative-event-time indicators with each unit's eventual cumulative dose. The pooled-cohort versions of these figures are collected here. (Per-cell decomposition versions appear in Appendix B.4.)

A.5 Binary ES benchmark diagnostics

The binary-event analog of the per-pp pooled specification (Equation 3) replaces each $\Delta\tau_{it}^{s}$ dose interaction with a post-treatment indicator that equals one in the event year and after. The four reported coefficients per regression are $\{\text{everInc}^{M}, \text{everDec}^{M}, \text{everInc}^{SD}, \text{everDec}^{SD}\}$. Sign-wise the binary estimates agree with the per-pp pooled headline at every level: municipality-rate increases reduce residence-side outcomes; SD-rate coefficients are small or sign-mixed. The binary analog complements the per-pp pooled specification by not weighting events by dose magnitude; cohort-event-time IW pooled estimates are reported by estimating unit below.

Table 21: Pooled Inc/Dec headline — block level (binary event)

Outcome	Log(workers live)	Log(workers work)	Share live+work in muni
Municipality rate (τ^M)			
$\Delta\tau > 0$ (Inc)	-0.052**	-0.038**	-0.018**
	(0.025)	(0.016)	(0.008)
$\Delta\tau < 0$ (Dec)	-0.050***	-0.028	0.013***
	(0.017)	(0.053)	(0.004)
School district rate (τ^{SD})			
$\Delta\tau > 0$ (Inc)	0.003	-0.000	0.001
	(0.006)	(0.007)	(0.001)
$\Delta\tau < 0$ (Dec)	0.008	0.066***	0.005***
	(0.006)	(0.015)	(0.001)
N	23,567,655	8,751,581	23,567,655
R^2	0.858	0.865	0.563
Within R^2	0.000	0.000	0.000
Dep. mean	2.238	2.318	0.121
Dep. SD	1.186	1.655	0.184

Notes: Binary post-treatment indicator: 1 in the event year and after. Cluster-robust standard errors in parentheses. Each regression includes both rate-setters' direction-pooled treatments jointly; the four reported coefficients are β^M_{Inc}, β^M_{Dec}, β^{SD}_{Inc}, β^{SD}_{Dec}. $^*p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 22: Pooled Inc/Dec headline — municipality level (binary event)

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Municipality rate (τ^M)				
$\Delta\tau > 0$ (Inc)	-0.117***	-0.099***	-0.000	-0.012***
	(0.034)	(0.034)	(0.020)	(0.004)
$\Delta\tau < 0$ (Dec)	-0.141*	-0.049	0.038	-0.010**
	(0.074)	(0.038)	(0.058)	(0.005)
School district rate (τ^{SD})				
$\Delta\tau > 0$ (Inc)	0.031	0.062**	0.015***	0.001
	(0.020)	(0.026)	(0.006)	(0.001)
$\Delta\tau < 0$ (Dec)	-0.007	0.023	-0.013	-0.002
	(0.078)	(0.114)	(0.016)	(0.003)
N	218,677	239,023	242,722	239,033
R^2	0.964	0.959	0.981	0.823
Within R^2	0.001	0.002	0.001	0.001
Dep. mean	3.677	5.847	6.697	0.093
Dep. SD	1.934	1.887	1.299	0.076

Notes: Binary post-treatment indicator: 1 in the event year and after. Cluster-robust standard errors in parentheses. Each regression includes both rate-setters' direction-pooled treatments jointly; the four reported coefficients are β^M_{Inc}, β^M_{Dec}, β^{SD}_{Inc}, β^{SD}_{Dec}. $^*p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 23: Pooled Inc/Dec headline — school district level (binary event)

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
School district rate (τ^{SD})				
$\Delta\tau > 0$ (Inc)	0.089*	0.091**	0.129***	0.000
	(0.051)	(0.046)	(0.035)	(0.003)
$\Delta\tau < 0$ (Dec)	-0.007	-0.072***	-0.069***	0.008*
	(0.042)	(0.020)	(0.023)	(0.004)
N	51,440	51,474	51,474	51,474
R^2	0.986	0.987	0.987	0.953
Within R^2	0.004	0.007	0.006	0.006
Dep. mean	6.260	8.335	8.768	0.125
Dep. SD	1.298	1.202	0.909	0.068

Notes: Binary post-treatment indicator: 1 in the event year and after. Cluster-robust standard errors in parentheses. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

B Binding-State Decomposition

This appendix collects the full set of binding-state decomposition results and event-study diagnostics by estimating unit. The decomposition splits each rate-setter into the six cells (Inc-bind, Inc-becomes, Inc-nobind, Dec-bind, Dec-becomes-non, Dec-nobind) described in Section 3.

B.1 Block-level decomposition

The block-level binding-state decomposition is the preferred causal estimate of the paper. The main-text reports Table 14 (municipality-rate cells) and Table 15 (SD-rate cells) on the count outcomes `log_home_count` and `log_work_count` in Section 6.2.1. A 1 pp Inc-bind increase in the municipality resident rate reduces log(home_count) by -0.208^{***} ($\approx 19\%$), with a clean realized-dose pre-trend ($\beta_{\text{pre}} = -0.027$, $|t| = 0.51$; Appendix B.4) and a null Inc-nobind within-design placebo on the same count outcome.

Block-level share outcomes (supporting). The main paper reports block-level count outcomes (`log_home_count`, `log_work_count`) as the primary residential-sorting evidence. Block-level share outcomes are reported here for completeness. Block geography has small and sometimes sparse denominators, and both the numerator and denominator can move with the same residential decision, which makes share estimates mechanically fragile at this scale. I do not use the share-outcome coefficients to identify the headline causal effect; the count coefficients in the main text are the preferred estimands.

Table 24: Block-level dose-response: muni resident-rate cells (share outcomes, supporting)

Outcome	Share live+work in muni
Inc-bind	-0.035
	(0.031)
Inc-becomes	-0.043***
	(0.010)
Inc-nobind	-0.037**
	(0.015)
Dec-bind	-0.049
	(0.031)
Dec-becomes-non	-0.080**
	(0.033)
Dec-nobind	0.006
	(0.006)
N	23,567,655
R^2	0.563
Within R^2	0.000
Dep. mean	0.121
Dep. SD	0.184

Notes: Cluster-robust standard errors in parentheses. Continuous stacked-DiD with block-by-stack and year-by-stack fixed effects; standard errors two-way clustered at the muni and SD levels. Each block is exposed to one municipality and one school district; both rate-setters' policy variation enters at its native level. Block-level share outcomes are reported for completeness. Sparse denominators and endogenous denominator dynamics make these estimates mechanically fragile at block geography; they are not used to identify the headline causal effect. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 25: Block-level dose-response: SD resident-rate cells (share outcomes, supporting)

Outcome	Share live+work in muni
Inc-bind	0.006*
	(0.003)
Inc-becomes	-0.002
	(0.006)
Dec-bind	-0.018**
	(0.007)
Dec-becomes-non	-0.007***
	(0.003)
Dec-nobind	-0.034**
	(0.017)
N	23,567,655
R^2	0.563
Within R^2	0.000
Dep. mean	0.121
Dep. SD	0.184

Notes: Cluster-robust standard errors in parentheses. Continuous stacked-DiD with block-by-stack and year-by-stack fixed effects; standard errors two-way clustered at the muni and SD levels. Each block is exposed to one municipality and one school district; both rate-setters' policy variation enters at its native level. Block-level share outcomes are reported for completeness; not used for headline identification. Rows omitted (no estimable coefficients in any outcome column): Inc-nobind. These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

B.2 Municipality-level decomposition

The municipality-level binding-state decomposition is the policy-jurisdiction calibration estimate. The main-text reports Table 16 (municipality-rate cells) and Table 17 (SD-rate

cells) in Section 6.2.2; this subsection retains the full detailed Inc-cell coefficients with regression diagnostics, the thin Dec-cell coefficients, the per-outcome diagnostics table, the full single-regression specification (with all controls), and heterogeneity by municipality type.

Full Inc-cell coefficients with diagnostics. Table 26 reports the municipality-resident-rate Inc-cell coefficients across the eight outcomes. Each row is a separate regression of the form in Equation (4); the displayed columns are the municipality rate-setter Inc-bind, Inc-becomes, and Inc-nobind cells. SD and nonresident-rate cell coefficients are estimated jointly in the same regression and are reported in Table 29. Dec-cell coefficients are in Table 27.

Table 26: Headline stacked-continuous DiD: municipality-rate Inc cells with diagnostics

Outcome	Inc-bind (n_{muni}=40)	Inc-becomes (n_{muni}=58)	Inc-nobind (placebo) (n_{muni}=26)
Coefficients (per 1 pp of $\Delta\tau^{muni}$)			
log(live_and_work)	-0.402^{***}	-0.181^{**}	-0.062
	(0.096)	(0.072)	(0.169)
log(in_commuters)	-0.213^{**}	-0.197^{***}	-0.187
	(0.083)	(0.054)	(0.204)
log(out_commuters)	-0.053	-0.042	$+0.217^{**}$
	(0.038)	(0.029)	(0.091)
log(total_jobs)	-0.270^{***}	-0.212^{***}	-0.201
	(0.086)	(0.054)	(0.203)
log(total_residents)	-0.100^{***}	-0.068^{*}	$+0.194^{**}$
	(0.033)	(0.040)	(0.095)
share_live_in	-0.045^{***}	-0.012	-0.012
	(0.010)	(0.009)	(0.008)
share_work_outside	$+0.041^{***}$	$+0.020^{***}$	$+0.017$
	(0.016)	(0.007)	(0.014)
Diagnostics (log_live_and_work regression)			
N	218,685		
R^2	0.964		
Within R^2	0.0007		
Dep. var. mean (log_live_and_work)	3.677		
Dep. var. SD	1.934		
Fixed effects	stack-municipality + stack-year		
Clusters (municipality level)	1,945		

Notes: Stacked continuous DiD, Equation (4), drop-incomplete sample with NC controls strictly defined as municipalities with no rate change on the relevant rate-setter. Each municipality in a treated cell appears in one sub-experiment (its own first-event cohort); NC municipalities appear in every sub-experiment. Only municipality rate-setter Inc cells are displayed here; SD and nonresident cells are estimated jointly in the same regression and reported in Table 29. Each regression also includes the share-based control vector (police, streets, debt service / total expenditures; EIT, real estate, intergov

Inc-bind: the binding-cell effect. A 1 percentage-point increase in the municipality resident rate in the Inc-bind cell (where the bill rises by the full $\Delta\tau$ for in-municipality residents) reduces `live_and_work` by roughly 33 percent (a -0.40 log-point coefficient), `in_commuters` by roughly 19 percent (-0.21 log points), `total_jobs` by roughly 24 percent (-0.27 log points), and `total_residents` by roughly 10 percent (-0.10 log points). The dependent variable mean on `log(live_and_work)` is 3.68, so the typical municipality has $e^{3.68} \approx 40$ live-and-work workers; a 1 pp hike at this mean predicts $40 \cdot (e^{-0.402} - 1) \approx -13$ live-and-work workers. Actual rate hikes in the Inc-bind cell average 0.27 pp, which scales the predicted effect down to roughly $0.27 \cdot 40 \cdot (e^{-0.402} - 1) \approx -3.6$ workers per typical event over 5 years.

Inc-becomes: the bill-flip cell. Inc-becomes events show smaller but still significant effects: -0.181 on `live_and_work`, -0.212 on `total_jobs`, -0.068 on `total_residents`. The attenuation relative to Inc-bind is consistent with the bill rising by less than the full $\Delta\tau$ in this cell—only the portion of the rate change above the work nonresident rate is incident on the worker.

Dec cells: thin and noisy. Dec-cell coefficients (Table 27) are reported for completeness but are not robust to specification. Dec-bind has 11 municipalities; Dec-becomes-non has 4; Dec-nobind has 15. I do not draw asymmetric-response conclusions from these cells.

Table 27: Headline stacked-continuous DiD: municipality-rate Dec cells (thin)

Outcome	Dec-bind (n_{muni}=11)	Dec-becomes-non (n_{muni}=4)	Dec-nobind (n_{muni}=15)
log(live_and_work)	$+0.241^{***}$ (0.078)	-0.260 (0.272)	$+0.332^{**}$ (0.165)
log(in_commuters)	$+0.052$ (0.080)	$+0.496$ (0.339)	$+0.013$ (0.121)
log(out_commuters)	-0.068 (0.061)	-0.628 (0.416)	$+0.132$ (0.101)
log(total_jobs)	$+0.111$ (0.073)	$+0.420$ (0.265)	$+0.022$ (0.116)
log(total_residents)	-0.056 (0.043)	-0.651 (0.420)	$+0.135$ (0.099)
share_live_in	$+0.054^{*}$ (0.030)	-0.071 (0.095)	$+0.012$ (0.030)
share_work_outside	-0.011 (0.013)	$+0.014$ (0.017)	-0.003^{*} (0.001)

Notes: Same regression as Table 26. Dec cells are thin and noisy; I do not draw causal conclusions from these coefficients individually. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Per-outcome diagnostics. Table 28 reports N, R^2, within-R^2, and the dependent-variable mean for each headline regression.

Table 28: Headline regressions: per-outcome diagnostics

Outcome	N	R^2	Within R^2	Dep. mean	Dep. SD
log(live_and_work)	218,685	0.964	0.0007	3.677	1.934
log(in_commuters)	239,031	0.959	0.0011	5.847	1.887
log(out_commuters)	242,730	0.979	0.0009	6.636	1.263
log(total_jobs)	239,041	0.961	0.0011	5.948	1.901
log(total_residents)	242,730	0.981	0.0010	6.697	1.299
share_live_in	239,041	0.823	0.0010	0.093	0.076
share_work_outside	242,730	0.945	0.0008	0.943	0.066
share_nonres	239,041	0.823	0.0010	0.907	0.076

Notes: Same regression structure as Table 26. R^2 is the standard within-stack coefficient of determination; within-R^2 is the variation explained by the cell×dose interactions and controls, net of the municipality-stack and year-stack fixed effects. Dep. mean and SD computed over the regression sample.

Full main specification with control estimates. Table 29 reports all coefficients from the headline regression on `log(live_and_work)`: municipality rate-setter cells, SD cells, nonresident cells, and control coefficients.

Table 29: Full main spec: log(live_and_work) on all treatments + controls

Term	Coefficient	SE
Muni resident-rate cells (vs NC^{muni})		
Inc-bind $\times\Delta\tau^{\text{muni}}$	-0.402^{***}	0.096
Inc-becomes $\times\Delta\tau^{\text{muni}}$	-0.181^{**}	0.072
Inc-nobind $\times\Delta\tau^{\text{muni}}$	-0.062	0.169
Dec-bind $\times\Delta\tau^{\text{muni}}$	$+0.241^{***}$	0.078
Dec-becomes-non $\times\Delta\tau^{\text{muni}}$	-0.260	0.272
Dec-nobind $\times\Delta\tau^{\text{muni}}$	$+0.332^{**}$	0.165
School district rate cells (vs NC^{SD})		
Inc-bind $\times\Delta\tau^{\text{SD}}$	$+0.177^{*}$	0.099
Inc-becomes $\times\Delta\tau^{\text{SD}}$	$+0.103$	0.090
Dec-bind $\times\Delta\tau^{\text{SD}}$	-0.721^{***}	0.191
Dec-becomes-non $\times\Delta\tau^{\text{SD}}$	-0.223^{**}	0.109
Dec-nobind $\times\Delta\tau^{\text{SD}}$	$+0.048$	0.179
Share-based controls		
share_police	$+0.097$	0.100
share_streets	$+0.049$	0.039
share_debt_svc	-0.058	0.051
share_eit	-0.022	0.085
share_re	-0.182^{**}	0.087
share_intergov	$+0.016$	0.044
Diagnostics		
N	218,685	
R^2	0.964	
Within R^2	0.0007	
Dep. mean (log_live_and_work)	3.677	
Dep. SD	1.934	
Fixed effects	stack-municipality + stack-year	
Clusters (municipality)	1,945	

Notes: All coefficients from a single regression of log(live_and_work) on cell-specific dose interactions

Heterogeneity by municipality type. The headline municipality-level coefficient is concentrated in 8 post-industrial PA cities (Wilkes-Barre, Hazleton, Johnstown, Jeannette, Coatesville, Warren, Carbondale, Farrell). Table 30 reports the municipality-Inc-bind cell coefficient on `log(live_and_work)` by municipality type.

Table 30: Inc-bind cell by municipality type (drop sample, post-NC-fix)

Type	n	Mean residents	Inc-bind coef	Sig
City	8	6,904	−0.455	** (p=0.021)
Non-city (borough + township)	32	3,207	−0.194	— (p=0.189)
Borough	10	2,631	−0.213	— (p=0.173)
Township	22	3,468	−0.161	— (p=0.438)
Pooled (headline)	40	—	−0.402***	—

Notes: Stacked-continuous DiD with same controls and FE structure as the municipality-level headline. Each row restricts the sample to its own municipality type; NC controls are restricted to the same type. Cluster-robust SEs by municipality. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

The 8 cities drive the municipality-level headline. The 32 borough and township Inc-bind municipalities show direction-consistent but individually insignificant effects.

B.3 School-district-level decomposition

Table 31 reports the SD-level dose-response coefficients. The unit of observation is SD-year. The SD's own resident rate enters the regression with binding-cell decomposition; the within-SD aggregate of constituent municipality rates enters as a control; and SD-level fiscal share controls constructed from Census F-33 data control for SD fiscal posture. Binding-state classification at the SD level uses the within-SD mean of `muni_nonres_eit` (the work-side nonresident rate) as the comparator, aggregating the institutional rate field across the SD's constituent municipalities (the previous draft used the within-SD mean resident rate as a

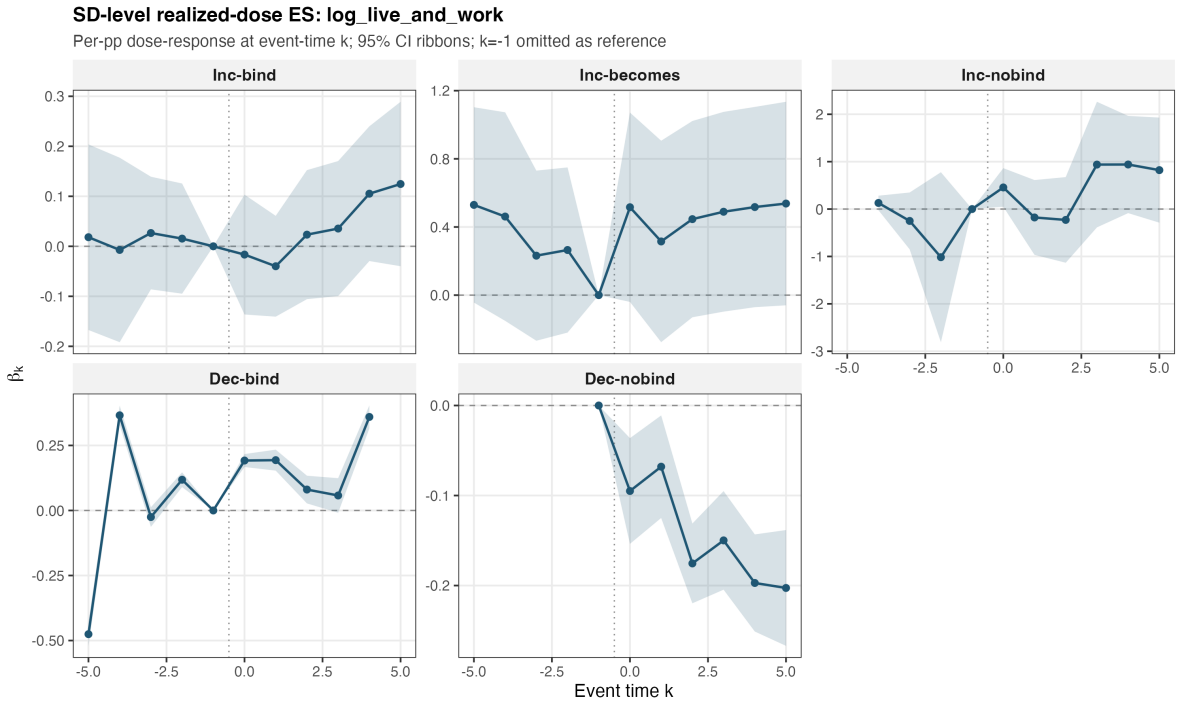
proxy; that proxy has been retired and the cell labels reproduce essentially unchanged under the corrected definition).

Table 31: School-district-level dose-response by binding-state cell

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind	0.027	0.057	0.106***	0.003
	(0.038)	(0.046)	(0.030)	(0.004)
Inc-becomes	0.271*	0.184	0.318**	0.004
	(0.154)	(0.137)	(0.131)	(0.006)
Inc-nobind	0.793	1.149*	0.422***	-0.035
	(0.506)	(0.633)	(0.148)	(0.023)
Dec-bind	0.060	0.050	0.178***	0.000
	(0.077)	(0.111)	(0.045)	(0.004)
Dec-nobind	-0.126***	0.154***	0.076***	-0.032***
	(0.025)	(0.028)	(0.016)	(0.002)
N	51,420	51,454	51,454	51,454
R^2	0.986	0.987	0.987	0.953
Within R^2	0.004	0.004	0.005	0.004
Dep. mean	6.260	8.335	8.769	0.125
Dep. SD	1.298	1.202	0.909	0.068

Notes: Cluster-robust standard errors in parentheses. Continuous stacked-DiD with SD-by-stack and year-by-stack fixed effects; standard errors clustered at the school-district level. Unit of observation is SD×year. Binding-cell classification uses the within-SD mean of $\tau^{\text{nr,muni}}$ as the work-side comparator, aggregating the institutional nonresident-rate field over the SD's constituent munis. Rows omitted (no estimable coefficients in any outcome column): Dec-becomes-non. These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Figure 19: SD-level realized-dose ES on `log(live_and_work)` by cell.



Notes: per-pp realized-dose coefficient at each event-time, SD-level stacked-continuous DiD. The SD-level results are reported for completeness. They show that sorting is not robustly detected at the school-district aggregation: the SD-resident-rate Inc-bind coefficient is positive in this aggregation, sign-reversed relative to the municipality- and block-level findings, and the pre-event lead is marginally non-zero on `log(live_and_work)` ($\beta_{\text{pre}} = +0.061$, $|t| = 1.74$). These estimates should not be interpreted as standalone evidence of an SD-channel residential response.

B.4 Realized-dose ES decomposition diagnostics

Block level, by cell. The block-level realized-dose ES on the log count of workers living in the block (Figure 15 in the main text) is the diagnostic on which the headline causal claim rests; the remaining block-level outcomes are reported here.

Figure 20: Block-level realized-dose ES on the log count of workers working in the block, faceted by binding-state cell; municipality-cohort (red) and SD-cohort (blue).

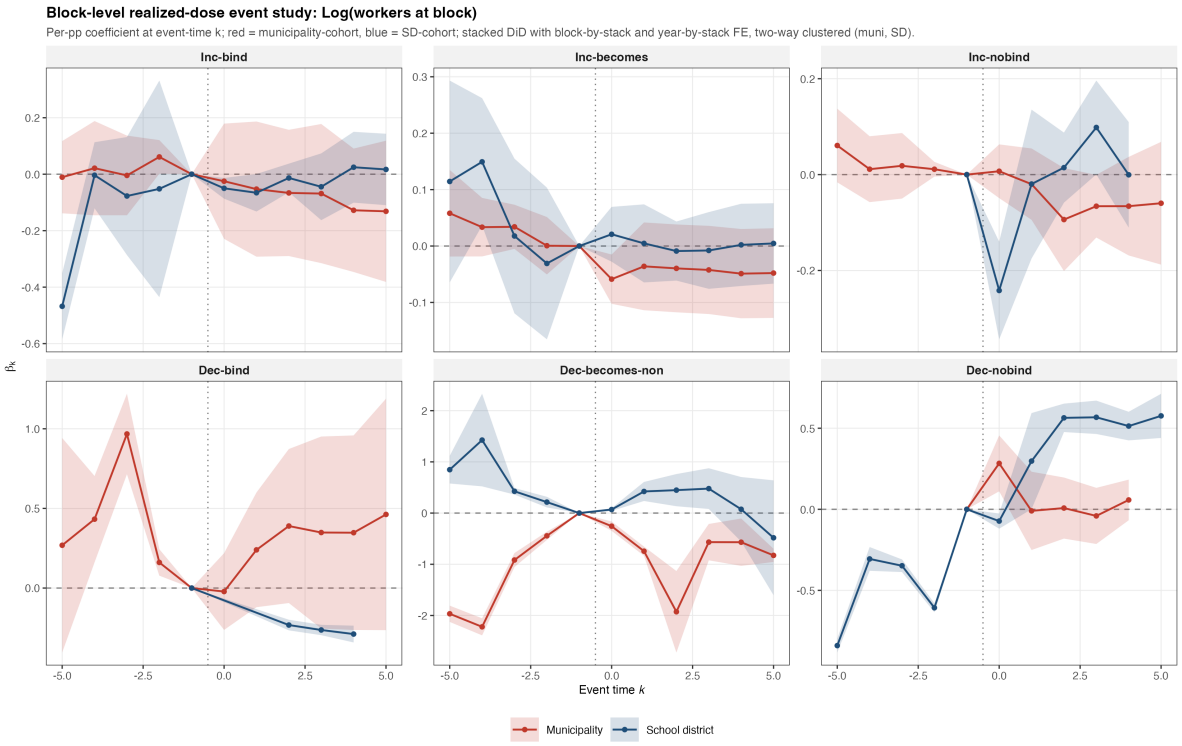
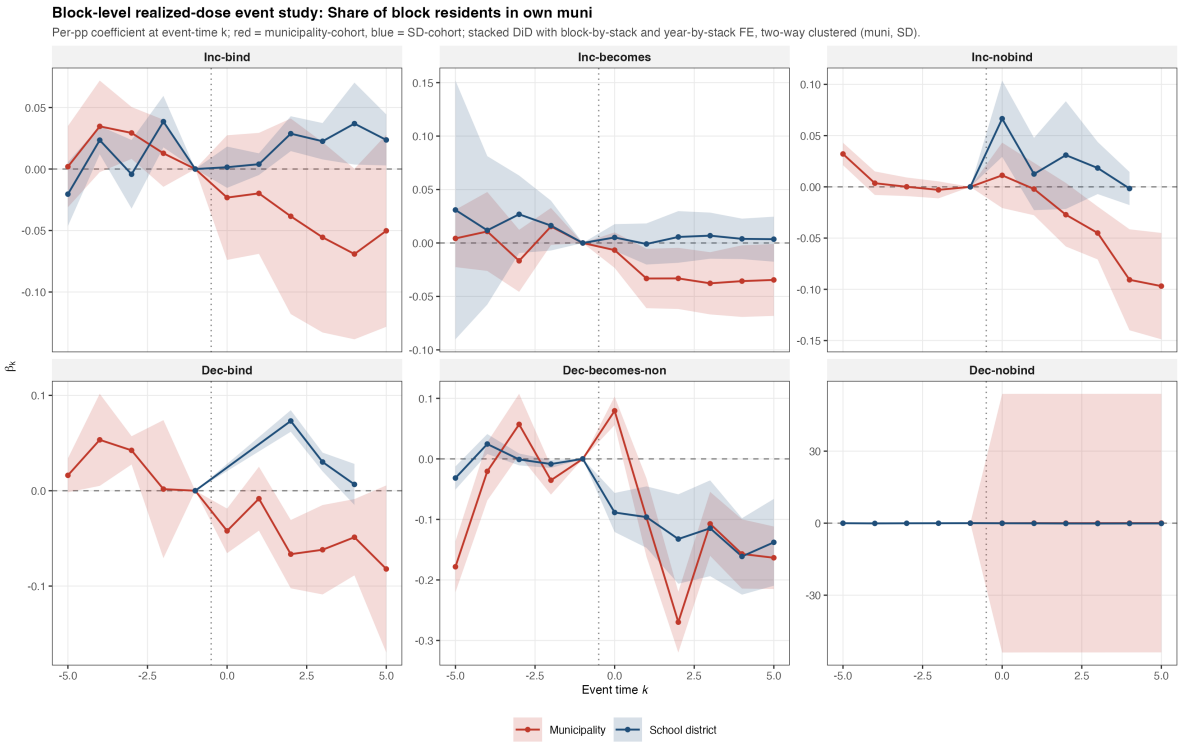


Figure 21: Block-level realized-dose ES on the share of block residents working in their own municipality, faceted by binding-state cell.



Municipality level, by cell. The municipality-level live-and-work figure (Figure 17) appears in the main text; the remaining residence-side and commuter outcomes are reported here.

Figure 22: Muni-level realized-dose ES on log(total residents), by cell.

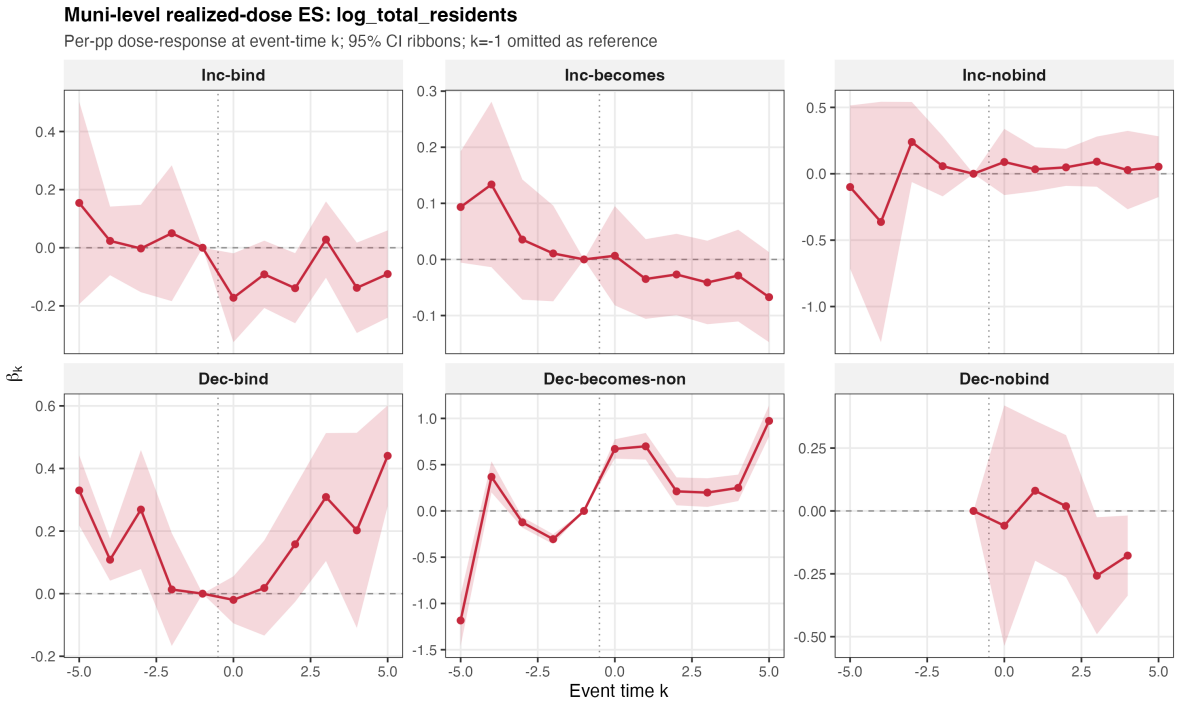


Figure 23: Muni-level realized-dose ES on log(in-commuters), by cell.

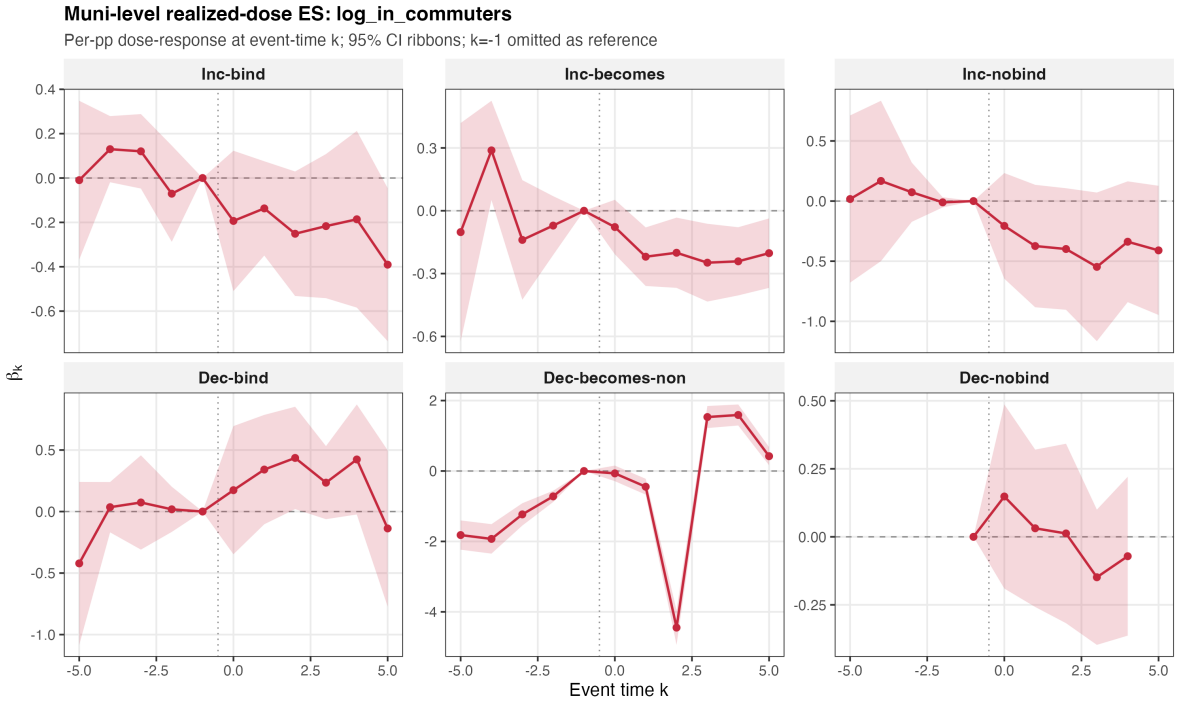
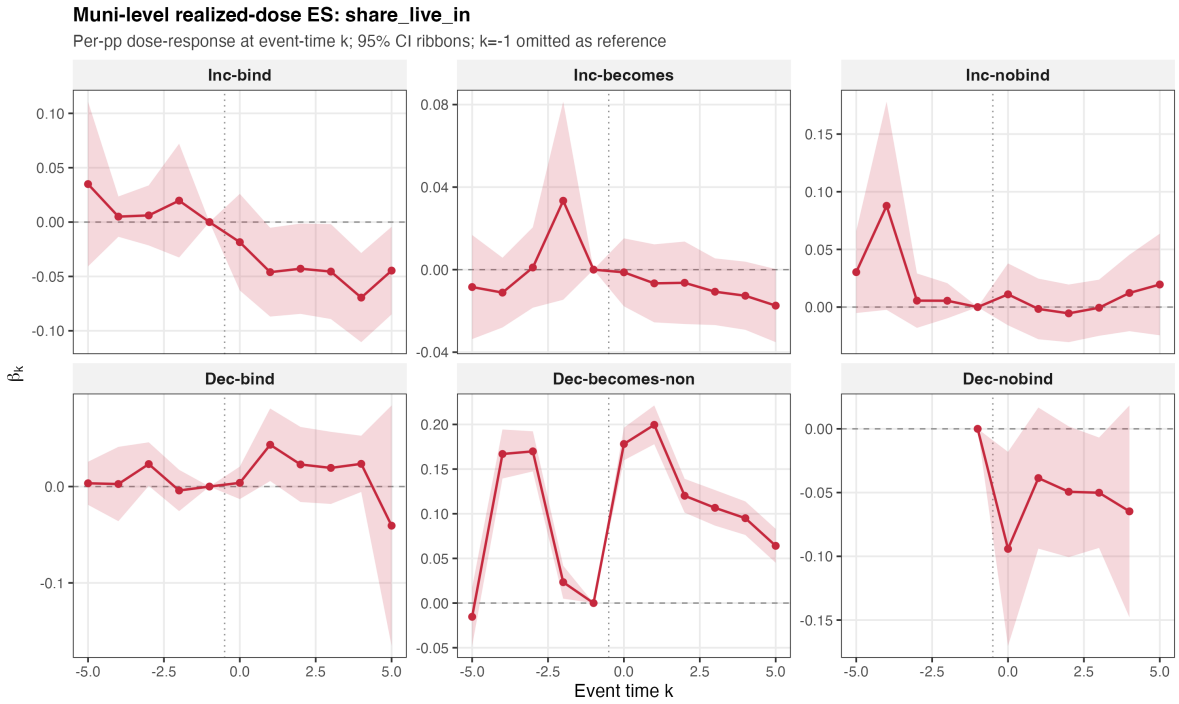


Figure 24: Muni-level realized-dose ES on the share of workers living in their own municipality, by cell.



School-district level, by cell. The SD-level realized-dose ES, run on the same residence-side outcomes aggregated to the school district, is the input to the SD aggregation diagnostic in Section 6.2.3.

Figure 25: SD-level realized-dose ES on log(live_and_work), by cell.

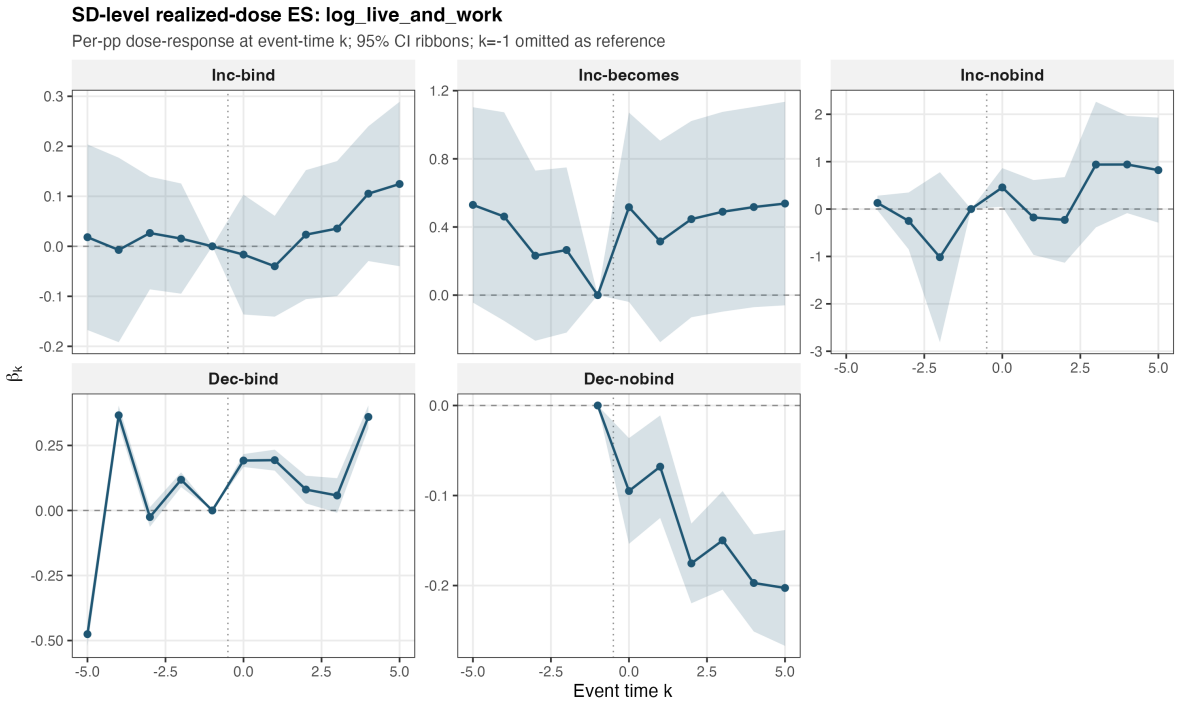


Figure 26: SD-level realized-dose ES on log(total residents), by cell.

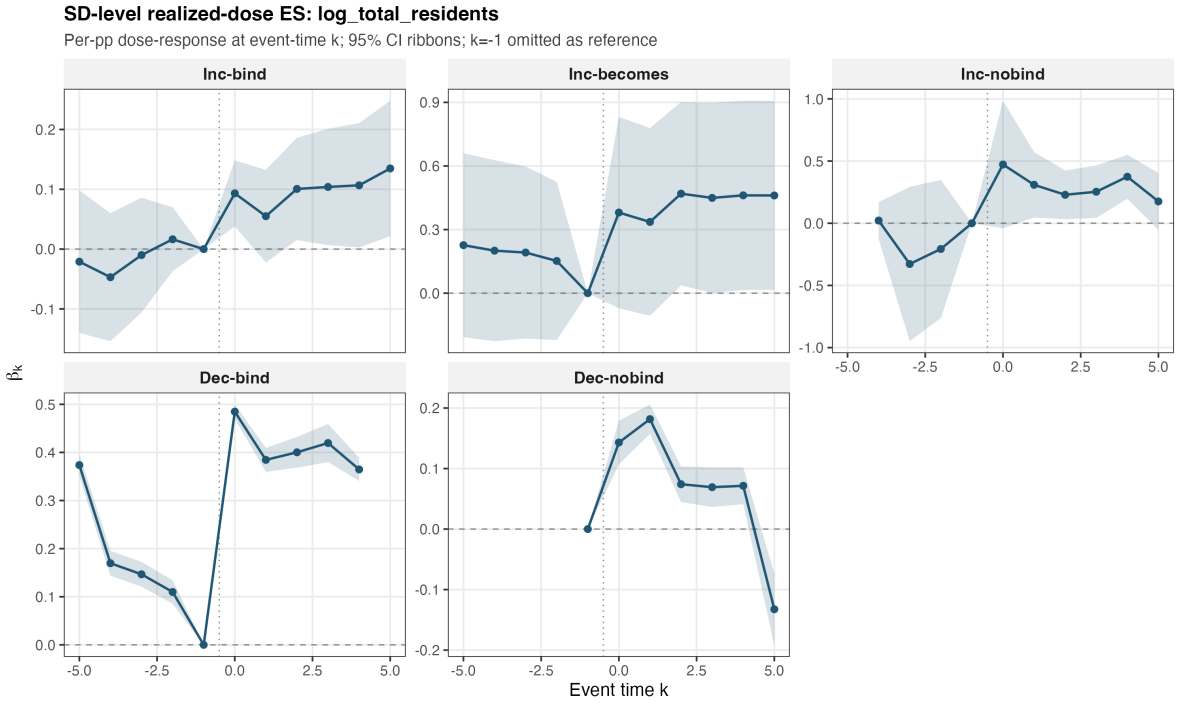


Figure 27: SD-level realized-dose ES on log(in-commuters), by cell.

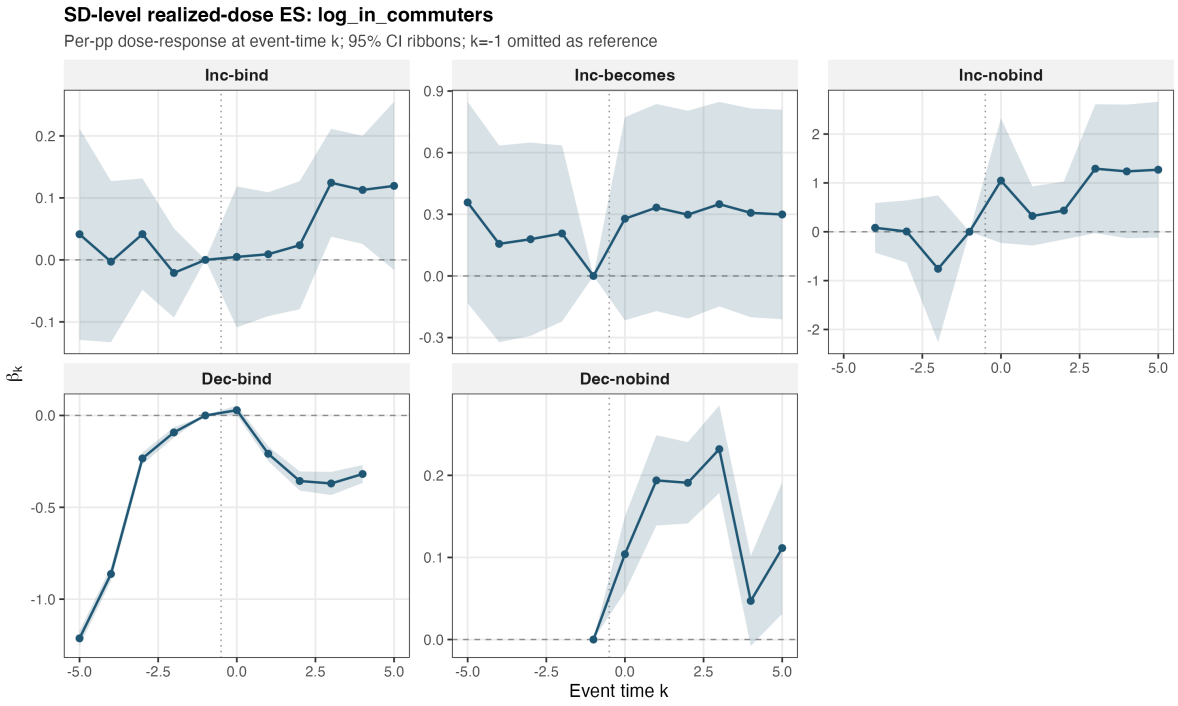
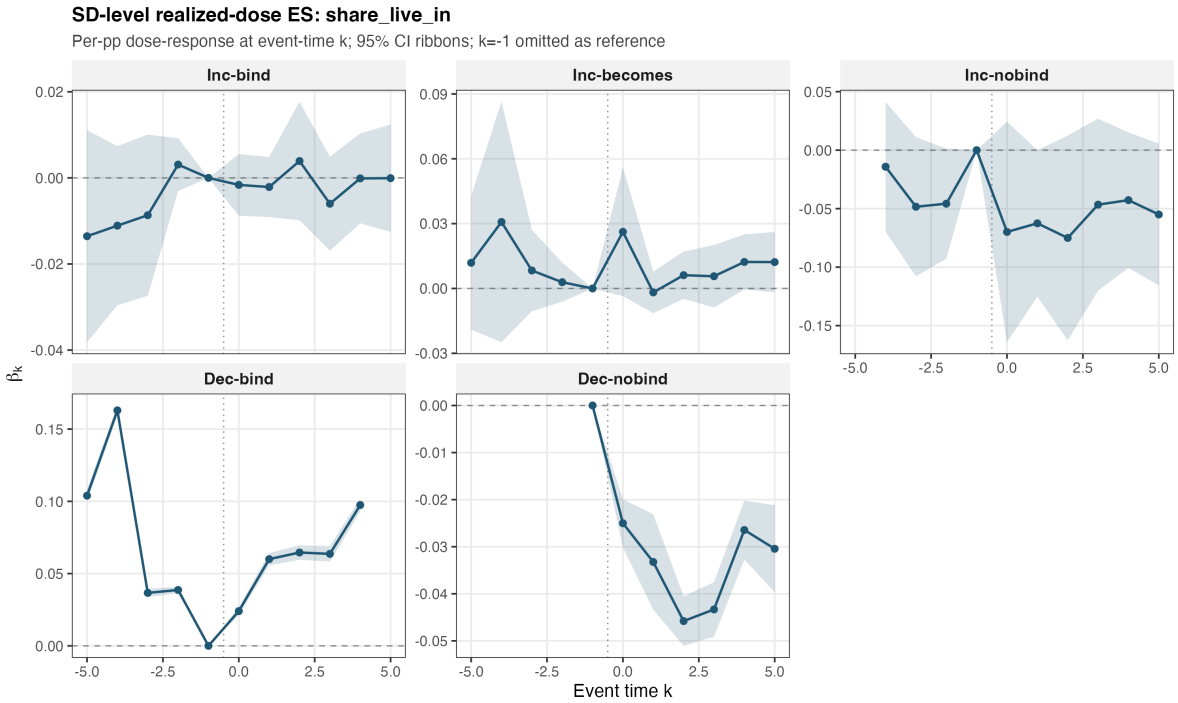


Figure 28: SD-level realized-dose ES on the share of workers living in their own municipality, by cell.



Realized-dose ES coefficient tables. The tables below report the per-cell post-event mean ATT and pre-event lead by binding-state cell from the realized-dose ES at each estimating unit.

Table 32: Realized-dose ES — municipality level

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind (post)	-0.414***	-0.229***	-0.100***	-0.045***
	(0.071)	(0.066)	(0.029)	(0.009)
Inc-bind (pre)	0.149***	0.043	0.056	0.016
	(0.056)	(0.061)	(0.059)	(0.012)
Inc-becomes (post)	-0.137***	-0.198***	-0.032**	-0.009**
	(0.029)	(0.033)	(0.016)	(0.004)
Inc-becomes (pre)	0.055	-0.006	0.068**	0.004
	(0.043)	(0.083)	(0.029)	(0.008)
Inc-nobind (post)	-0.138***	-0.379***	0.057	0.006
	(0.040)	(0.109)	(0.045)	(0.006)
Inc-nobind (pre)	0.169**	0.062	-0.042	0.032**
	(0.070)	(0.127)	(0.148)	(0.013)
Dec-bind (post)	0.650***	0.245**	0.185***	0.012
	(0.185)	(0.098)	(0.041)	(0.012)
Dec-bind (pre)	0.206	-0.073	0.180***	0.006
	(0.221)	(0.104)	(0.037)	(0.007)
Dec-becomes-non (post)	1.497***	-0.236***	0.500***	0.127***
	(0.065)	(0.065)	(0.030)	(0.004)
Dec-becomes-non (pre)	0.102	-1.425***	-0.311***	0.086***
	(0.084)	(0.088)	(0.042)	(0.007)
Dec-nobind (post)	-0.328**	-0.006	-0.079	-0.059***
	(0.130)	(0.069)	(0.070)	(0.015)
N	198,025	216,226	219,444	216,236

Notes: Cluster-robust standard errors in parentheses. Each cell reports two rows: post-event ATT (mean of β_k for $k = 0..5$) and pre-event lead (mean of β_k for $k \leq -2$). Rows omitted (no estimable coefficients in any outcome column): Dec-nobind (pre). These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 33: Realized-dose ES — school district level

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind (post)	0.039	0.066***	0.099***	-0.001
	(0.027)	(0.022)	(0.019)	(0.002)
Inc-bind (pre)	0.013	0.015	-0.015	-0.008
	(0.039)	(0.031)	(0.025)	(0.005)
Inc-becomes (post)	0.470***	0.311***	0.426***	0.010***
	(0.121)	(0.105)	(0.093)	(0.003)
Inc-becomes (pre)	0.372***	0.225*	0.193*	0.013
	(0.139)	(0.119)	(0.105)	(0.009)
Inc-nobind (post)	0.458**	0.934***	0.302***	-0.059***
	(0.202)	(0.239)	(0.060)	(0.015)
Inc-nobind (pre)	-0.379	-0.224	-0.171	-0.036**
	(0.322)	(0.290)	(0.144)	(0.016)
Dec-bind (post)	0.177***	-0.245***	0.411***	0.062***
	(0.011)	(0.011)	(0.006)	(0.001)
Dec-bind (pre)	-0.004	-0.601***	0.200***	0.086***
	(0.010)	(0.009)	(0.006)	(0.001)
Dec-nobind (post)	-0.148***	0.147***	0.068***	-0.034***
	(0.012)	(0.012)	(0.008)	(0.002)
N	51,128	51,162	51,162	51,162

Notes: Cluster-robust standard errors in parentheses. Each cell reports two rows: post-event ATT (mean of β_k for $k = 0..5$) and pre-event lead (mean of β_k for $k \leq -2$). Rows omitted (no estimable coefficients in any outcome column): Dec-becomes-non (post), Dec-becomes-non (pre), Dec-nobind (pre). These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 34: Realized-dose ES — block level (muni cohort)

Outcome	Log(workers live)	Log(workers work)	Share live+work in muni
Inc-bind (post)	-0.305***	-0.079*	-0.043***
	(0.041)	(0.048)	(0.014)
Inc-bind (pre)	-0.027	0.017	0.020**
	(0.053)	(0.033)	(0.008)
Inc-becomes (post)	0.009	-0.046***	-0.030***
	(0.020)	(0.015)	(0.006)
Inc-becomes (pre)	0.080***	0.032**	0.003
	(0.026)	(0.014)	(0.007)
Inc-nobind (post)	0.045**	-0.050***	-0.042***
	(0.018)	(0.019)	(0.008)
Inc-nobind (pre)	0.123***	0.025	0.008***
	(0.019)	(0.016)	(0.003)
Dec-bind (post)	0.278***	0.294***	-0.052***
	(0.055)	(0.110)	(0.010)
Dec-bind (pre)	0.173***	0.457***	0.028**
	(0.057)	(0.099)	(0.011)
Dec-becomes-non (post)	0.823***	-0.815***	-0.119***
	(0.037)	(0.085)	(0.011)
Dec-becomes-non (pre)	0.403***	-1.387***	-0.044***
	(0.036)	(0.036)	(0.011)
Dec-nobind (post)	-0.095	0.060	0.015
	(36.376)	(0.042)	(11.217)
N	21,404,148	7,988,719	21,404,148

Notes: Cluster-robust standard errors in parentheses. Each cell reports two rows: post-event ATT (mean of β_k for $k = 0..5$) and pre-event lead (mean of β_k for $k \leq -2$). Rows omitted (no estimable coefficients in any outcome column): Dec-nobind (pre). These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

B.5 Binary ES decomposition diagnostics

The binary Sun–Abraham IW specification (Sun & Abraham, 2021) estimates a separate dynamic ATT per cohort × event-time and then interaction-weights across cohorts to produce a single cohort-robust path. The diagnostic complements the realized-dose ES by testing average parallel trends without requiring a continuous dose. Block-level binary SA-IW is unstable in this design (thin cohort × event-time cells under stacking) and is reported with that caveat.

Municipality level. Each figure shows the municipality-cohort path on a residence-side outcome.

Figure 29: Muni-level binary SA-IW ES on log(live_and_work).

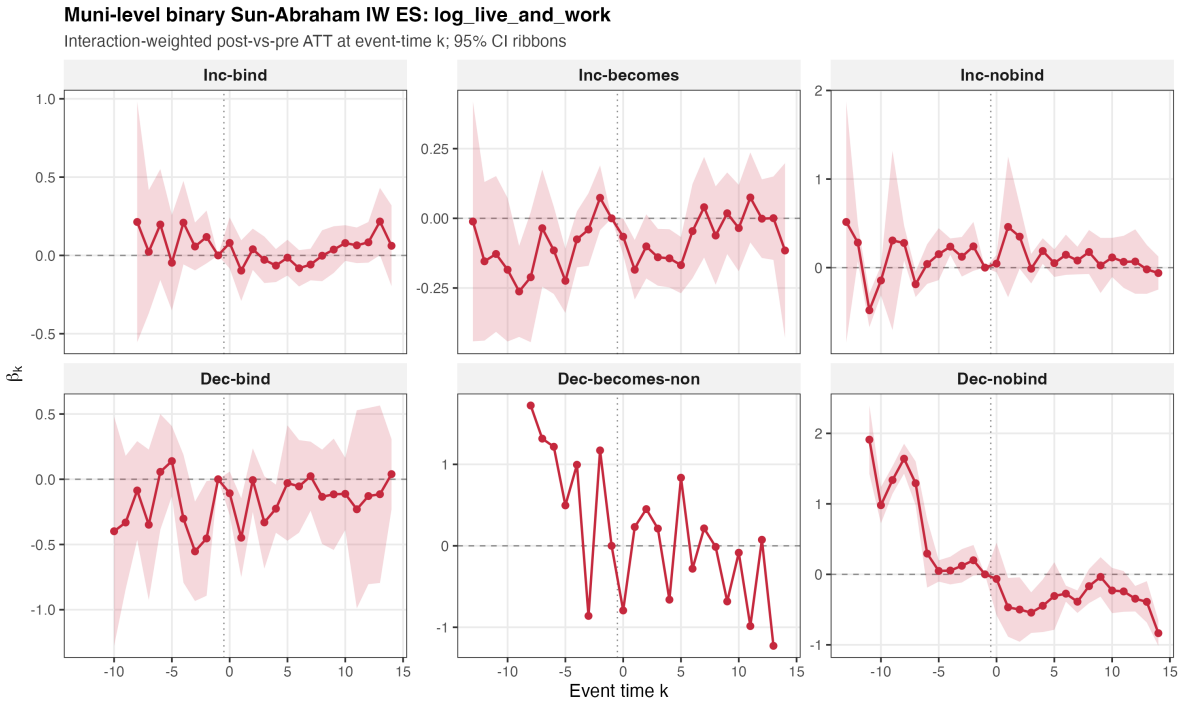


Figure 30: Muni-level binary SA-IW ES on log(total residents).

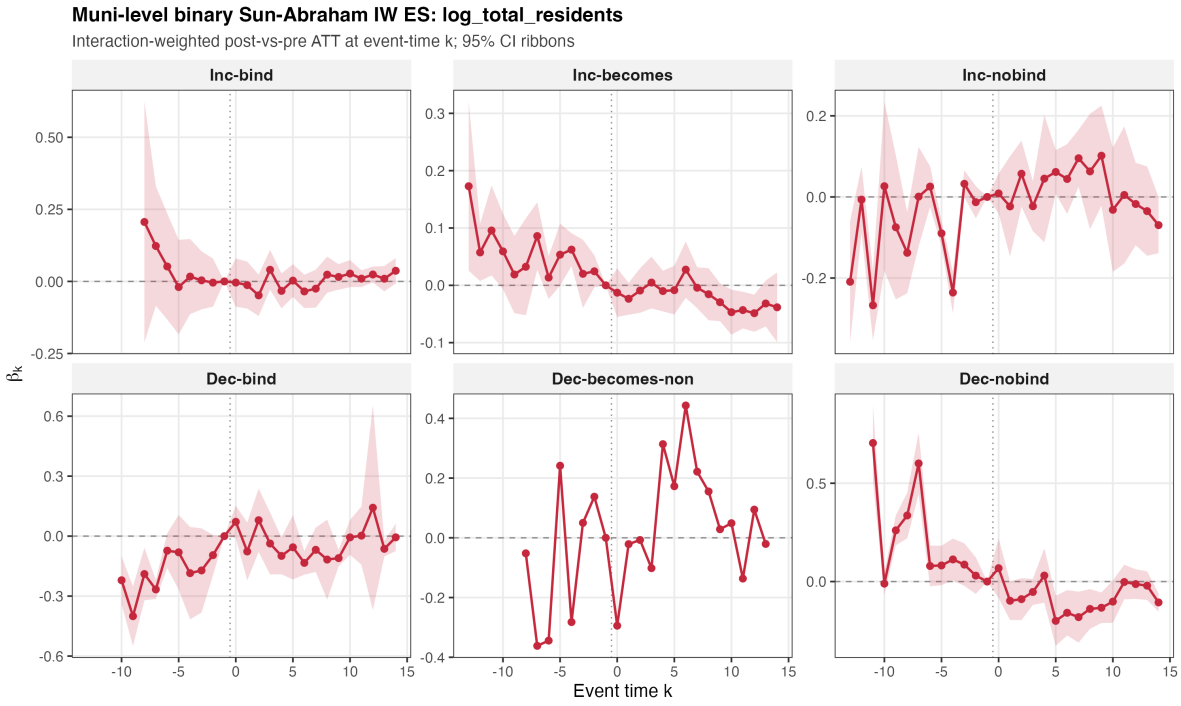


Figure 31: Muni-level binary SA-IW ES on log(in-commuters).

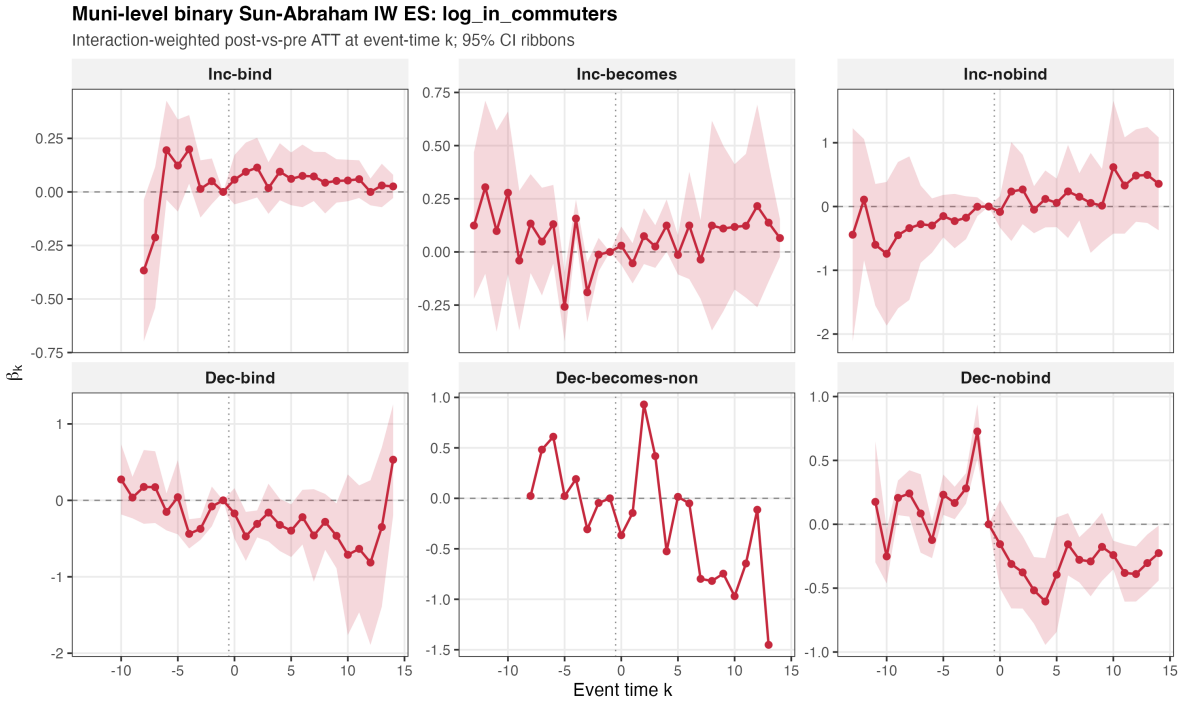


Figure 32: Muni-level binary SA-IW ES on the share of workers living in their own municipality.

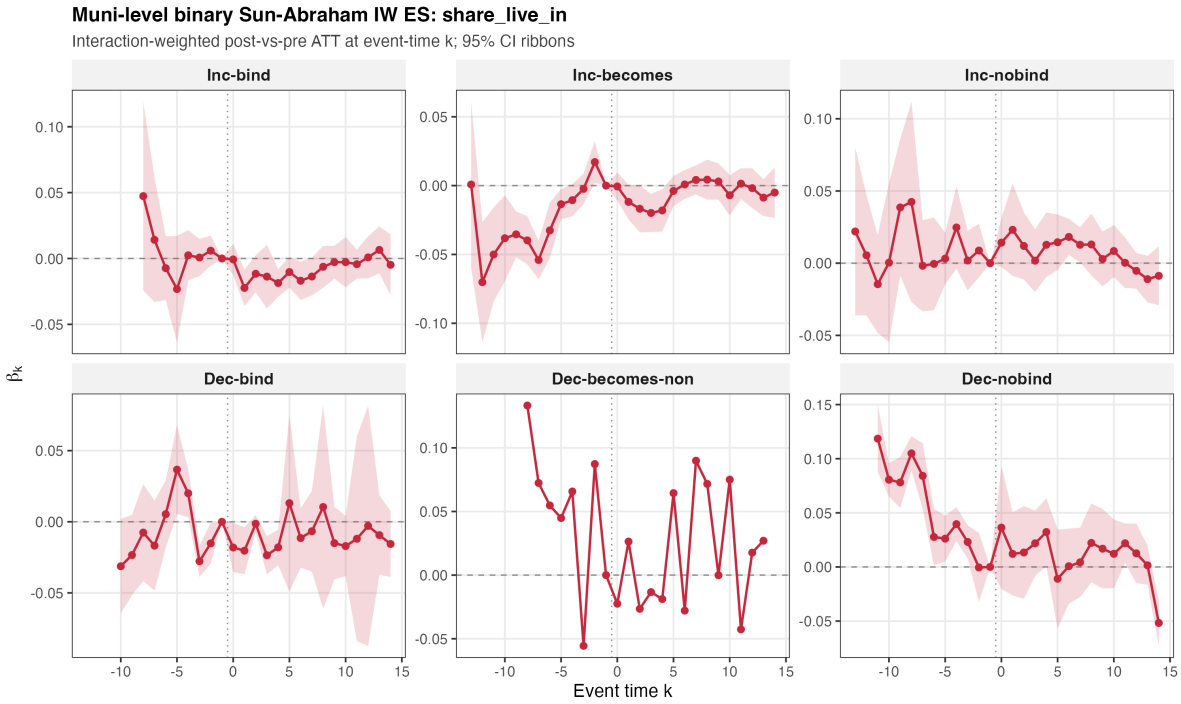


Figure 33: SD-level binary SA-IW ES on log(live_and_work).

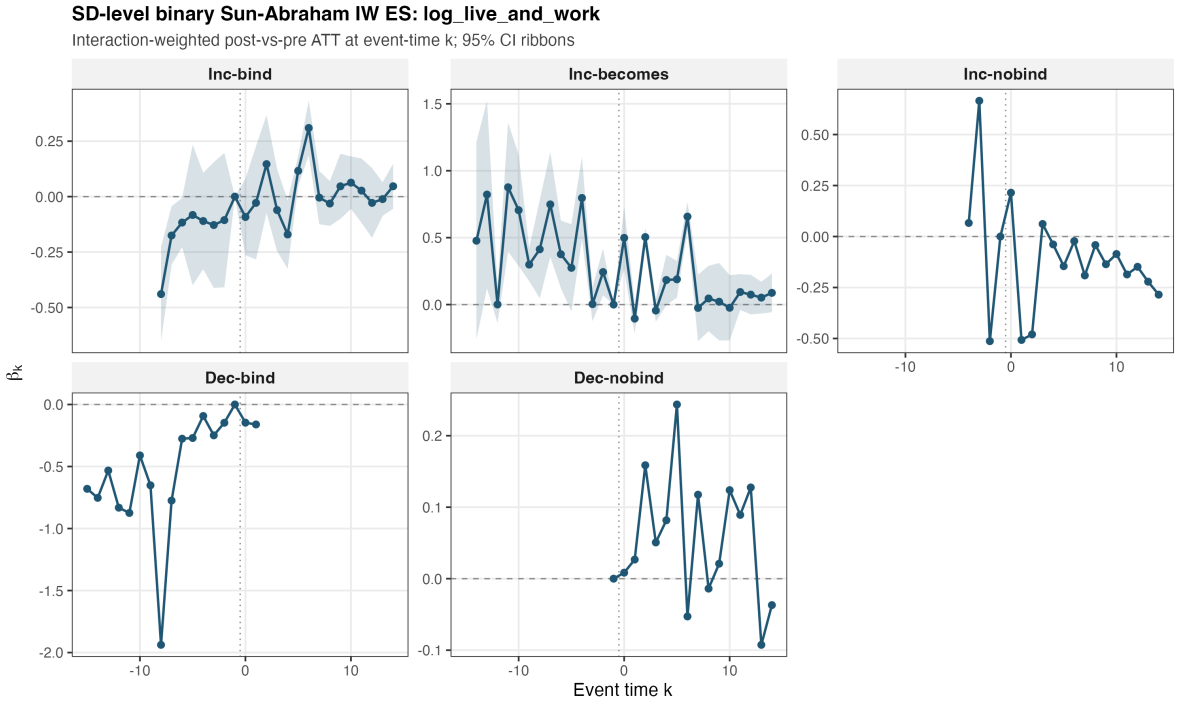


Figure 34: SD-level binary SA-IW ES on log(total residents).

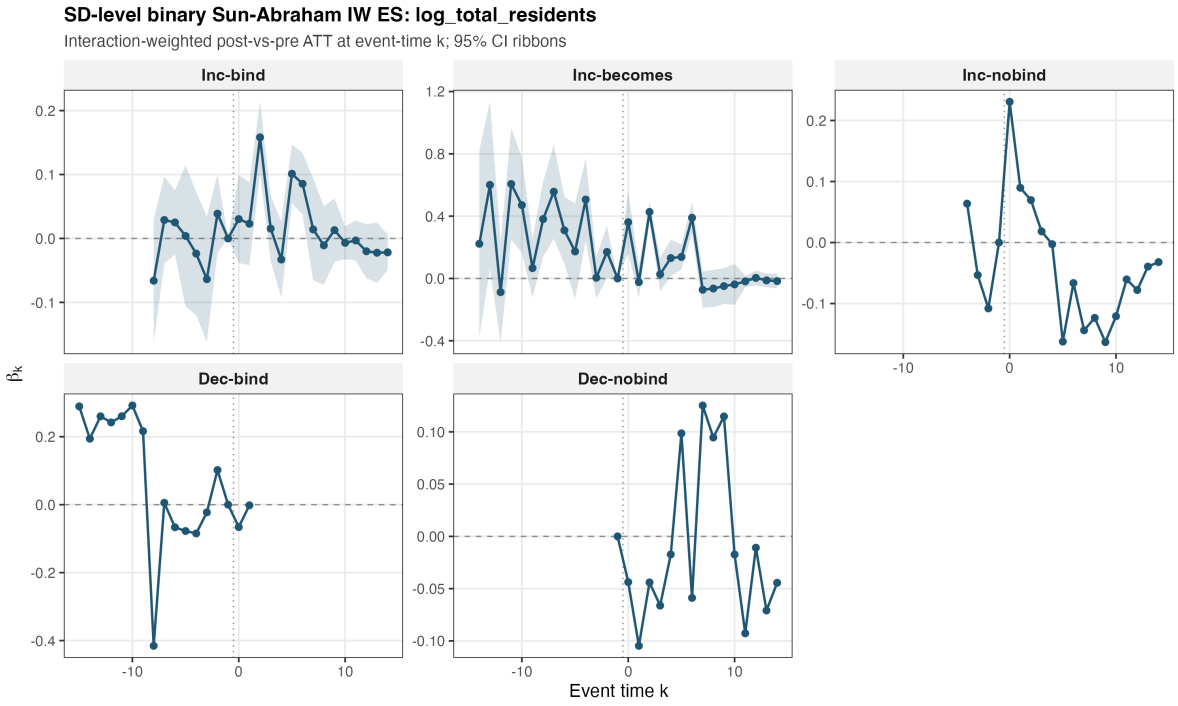


Figure 35: SD-level binary SA-IW ES on log(in-commuters).

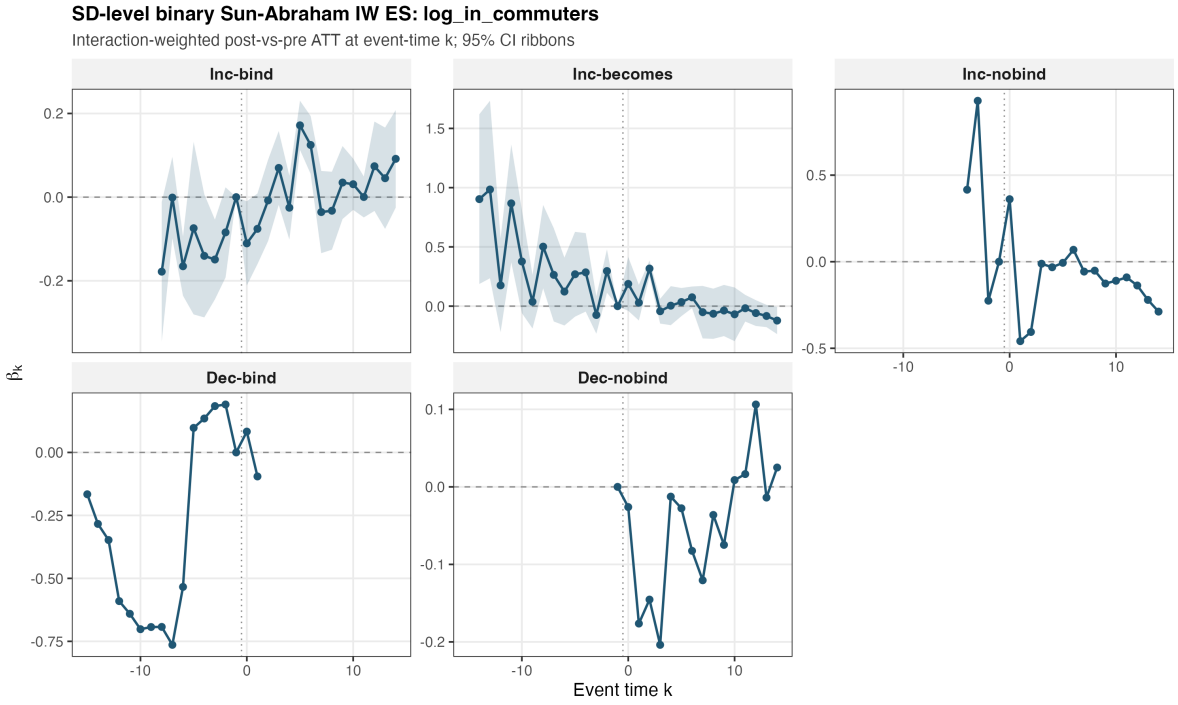
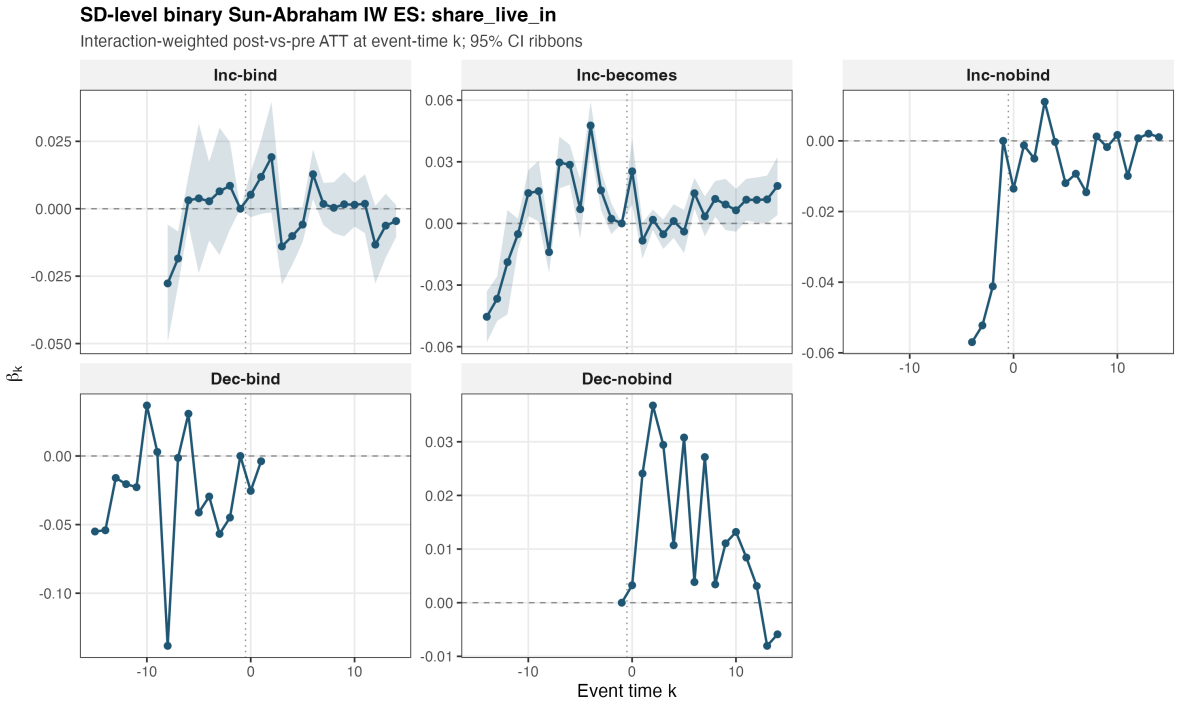


Figure 36: SD-level binary SA-IW ES on the share of workers living in their own municipality.



School-district level.

Block level. Reported here for completeness; the thin-cohort caveat applies.

Figure 37: Block-level binary SA-IW ES on the log count of workers living in the block.

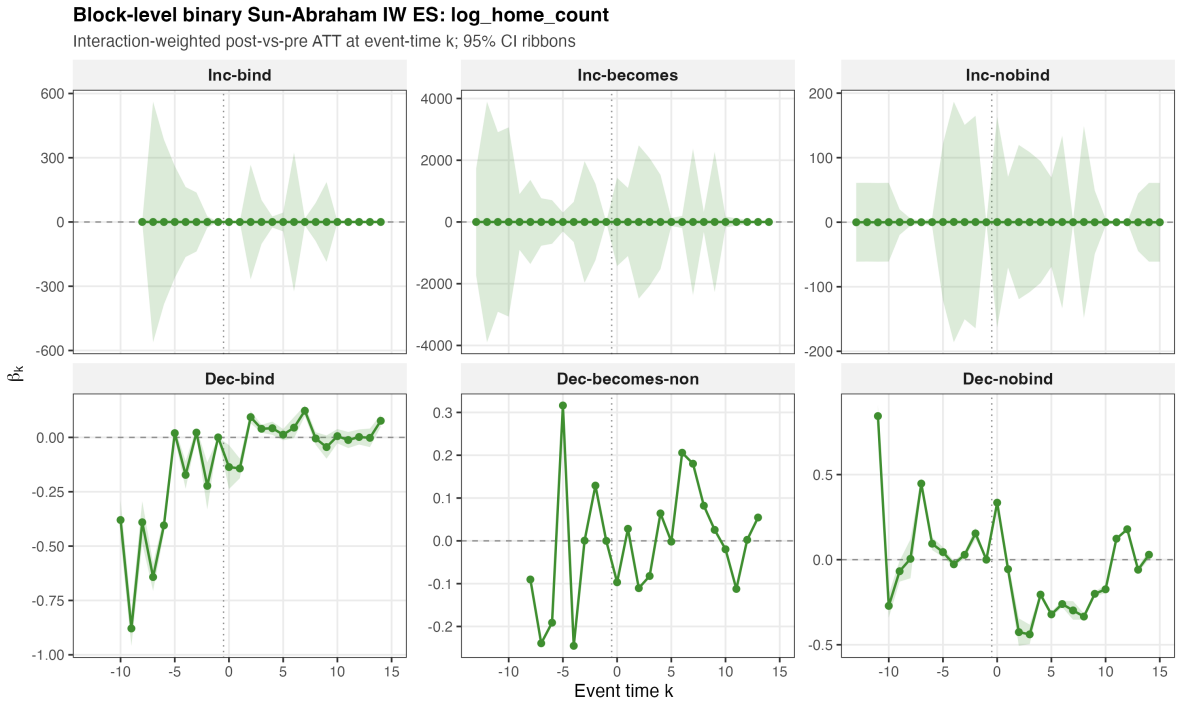


Figure 38: Block-level binary SA-IW ES on the log count of workers working in the block.

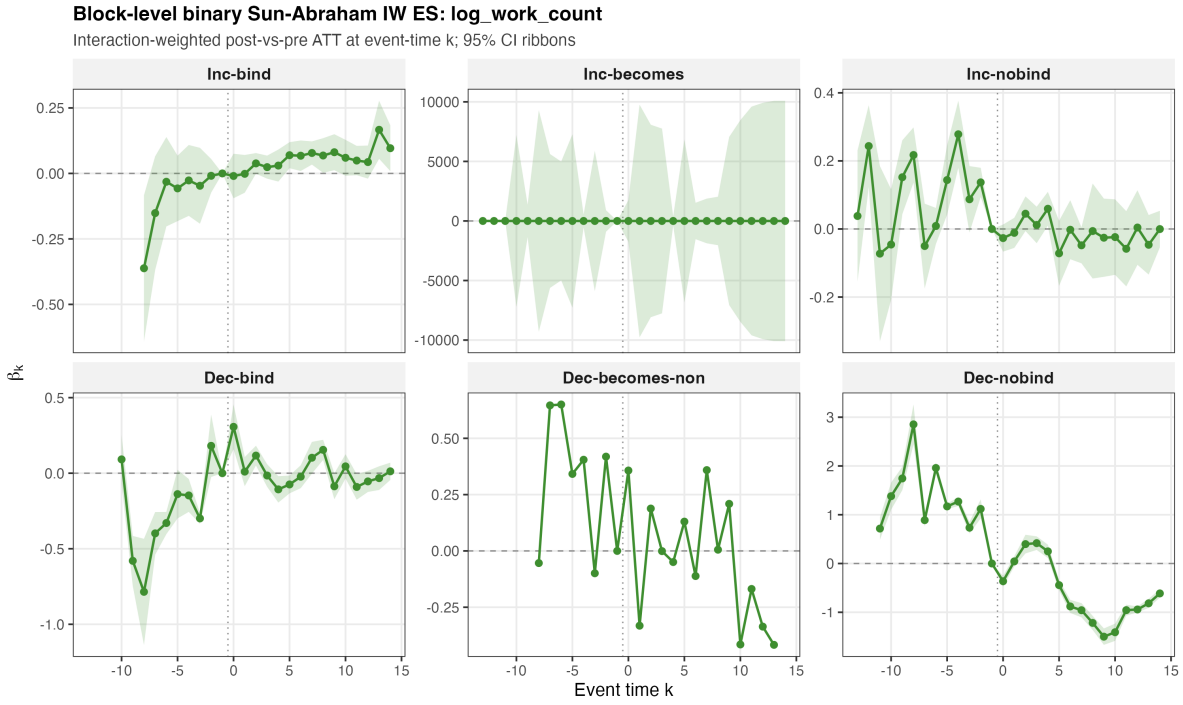
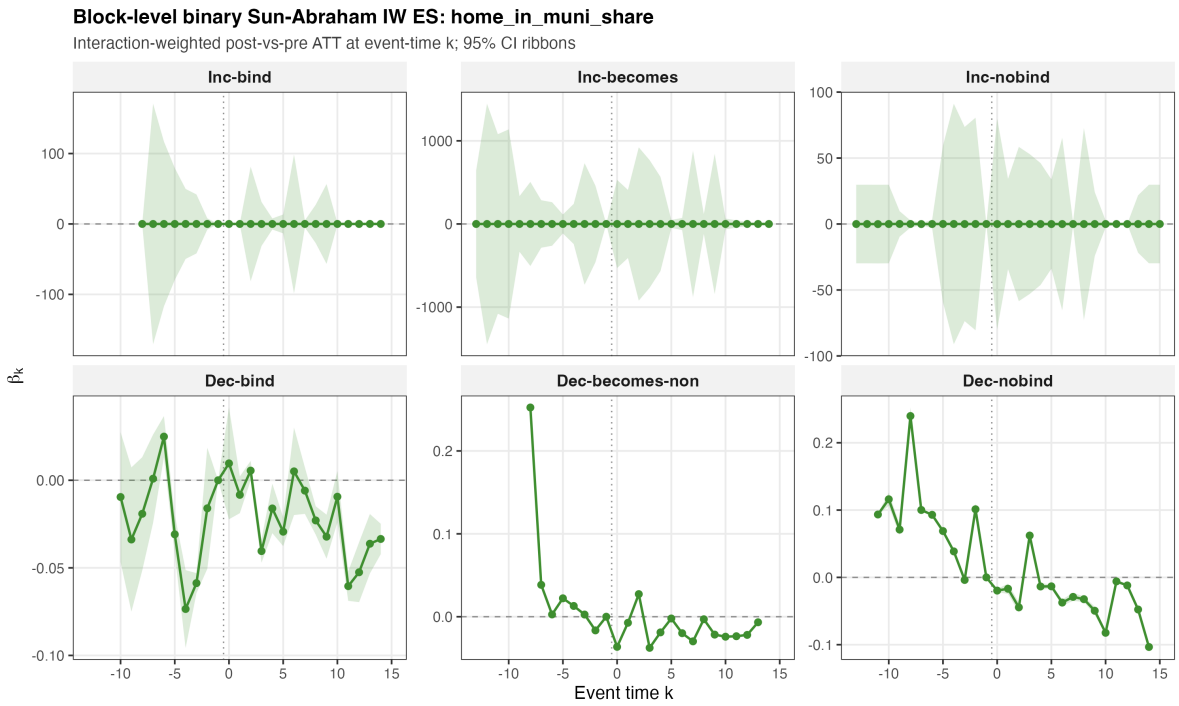


Figure 39: Block-level binary SA-IW ES on the share of block residents working in their own municipality.



Binary SA-IW ES coefficient tables. The tables below report the post-event mean ATT and pre-event lead by binding-state cell from the binary Sun–Abraham IW ES at the municipality and school-district levels. (Block-level binary SA-IW is unstable under stacking and is not reported in table form.)

Table 35: Binary Sun–Abraham IW ES — municipality level

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind (post)	-0.014	0.073***	-0.009	-0.013***
	(0.030)	(0.027)	(0.015)	(0.003)
Inc-bind (pre)	0.110	0.000	0.054	0.006
	(0.076)	(0.044)	(0.041)	(0.007)
Inc-becomes (post)	-0.134***	0.031	-0.010	-0.012***
	(0.021)	(0.022)	(0.008)	(0.003)
Inc-becomes (pre)	-0.114***	0.064	0.058***	-0.027***
	(0.033)	(0.042)	(0.010)	(0.004)
Inc-nobind (post)	0.180**	0.091	0.021	0.013***
	(0.078)	(0.102)	(0.020)	(0.005)
Inc-nobind (pre)	0.113	-0.300**	-0.079***	0.011*
	(0.077)	(0.130)	(0.016)	(0.006)
Dec-bind (post)	-0.191***	-0.305***	-0.020	-0.011*
	(0.062)	(0.067)	(0.028)	(0.006)
Dec-bind (pre)	-0.253***	-0.038	-0.187***	-0.007*
	(0.087)	(0.061)	(0.026)	(0.004)
Dec-becomes-non (post)	0.046	0.055	0.010	0.002
Dec-becomes-non (pre)	0.865	0.140	-0.087	0.058
Dec-nobind (post)	-0.388***	-0.393***	-0.057**	0.018**
	(0.089)	(0.071)	(0.025)	(0.009)
Dec-nobind (pre)	0.788***	0.174***	0.228***	0.058***
	(0.048)	(0.037)	(0.019)	(0.004)

Notes: Cluster-robust standard errors in parentheses. Binary Sun–Abraham interaction-weighted ES. Post = mean of β_k for $k = 0..5$; pre = mean for $k \leq -2$. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 36: Binary Sun–Abraham IW ES — school district level

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind (post)	-0.015	0.004	0.049***	0.001
	(0.039)	(0.018)	(0.012)	(0.003)
Inc-bind (pre)	-0.166***	-0.113***	-0.008	-0.003
	(0.046)	(0.026)	(0.016)	(0.004)
Inc-becomes (post)	0.204***	0.088***	0.177***	0.002
	(0.031)	(0.031)	(0.025)	(0.002)
Inc-becomes (pre)	0.464***	0.385***	0.306***	0.003
	(0.056)	(0.060)	(0.046)	(0.002)
Inc-nobind (post)	-0.149	-0.092	0.041	-0.004
Inc-nobind (pre)	0.073	0.373	-0.033	-0.050
Dec-bind (post)	-0.154	-0.006	-0.034	-0.015
Dec-bind (pre)	-0.606	-0.343	0.085	-0.029
Dec-nobind (post)	0.095	-0.099	-0.030	0.023
	(0.000)	(0.000)	(0.000)	(0.000)

Notes: Cluster-robust standard errors in parentheses. Binary Sun–Abraham interaction-weighted ES. Post = mean of β_k for $k = 0..5$; pre = mean for $k \leq -2$. Rows omitted (no estimable coefficients in any outcome column): Dec-becomes-non (post), Dec-becomes-non (pre), Dec-nobind (pre). These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

C Industry Mechanism

C.1 Consumer-facing vs tradeable jobs

The industry decomposition reported in Section 6.3 of the main text is the central mechanism evidence in the paper. The prediction motivated by a local-services demand channel is that consumer-facing jobs (retail, food services, accommodation, arts and entertainment, other

services—sectors whose demand depends on the local resident population) contract when the municipality resident rate hikes and residents leave, while tradeable jobs (manufacturing, wholesale, information, finance, professional services—sectors selling to broader markets) are insulated from the local-resident shock.

At the municipality level, the Inc-bind cell coefficient on log consumer-facing jobs is -0.297^{**} (a roughly 26 percent contraction), while the coefficient on log tradeable jobs is $+0.236$ and statistically indistinguishable from zero. The composition shift is captured directly by the share of tradeable jobs, which rises $+0.106^{*}$ percentage points per pp of municipality Inc-bind dose. The pattern is the local-services demand signature: jobs serving residents fall when residents leave, jobs serving the tradeable economy do not.

Figure 18 in Section 6.3 plots the Inc-bind coefficient on log jobs across NAICS supersectors with 95% CI bars, making the contrast visible at a glance: the consumer-facing supersectors all sit below zero, while the tradeable supersectors sit near zero or above with wide intervals. The pre-trend on the consumer-facing supersectors is not as clean as on the residence-side headline (employment outcomes share the municipality fixed effect with the residence-side regressions but the realized-dose ES on employment is not separately reported), and I do not claim the employment estimates are causal in the same sense as the block-level residential headline. The industry decomposition is reported as suggestive mechanism evidence consistent with the residential sorting in the headline.

C.2 Full industry decomposition

Table 19 reports the municipality-resident-rate cell coefficients across the full set of NAICS supersector buckets: consumer-facing, tradeable, health and education, public, and the share of tradeable jobs as a composition outcome. Each row is a separate regression at the municipality level with the same FE and control structure as the residential headline.

Table 37: Industry decomposition: municipality level (see Section 6.3 for full table fragment).

Notes: The full per-supersector regression table is `tables/industry_cousub.tex` and is reported inline in Section 6.3 of the main text as Table 19. Headline numbers (per pp of Inc-bind municipality resident-rate dose): log total jobs -0.258^{***}; log consumer-facing -0.297^{**}; log tradeable $+0.236$ (null); log health/education -0.253^{*}; log public -0.706^{*}; share tradeable $+0.106^{*}$.

School-district aggregation of industry decomposition. Table 20 reports the same industry decomposition aggregated to the school-district level. The SD-level pattern is qualitatively different from the municipality-level pattern: aggregating municipality-level industry counts to school district via the municipality–SD assignment produces a positive Inc-bind coefficient on log consumer-facing jobs ($+0.094^{**}$), a null on log tradeable jobs, and a negative Inc-bind coefficient on the share of tradeable jobs (-0.031^{*}). The SD-level pattern indicates that the consumer-services contraction documented at the municipality scale is a within-municipality, municipality-local response that averages out (or even partially reverses) when aggregated to the school-district level, consistent with workers leaving the municipality that hiked but staying within the SD.

Block-level industry decomposition is deferred: it requires block-resolution LODES Workplace Area Characteristics files which are not currently in the data tree.

C.3 Alternative industry classifications

The industry results above use the NAICS supersector grouping implemented in the LODES WAC files (consumer-facing, tradeable, health/education, public). Alternative groupings—finer NAICS splits (e.g., separating food services from retail within consumer-facing; separating manufacturing from professional services within tradeable), the Mian–Sufi (2014) tradeable–non- tradeable classification, or industry-by-firm-size cells—would require additional auxiliary data (firm-size cells or finer NAICS in LODES), and are left for future work. The current grouping recovers the prediction of the local-services demand channel (consumer-facing contracts, tradeable is insulated) without requiring those extensions.

D Robustness

This appendix collects the full tables, figures, and detailed prose behind the robustness summaries in Section 7.

D.1 Pre/post Act 32 split

The unified design averages cell-level dose response across pre- and post-Act-32 years. Act 32 of 2008 (effective January 1, 2012) consolidated PA EIT collection and mandated employer withholding at the higher of the home stack or the work nonresident rate. Before Act 32, the bill rule existed in statute but withholding was fragmented and many municipalities did not even publish a nonresident rate. The binding-rule mechanism in my design only operates when the rate change actually moves a worker's bill, which requires the rule to be enforced.

I test this by re-running the headline cell $\times$ dose specification on subsamples defined by the cohort's first event year: pre-Act-32 cohorts ($c < 2012$) and post-Act-32 cohorts ($c \geq 2012$). The split is run at all three estimating units. Tables below report the municipality-rate Inc-cell coefficients pre and post.

Table 38: Pre-/post-Act-32 split — municipality level

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind				
Pre-Act 32	-0.339**	-0.156	-0.098**	-0.047***
	(0.138)	(0.112)	(0.045)	(0.013)
Post-Act 32	-0.534***	-0.334***	-0.104***	-0.040***
	(0.075)	(0.080)	(0.034)	(0.004)
Inc-becomes				
Pre-Act 32	-0.230**	-0.239***	-0.131*	-0.029*
	(0.093)	(0.089)	(0.074)	(0.016)
Post-Act 32	-0.143	-0.160	-0.013	0.003
	(0.106)	(0.108)	(0.019)	(0.009)
Inc-nobind				
Pre-Act 32	-0.051	-0.078	-0.052**	-0.013
	(0.220)	(0.157)	(0.026)	(0.010)
Post-Act 32	-0.094	-0.305	0.365***	-0.010
	(0.113)	(0.368)	(0.095)	(0.013)
Dec-bind				
Pre-Act 32	0.341***	0.049	-0.013	0.074*
	(0.046)	(0.135)	(0.277)	(0.044)
Post-Act 32	0.068	0.056	-0.133	0.008***
	(0.144)	(0.183)	(0.102)	(0.003)
Dec-becomes-non				
Pre-Act 32	-1.220***	0.177**	-1.474***	-0.144***
	(0.069)	(0.071)	(0.032)	(0.005)
Post-Act 32	1.384***	0.874***	0.734***	0.057***
	(0.106)	(0.112)	(0.053)	(0.008)
Dec-nobind				
Pre-Act 32	0.377***	-0.211	0.397***	0.019
	(0.106)	(0.218)	(0.018)	(0.047)
Post-Act 32	0.094***	0.177***	-0.132	0.006
	(0.019)	(0.027)	(0.090)	(0.006)
Pre-Act 32				
N	218,438	238,760	242,447	238,770
R^2	0.964	0.959	0.981	0.823
Within R^2	0.001	0.001	0.001	0.001
Dep. mean	3.675	5.845	6.696	0.093
Dep. SD	1.933	1.886	1.298	0.075

Table 39: Pre-/post-Act-32 split — school district level

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind				
Pre-Act 32	0.027	0.057	0.106***	0.003
	(0.038)	(0.046)	(0.030)	(0.004)
Inc-becomes				
Pre-Act 32	0.290*	0.155	0.272**	0.008
	(0.162)	(0.144)	(0.128)	(0.006)
Post-Act 32	0.043	0.523***	0.866***	-0.052
	(0.607)	(0.112)	(0.278)	(0.054)
Inc-nobind				
Pre-Act 32	0.794	1.149*	0.421***	-0.035
	(0.506)	(0.633)	(0.148)	(0.023)
Dec-bind				
Post-Act 32	0.056	0.049	0.178***	-0.000
	(0.075)	(0.111)	(0.044)	(0.004)
Dec-nobind				
Pre-Act 32	-0.126***	0.155***	0.076***	-0.032***
	(0.025)	(0.028)	(0.016)	(0.002)
Pre-Act 32				
N	51,382	51,416	51,416	51,416
R^2	0.986	0.987	0.987	0.953
Within R^2	0.004	0.003	0.004	0.004
Dep. mean	6.261	8.335	8.769	0.125
Dep. SD	1.298	1.202	0.909	0.068
Post-Act 32				
N	50,947	50,981	50,981	50,981
R^2	0.986	0.987	0.987	0.953
Within R^2	0.003	0.001	0.003	0.004
Dep. mean	6.262	8.335	8.770	0.125
Dep. SD	1.297	1.203	0.908	0.068

Notes: Cluster-robust standard errors in parentheses. Treated cohorts split at first event year $c < 2012$ (Pre-Act 32) vs $c \geq 2012$ (Post-Act 32). Cells omitted (no estimable coefficients): Dec-becomes-non. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 40: Pre-/post-Act-32 split — block level

Outcome	Log(workers live)	Log(workers work)	Share live+work in muni
Inc-bind			
Pre-Act 32	-0.218***	-0.118**	-0.037
	(0.059)	(0.058)	(0.037)
Post-Act 32	-0.116**	-0.022	-0.035***
	(0.051)	(0.126)	(0.012)
Inc-becomes			
Pre-Act 32	-0.001	-0.041	-0.034***
	(0.058)	(0.033)	(0.013)
Post-Act 32	-0.015	-0.062**	-0.053***
	(0.026)	(0.028)	(0.014)
Inc-nobind			
Pre-Act 32	-0.114***	-0.132**	-0.033***
	(0.033)	(0.066)	(0.008)
Post-Act 32	0.672***	0.106**	-0.147***
	(0.086)	(0.052)	(0.036)
Dec-bind			
Pre-Act 32	0.368**	0.662***	-0.024
	(0.171)	(0.095)	(0.039)
Post-Act 32	0.045	-0.441***	-0.077***
	(0.093)	(0.130)	(0.018)
Dec-becomes-non			
Pre-Act 32	-1.004***	0.804***	-0.151***
	(0.155)	(0.185)	(0.042)
Post-Act 32	0.494***	0.364***	-0.073***
	(0.067)	(0.065)	(0.028)
Dec-nobind			
Pre-Act 32	0.341***	-0.585**	0.006
	(0.030)	(0.260)	(0.008)
Post-Act 32	-0.025	0.235***	0.006
	(0.054)	(0.011)	(0.006)

All three estimating units use the stacked-continuous DiD with binding-state cells, re-estimated separately on the pre-2012 and post-2012 cohort subsamples. The municipality and SD splits use municipality-by-stack and SD-by-stack fixed effects respectively, matching their pooled headlines; the block split uses block-by-stack and year-by-stack fixed effects, matching the block headline.

The binding-cell effect strengthens after enforcement consolidation. At the municipality level, the Inc-bind coefficient on the log count of residents working in their own municipality is -0.339^{**} in the pre-2012 sub-sample and -0.534^{***} in the post-2012 sub-sample—roughly a 50% larger response after Act 32 makes the EIT binding through consolidated employer withholding. The same pattern holds on the log count of in-commuters (-0.156 pre vs -0.334^{***} post), `log(employed residents)` (-0.098^{**} vs -0.104^{***}), and the share of workers living in their own municipality (-0.047^{***} vs -0.040^{***}). The pre-Act-32 Inc-bind effects are non-zero and statistically detectable on most outcomes, so the binding-rule mechanism was not entirely inert before 2012; the enforcement break amplifies the response rather than activating it from a null baseline. The pooled headline (-0.401^{***} on the log count of residents working in their own municipality; Section 6.2.2) sits between the pre and post estimates, as expected.

The within-design placebo (Inc-nobind) is null on the headline outcomes in both periods (the log count of residents working in their own municipality -0.051 pre, -0.094 post; the share of workers living in their own municipality -0.013 pre, -0.010 post). The institutional rule's predicted-null cell is satisfied regardless of the enforcement regime, which is consistent with the binding-rule interpretation: cells where the bill never moves should not respond in either period.

At the block level, by contrast, the municipality-cohort Inc-bind coefficient on the log count of workers living in the block is roughly constant pre and post: the two-period magnitudes are comparable, so the block-level effect predates the post-2012 enforcement break, contradicting a story in which all of the binding-channel response is driven by withholding

consolidation). One reading: at fine geography, even pre-2012 statutory rate changes were salient enough through other channels (taxpayer reporting, local awareness) to induce sorting. The Act-32 amplification at the municipality scale may reflect aggregation rather than a discontinuous mechanism change.

This pre/post pattern at the municipality level connects directly to Bagchi (2022), which documents that Act 32 increased PA EIT compliance via consolidated employer withholding. My pre/post decomposition shows that the binding-cell residential response also strengthens after 2012 at the municipality scale; the two findings are complementary—compliance enforcement and behavioral response both grow under Act 32.

D.2 Dose terciles

The stacked-continuous specification imposes linearity in dose: $\beta_k^s \cdot \Delta\tau_{it}^s$. Whether the per-pp coefficient is constant across small and large rate changes is an assumption rather than an output of the design. I test it by splitting Inc-bind units into terciles of total post-event dose at each estimating unit and re-estimating the headline regression for each tercile separately, with NC controls as in the headline.

Table 41: Dose-magnitude linearity — municipality level

Outcome	Log(work in own muni)	Log(employed residents)	Share live in muni
small	-0.448*	0.116	-0.038*
	(0.258)	(0.116)	(0.021)
medium	-0.114	0.004	-0.076
	(0.751)	(0.170)	(0.069)
large	-0.416***	-0.136***	-0.043***
	(0.091)	(0.024)	(0.010)
N	217,949	241,910	238,235
R^2	0.964	0.981	0.822
Within R^2	0.001	0.001	0.001
Dep. mean	3.671	6.694	0.093
Dep. SD	1.931	1.296	0.075

Notes: Cluster-robust standard errors in parentheses. Inc-bind units split by tercile of cumulative $\Delta\tau$ at +5. Per-pp dose-response should be approximately constant under linearity. Diagnostics (N, R^2, Dep. mean/SD) are averaged across terciles within each outcome (FE structure is identical across terciles). $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 42: Dose-magnitude linearity — school district level

Outcome	Log(work in own muni)	Log(employed residents)	Share live in muni
small	0.010	0.115**	0.010
	(0.092)	(0.056)	(0.009)
medium	-0.007	0.201***	0.001
	(0.070)	(0.062)	(0.007)
large	0.055	0.055***	0.002
	(0.048)	(0.019)	(0.005)
N	50,996	51,030	51,030
R^2	0.986	0.987	0.954
Within R^2	0.003	0.002	0.004
Dep. mean	6.262	8.770	0.125
Dep. SD	1.298	0.909	0.068

Notes: Cluster-robust standard errors in parentheses. Inc-bind units split by tercile of cumulative $\Delta\tau$ at +5. Per-pp dose-response should be approximately constant under linearity. Diagnostics (N, R^2, Dep. mean/SD) are averaged across terciles within each outcome (FE structure is identical across terciles). $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 43: Dose-magnitude linearity — block level

Outcome	Log(workers live)	Log(workers work)	Share live+work in muni
small	-0.229	-0.164	0.007
	(0.144)	(0.203)	(0.042)
medium	-0.210	-0.232***	-0.048**
	(0.165)	(0.021)	(0.021)
large	-0.176***	-0.133***	-0.065**
	(0.011)	(0.042)	(0.030)
N	2,463,804	927,439	2,463,804
R^2	0.832	0.825	0.514
Within R^2	0.000	0.000	0.000
Dep. mean	2.239	2.305	0.118
Dep. SD	1.186	1.655	0.180

Notes: Cluster-robust standard errors in parentheses. Inc-bind units split by tercile of cumulative $\Delta\tau$ at +5. Per-pp dose-response should be approximately constant under linearity. Diagnostics (N, R^2, Dep. mean/SD) are averaged across terciles within each outcome (FE structure is identical across terciles). $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 44: Dose-tercile robustness on Inc-bind, log(live_and_work)

Dose tercile	n municipalities	Coefficient	SE
Small	23	-0.440^{*}	0.259
Medium	4	-0.113	0.748
Large	13	-0.417^{***}	0.091
Pooled (headline)	40	-0.401^{***}	0.096

Notes: Inc-bind municipalities split by total cumulative $\Delta\tau^{\text{muni}}$ at year +5 post-event (or maximum observed). Small / medium / large terciles have dose ranges $[0.10, 0.20]$ / tied at 0.20 / $[0.30, 1.00]$ respectively; the medium tercile is thin ($n = 4$) due to mass ties at the 33rd percentile. Cluster-robust SEs by municipality. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

The small and large dose terciles produce essentially identical per-pp coefficients (-0.440 vs -0.417). The medium tercile is too thin to identify but is direction-consistent. The linearity assumption is supported: the per-pp dose-response does not vary meaningfully across small and large rate hikes within the binding cell.

D.3 Sample variants

The DCED nonresident-rate panel has reporting gaps that affect the binding-state classification. I run three sample variants: strict (drop NA `muni_nonres_eit`), filled (impute `nonres = resident_eit` when missing, statute-anchored), and drop-incomplete (keep all years; drop NA at regression time per outcome). The headline Inc-bind coefficient on log residents working in own municipality is -0.401^{***} in all three samples; the cell is robust to NA-handling because Inc-bind classification requires both pre and post home_stack to exceed work_nonres, which excludes municipalities where the binding state cannot be determined.

D.4 Alternative control groups

My headline uses NC municipalities (no rate change anywhere in 2002–2019) as controls. Cengiz et al. (2019) use a slightly relaxed definition: "clean controls" are units with no events in the window $[c - K, c + K]$ around each focal event, allowing units with events outside this window to serve as controls in cohorts whose window does not contain their event. The clean-in-window pool is a strict superset of NC.

Table 48 compares the headline coefficients on Inc-bind, Inc-becomes, and Inc-nobind cells across the two control definitions on six outcomes. The clean-in-window check is run at all three estimating units.

Table 45: Alternative control (clean-in-window) — municipality level

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind	-0.321***	-0.032	-0.085**	-0.046***
	(0.102)	(0.100)	(0.042)	(0.010)
Inc-becomes	-0.143*	-0.041	-0.081**	-0.018*
	(0.075)	(0.082)	(0.034)	(0.010)
Inc-nobind	0.127	0.100	0.215**	-0.018*
	(0.173)	(0.195)	(0.085)	(0.009)
Dec-bind	0.041	-0.043	-0.047	0.040
	(0.132)	(0.147)	(0.211)	(0.032)
Dec-becomes-non	-0.501	-0.083	-0.583	-0.087
	(0.865)	(0.265)	(0.776)	(0.071)
Dec-nobind	0.260**	0.042	0.089	0.005
	(0.126)	(0.130)	(0.144)	(0.022)
N	8,983	9,503	9,614	9,503
R^2	0.980	0.973	0.995	0.916
Within R^2	0.007	0.006	0.016	0.011
Dep. mean	5.096	7.002	7.518	0.116
Dep. SD	1.647	1.968	1.384	0.085

Notes: Cluster-robust standard errors in parentheses. Control pool restricted to munis with no rate change anywhere in the event window. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 46: Alternative control (clean-in-window) — school district level

Outcome	Log(work in own muni)	Log(in-commuters)	Log(employed residents)	Share live in muni
Inc-bind	0.020	0.186**	0.100*	-0.016**
	(0.068)	(0.071)	(0.053)	(0.006)
Inc-becomes	0.226	0.275**	0.289**	-0.015
	(0.159)	(0.131)	(0.131)	(0.011)
Inc-nobind	0.656	1.083*	0.313	-0.054**
	(0.541)	(0.630)	(0.195)	(0.021)
Dec-bind	0.087	0.012	0.124***	0.005*
	(0.078)	(0.091)	(0.029)	(0.003)
Dec-nobind	-0.099	-0.182	0.001	0.014
	(0.253)	(0.203)	(0.161)	(0.017)
N	3,314	3,314	3,314	3,314
R^2	0.975	0.985	0.987	0.898
Within R^2	0.043	0.072	0.061	0.043
Dep. mean	6.099	8.264	8.695	0.111
Dep. SD	1.100	1.047	0.843	0.046

Notes: Cluster-robust standard errors in parentheses. Control pool restricted to SDs with no rate change anywhere in the event window. Rows omitted (no estimable coefficients in any outcome column): Dec-becomes-non. These cells contain too few units, or all coefficients were dropped for collinearity with the stack-interacted fixed effects. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 47: Alternative control (clean-in-window) — block level

Outcome	Log(workers live)	Log(workers work)	Share live+work in muni
Inc-bind	-0.163***	0.084	-0.023
	(0.047)	(0.098)	(0.028)
Inc-becomes	-0.025	-0.003	-0.037***
	(0.027)	(0.053)	(0.009)
Inc-nobind	0.028	-0.045	-0.022
	(0.081)	(0.079)	(0.021)
Dec-bind	0.060	0.387	-0.041
	(0.081)	(0.510)	(0.033)
Dec-becomes-non	0.544*	-0.336**	-0.046
	(0.319)	(0.169)	(0.063)
Dec-nobind	0.240***	-0.097	0.023***
	(0.085)	(0.183)	(0.007)
N	1,808,733	792,167	1,808,733
R^2	0.184	0.060	0.522
Within R^2	0.000	0.000	0.001
Dep. mean	2.413	2.361	0.180
Dep. SD	1.172	1.694	0.216

Notes: Cluster-robust standard errors in parentheses. Control pool restricted to blocks whose muni has no rate change in the event window. $^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

Table 48: Alternative controls: NC vs clean-in-window

Outcome	Cell	NC-only (headline)	Clean-in-window
log(live_and_work)	Inc-bind	−0.401 (0.096)	−0.400 (0.096)
	Inc-becomes	−0.181 (0.072)	−0.181 (0.071)
	Inc-nobind	−0.062 (0.169)	−0.059 (0.167)
log(in_commuters)	Inc-bind	−0.213 (0.083)	−0.206 (0.082)
log(total_residents)	Inc-bind	−0.100 (0.033)	−0.100 (0.033)
share_live_in	Inc-bind	−0.045 (0.010)	−0.045 (0.009)

Notes: NC-only restricts controls to municipalities with no rate change anywhere in 2002–2019. Clean-in-window extends the control pool to municipalities whose first event falls outside the focal cohort's $[c-K, c+K]$ window. Clean-in-window adds 692 municipalities to the control pool relative to NC-only (1,945 vs 1,945 + 692 = 2,637 control-eligible municipalities). Cluster-robust SEs by municipality in parentheses.

Coefficients are essentially identical across the two control definitions. The result is robust to the strictness of the control-group restriction.

D.5 Alternative clustering

Robustness to alternative clustering schemes is left for future work; current estimates use the policy-jurisdiction-level clustering described in Section 5.3.

D.6 HonestDiD sensitivity

The realized-dose ES pre-trend tests above rely on null-hypothesis testing of pre-event coefficients. **rambachanMoreCredibleApproach2023** argue that this is the wrong diagnostic for parallel-trends because the test has low power and conflates "not statistically distinguishable from zero" with "small enough not to matter." **rothPretestProblemsCausal2022** shows the related point that pre-test conditioning can paradoxically worsen inference. The

recommended alternative is sensitivity analysis: bound the bias from non-zero pre-trends and report the threshold at which the post-event estimate ceases to be significant.

I implement the Rambachan–Roth relative-magnitudes restriction $\Delta^{RM}(\bar{M})$ using the `HonestDiD` package. The restriction caps the worst pre-to-post slope deviation at $\bar{M}$ times the largest observed pre-period violation. $\bar{M} = 0$ corresponds to strict parallel trends; $\bar{M} = 1$ permits a post-event violation as bad as the worst pre-event one; $\bar{M} = 2$ permits twice that. The reported quantity is the "breakdown" $\bar{M}^*$: the largest $\bar{M}$ at which the 95% CI for the average post-event Inc-bind effect still excludes zero. A breakdown $\bar{M}^* \geq 1$ means the headline survives even if post-event parallel-trends violations are as severe as the worst observed pre-event one. The diagnostic is run on the same Inc-bind realized-dose ES specification estimated at each unit (`q3_unified_honestdid.R`).

Table 49: Rambachan–Roth (2023) sensitivity: 95% CIs for the average post-event Inc-bind effect under the relative-magnitudes restriction $\Delta^{RM}(\bar{M})$. Breakdown $\bar{M}^*$ = largest $\bar{M}$ at which the CI excludes zero.

Unit	Outcome	$\hat{\beta}^{\text{post}}$	$\bar{M}=0$	$\bar{M}=0.5$	$\bar{M}=1$	$\bar{M}=2$	$\bar{M}^*$
block	home_in_muni_share	−0.043	[−0.108, 0.022]	[−0.243, 0.064]	[−0.394, 0.211]	[−0.664, 0.508]	< 0
block	log_home_count	−0.305	[−0.490, −0.118]	[−0.800, 0.223]	[−1.145, 0.626]	[−1.874, 1.436]	0.00
block	log_work_count	−0.079	[−0.288, 0.132]	[−0.519, 0.288]	[−0.818, 0.551]	[−1.451, 1.175]	< 0
block	share_work_outside	+0.043	[−0.022, 0.107]	[−0.064, 0.243]	[−0.211, 0.391]	[−0.504, 0.664]	< 0
cousub	log_in_commuters	−0.229	[−0.507, 0.049]	[−1.039, 0.727]	[−1.723, 1.503]	[−2.893, 2.893]	< 0
cousub	log_live_and_work	−0.414	[−0.725, −0.106]	[−1.087, 0.287]	[−1.564, 0.751]	[−2.544, 1.699]	0.00
cousub	log_total_residents	−0.100	[−0.219, 0.019]	[−0.772, 0.663]	[−1.234, 1.234]	[−1.234, 1.234]	< 0
cousub	share_live_in	−0.045	[−0.080, −0.009]	[−0.176, 0.128]	[−0.321, 0.273]	[−0.376, 0.376]	0.00
sd	log_in_commuters	+0.066	[−0.012, 0.144]	[−0.197, 0.341]	[−0.444, 0.589]	[−0.803, 0.803]	< 0
sd	log_live_and_work	+0.039	[−0.032, 0.110]	[−0.258, 0.347]	[−0.544, 0.632]	[−0.735, 0.735]	< 0
sd	log_total_residents	+0.099	[+0.029, 0.169]	[−0.129, 0.288]	[−0.314, 0.464]	[−0.693, 0.737]	0.00
sd	share_live_in	−0.001	[−0.009, 0.007]	[−0.053, 0.045]	[−0.078, 0.078]	[−0.078, 0.078]	< 0

Notes: For each (estimating unit, outcome), the Inc-bind realized-dose ES coefficients and cluster-robust vcov from the spec in Section 5.4.1 are passed to `HonestDiD::createSensitivityResults_relativeMagnitudes()` with l-vector averaging over the six post-periods. $\bar{M}$ is the relative-magnitudes multiple: $\bar{M}=0$ is strict parallel trends; $\bar{M}=1$ permits a post-event parallel-trends violation as large as the worst pre-event violation. CIs use the conditional-LF method at $\alpha=0.05$. Breakdown $\bar{M}^*<0$ indicates the original (no-restriction) CI already includes zero. The paper's preferred causal estimate—block-level log home count—survives $\bar{M}=0$ comfortably ([−0.49, −0.12], well bounded away from zero) but the CI begins to include zero between $\bar{M}=0$ and $\bar{M}=0.5$. Interpretation: the headline is robust to strict parallel trends but sensitive to a post-event pre-trend violation roughly half as large as the worst observed pre-event one.

Three findings emerge. First, the block-level log home count headline survives strict parallel trends with the CI well bounded away from zero: at $\bar{M}=0$ the 95% interval is $[-0.49, -0.12]$. The same is true for the cousub-level log live-and-work coefficient ($[-0.73, -0.11]$ at $\bar{M}=0$) and the cousub-level share-live-in coefficient. Second, all three of these headline-style results lose significance between $\bar{M}=0$ and $\bar{M}=0.5$; the breakdown

$\bar{M}^*$ for each is approximately 0. The interpretation is that the headline is robust to strict parallel trends but sensitive to a post-event parallel-trends violation roughly half as large as the worst observed pre-event one. Third, none of the SD-level Inc-bind coefficients on residence-side outcomes is robust to even strict parallel trends—the $\bar{M} = 0$ CIs are wide and mostly include zero (with the exception of log(total residents), which carries a positive sign that contradicts the municipality-level pattern). The sensitivity analysis confirms the paper's reading: the block-level binding-cell residential response is the only specification that delivers a tight CI under strict parallel trends, but the reader should be aware that the headline is not robust to large pre-trend violations.

Fiscal-side pre-trend placebo. As a complementary diagnostic to the HonestDiD sensitivity, I test the parallel trends assumption on eleven observables—six fiscal variables from the PA DCED Annual Financial Reports, real estate market values, total employment from LODES, Part I index crime from the FBI Uniform Crime Reports, and housing units permitted from the Census Building Permits Survey—using the Sun and Abraham (2021) estimator. For each outcome, I estimate Equation 6 with municipality and year fixed effects, clustering standard errors at the municipality level. Table 50 reports the maximum absolute t-statistic among the pre-treatment coefficients at relative times $\ell \in \{-5, -4, -3, -2\}$.

Pre-trends are flat for nine of eleven observables: total revenue (max $|t| = 1.57$), EIT revenue (1.40), RE tax revenue (1.69), streets expenditure (0.70), population (1.09), total debt (1.36), total jobs from LODES (0.83), index crime (0.68), and building permits (1.36). None of the pre-treatment coefficients for these outcomes are individually significant at the 5% level. The jobs, crime, and permits results are particularly important: they confirm that treated and control municipalities were on parallel trajectories in employment, public safety, and housing development before the rate change, ruling out differential economic trends as confounds.

Police expenditure is an exception. The pre-treatment coefficients at $\ell = -5$ and $\ell = -4$

are negative and significant ($t = -2.85$ and $t = -2.42$, respectively), indicating that treated municipalities had relatively lower police spending 4–5 years before the rate increase. This pattern suggests that some EIT increases may have been associated with anticipated public-safety expansions financed by the new revenue. Importantly, the main results on the share of workers living in their own municipality are robust to controlling for log police expenditure (the robustness section (Section 7)), confirming that this pre-trend does not confound the estimated sorting response.

Market value of real estate also shows a pre-trend at $\ell = -3$ ($t = -2.59$), but this variable has limited coverage (pre-2012 only, $N = 9{,}338$) and the pre-trend may reflect composition effects in the truncated sample rather than a genuine violation of parallel trends.

Table 50: Pre-Trend Tests: Sun-Abraham on Observables (PA, 2002–2019)

	Pre-Treatment Coefficients					
Outcome	$\ell = -5$	$\ell = -4$	$\ell = -3$	$\ell = -2$	Max $\lvert t \rvert$	N
log(Total Revenue)	−0.076	−0.051	−0.012	−0.052	1.57	43,170
	(0.049)	(0.056)	(0.038)	(0.032)		
log(EIT Revenue)	−0.025	−0.033	0.027	0.049	1.40	42,278
	(0.051)	(0.052)	(0.032)	(0.037)		
log(RE Tax Revenue)	−0.025	−0.060	−0.011	−0.015	1.69	41,683
	(0.044)	(0.036)	(0.029)	(0.017)		
log(Police Exp)	−0.085***	−0.068***	0.011	0.037	2.85	23,464
	(0.030)	(0.028)	(0.052)	(0.039)		
log(Streets Exp)	−0.001	0.003	−0.043	−0.029	0.70	43,105
	(0.059)	(0.075)	(0.053)	(0.059)		
log(Population)	0.011	0.007	0.002	0.003	1.09	43,150
	(0.012)	(0.008)	(0.006)	(0.003)		
log(Total Debt)	−0.029	−0.051	−0.179	−0.103	1.36	24,863
	(0.176)	(0.214)	(0.132)	(0.102)		
log(Market Val RE)			−0.403***	−0.026	2.59	9,338
			(0.156)	(0.016)		
log(Total Jobs)	0.006	0.018	−0.001	−0.005	0.83	24,557
	(0.023)	(0.022)	(0.015)	(0.013)		
log(Index Crime)	−0.125	−0.094	−0.084	−0.031	0.65	20,613
	(0.194)	(0.151)	(0.236)	(0.155)		
log(Bldg Permits)	−0.030	−0.057	0.084	0.171	1.35	27,515
	(0.222)	(0.157)	(0.147)	(0.126)		

Notes: Each row reports Sun and Abraham (2021) estimates of the pre-treatment coefficients $\hat{\mu}_\ell$ for $\ell \in \{-5, -4, -3, -2\}$, with $\ell = -1$ as the reference period. Standard errors (in parentheses) are clustered at the municipality level. Max $|t|$ is the largest absolute t-statistic among the pre-treatment coefficients. Pre-trends are flat for all outcomes except police expenditure (lower spending at $\ell = -5$ and $\ell = -4$) and market value of real estate (limited to pre-2012 data, $N = 9{,}338$). Total jobs from LODES; index crime from FBI UCR Offenses Known (93.5% agency-to-municipality match);

HonestDiD sensitivity plots. Each plot reports the 95% confidence interval for the average post-event Inc-bind effect under the **rambachanMoreCredibleApproach2023** relative-magnitudes restriction $\Delta^{RM}(\bar{M})$. The leftmost "Original" bar is the unrestricted CI; subsequent bars impose $\bar{M} \in \{0, 0.5, 1, 1.5, 2\}$. The breakdown $\bar{M}^*$ (Table 49) is the largest $\bar{M}$ at which the corresponding bar still excludes zero.

Figure 40: HonestDiD sensitivity, block-level log home count (the paper headline).

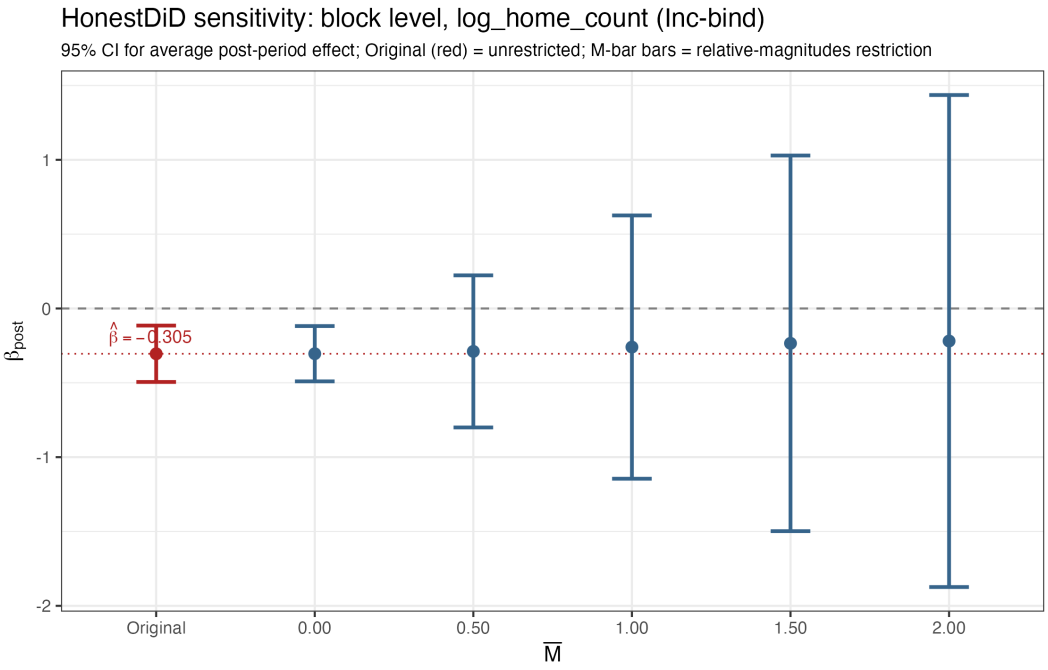


Figure 41: HonestDiD sensitivity, block-level log work count.

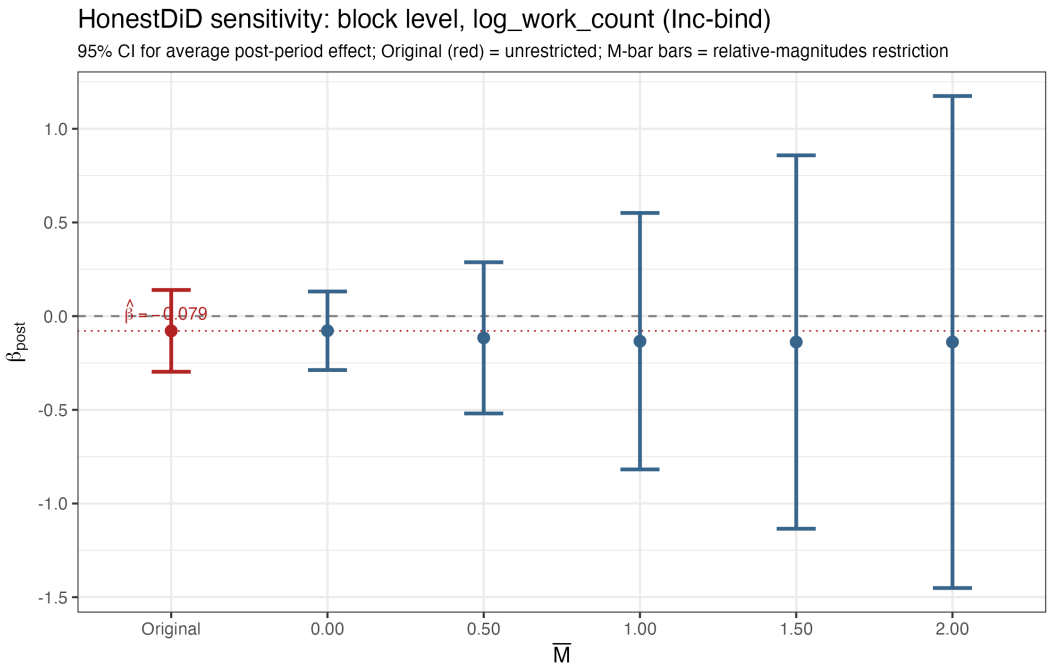
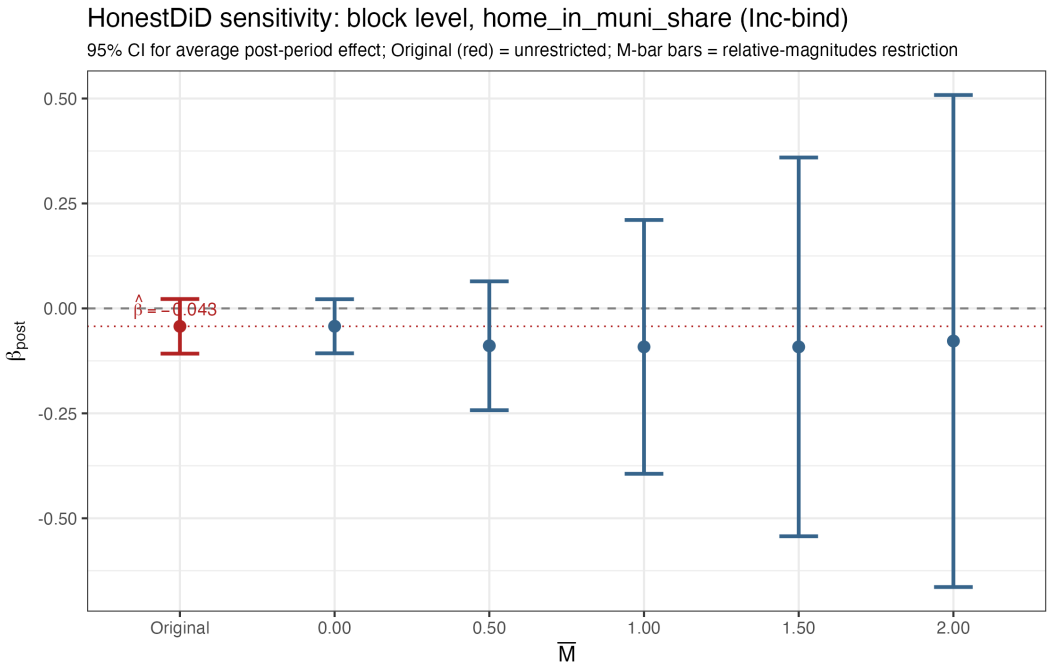


Figure 42: HonestDiD sensitivity, block-level share of residents working in their own municipality.



Block level.

Figure 43: HonestDiD sensitivity, municipality-level log live-and-work (the calibration estimate).

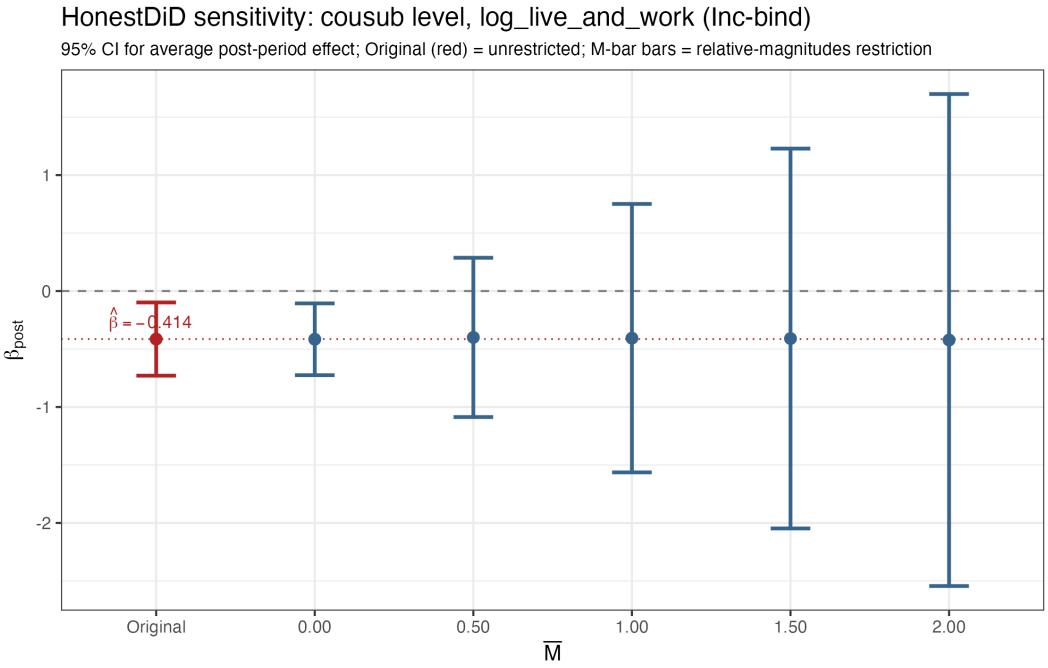


Figure 44: HonestDiD sensitivity, municipality-level log in-commuters.

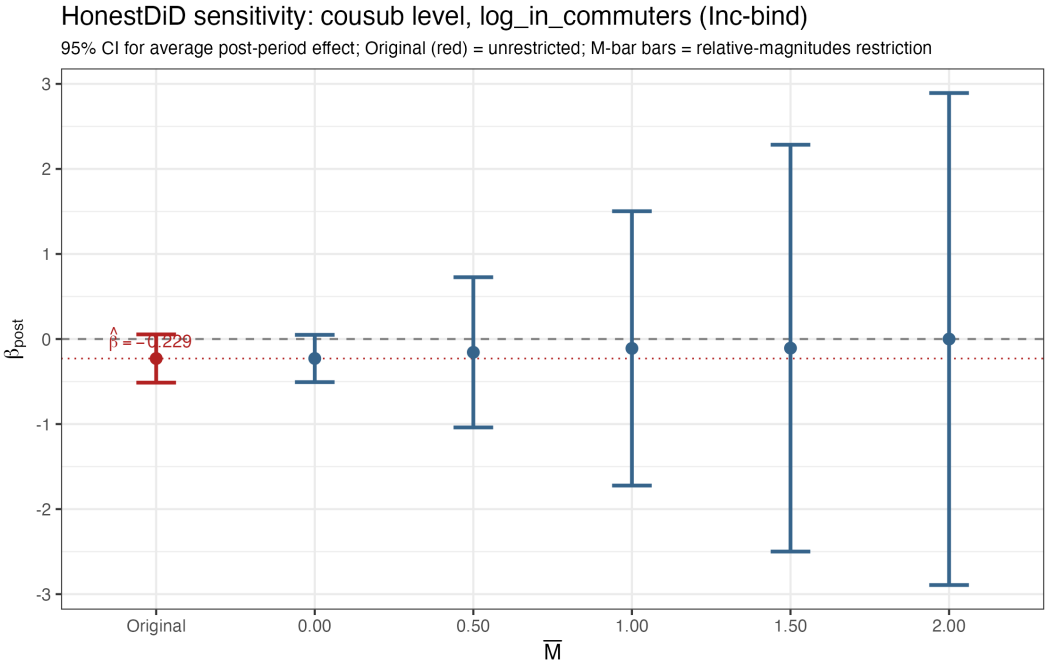


Figure 45: HonestDiD sensitivity, municipality-level log total residents.

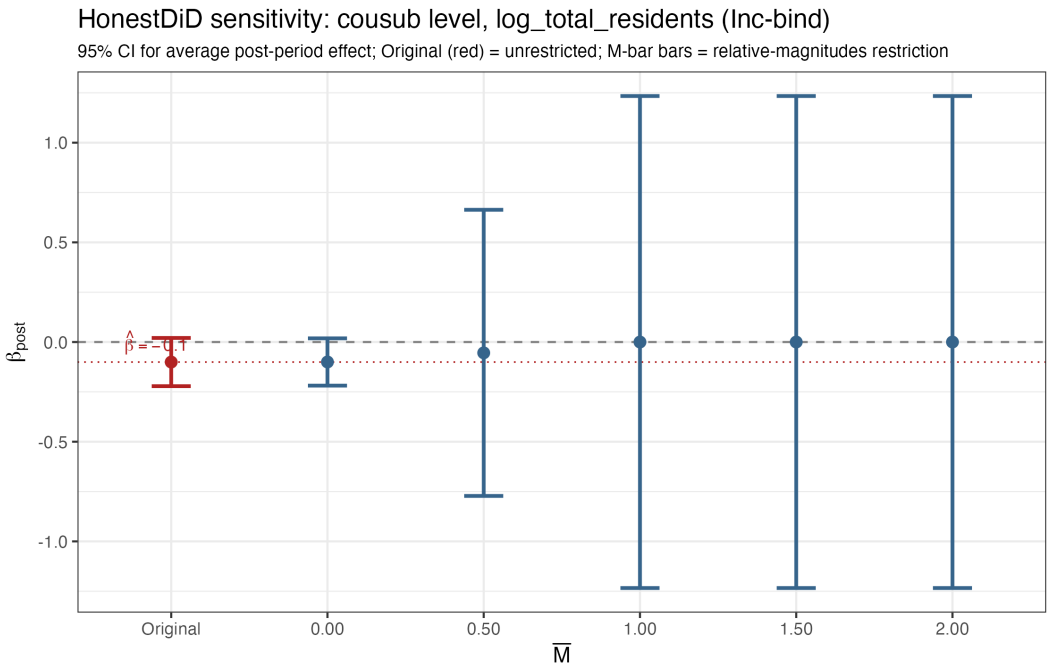
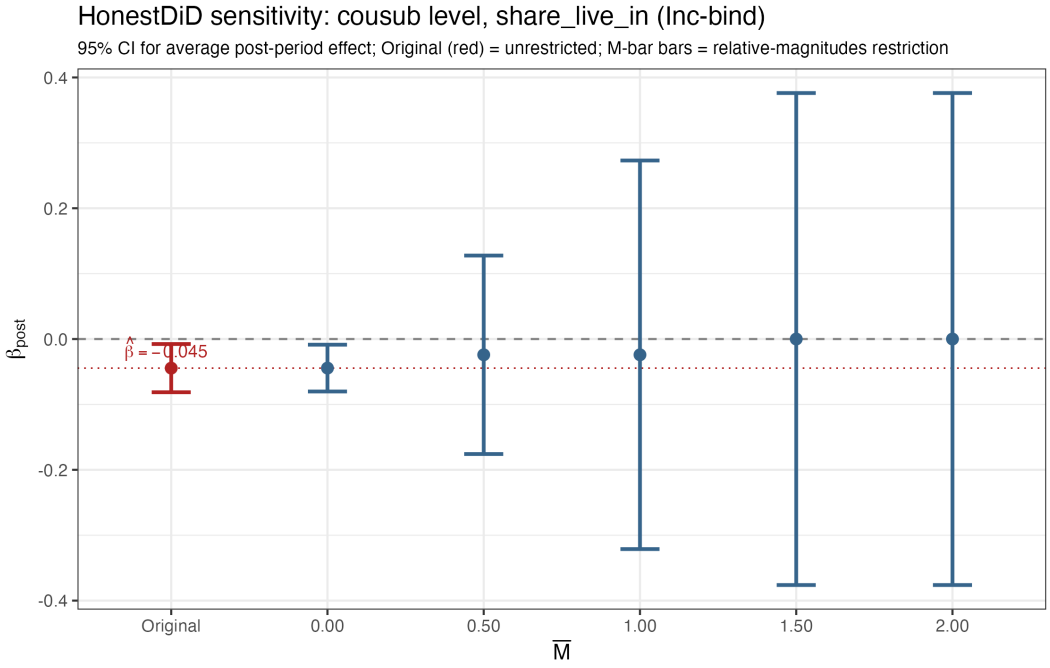


Figure 46: HonestDiD sensitivity, municipality-level share live-in.



Municipality level.

Figure 47: HonestDiD sensitivity, SD-level log live-and-work.

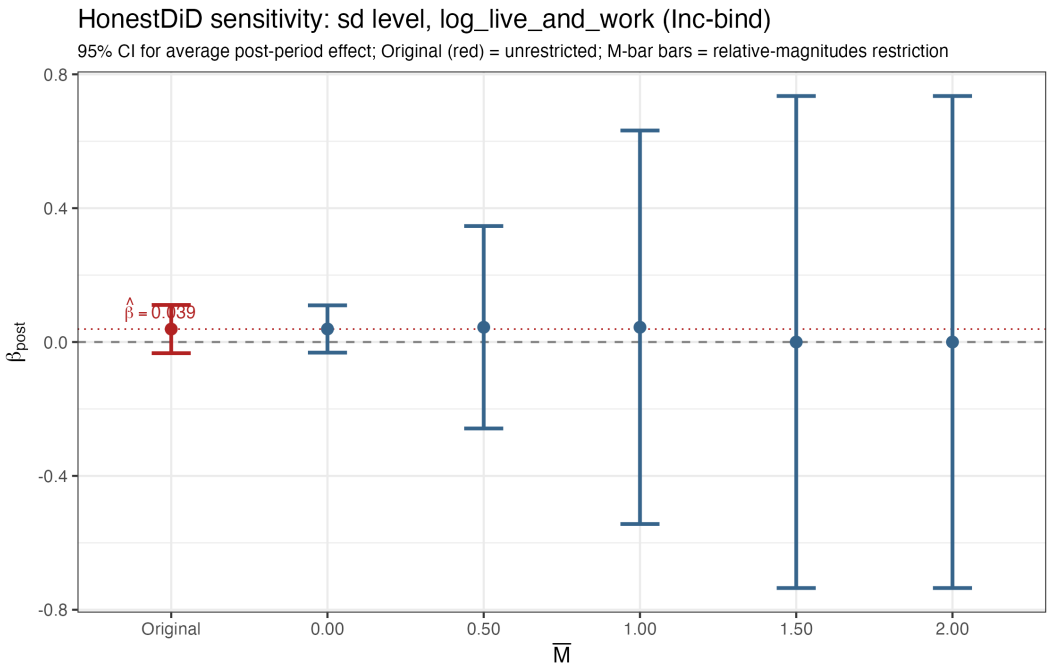


Figure 48: HonestDiD sensitivity, SD-level log in-commuters.

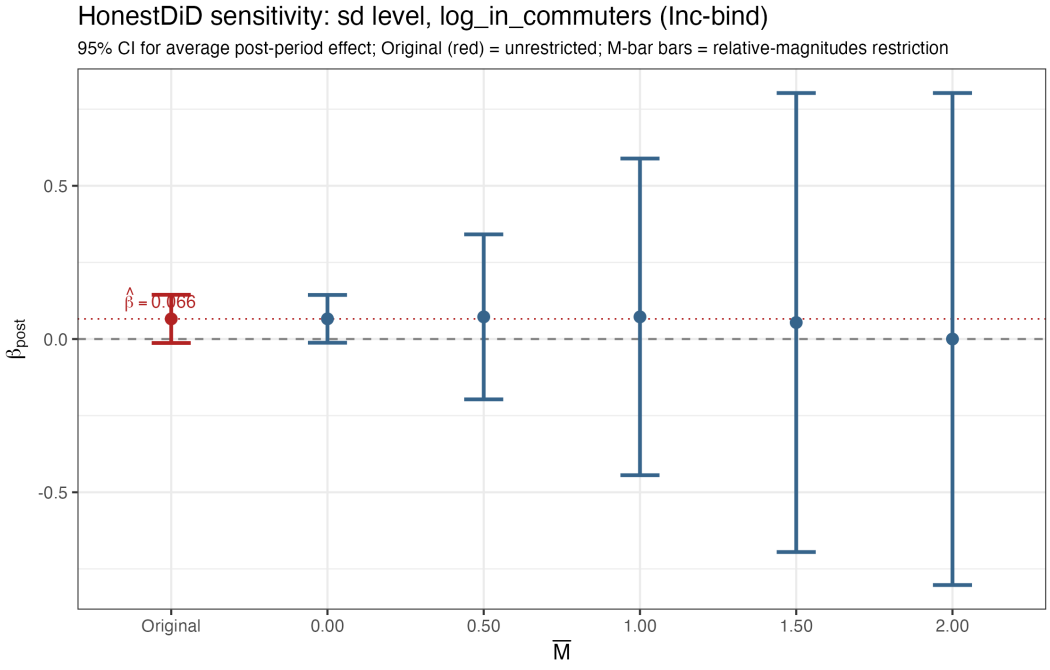


Figure 49: HonestDiD sensitivity, SD-level log total residents.

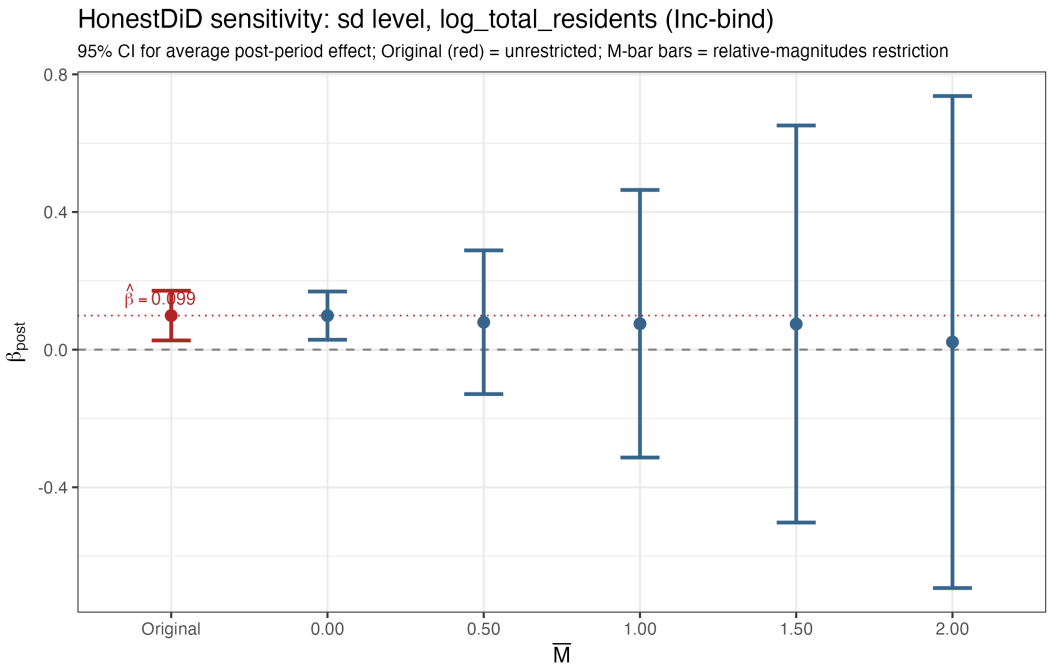
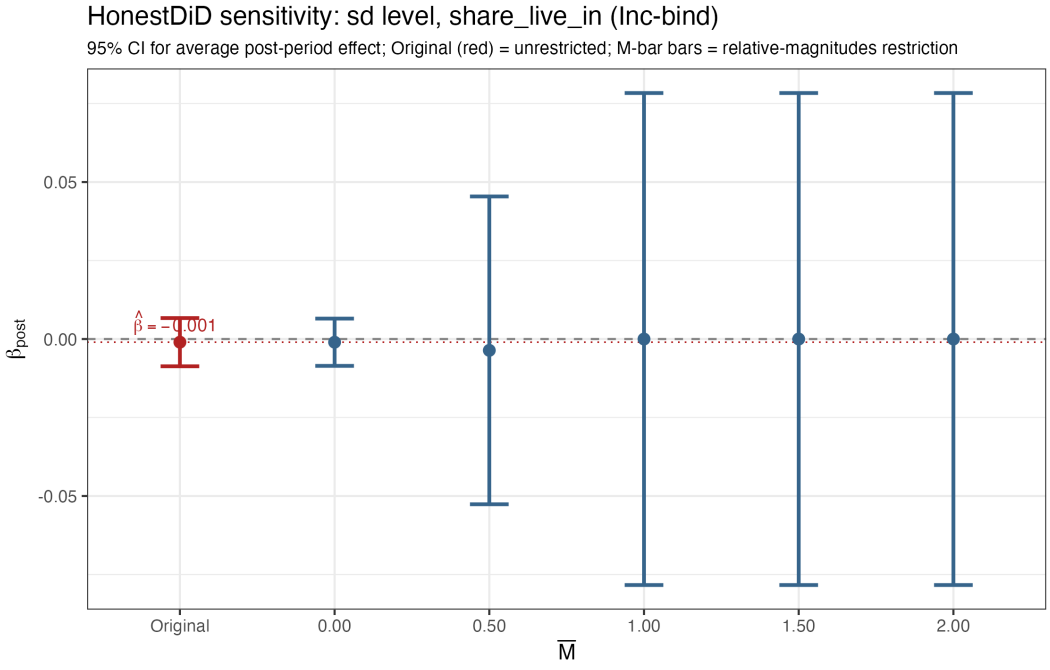


Figure 50: HonestDiD sensitivity, SD-level share live-in.



School-district level.

E Data Construction

E.1 DCED rate scraping

Municipal and school-district EIT rates come from the Pennsylvania Department of Community and Economic Development (DCED) MunStats portal (Pennsylvania Department of Community and Economic Development, 2024a). The scraper (`scrape_pa_dced_eit.py`) downloads annual XLS files from 2002 through 2023 and extracts three rate variables per municipality–year: the municipality resident rate $\tau^{\text{res,muni}}$, the municipality nonresident rate $\tau^{\text{nr,muni}}$, and the school-district rate τ^{SD} (the SD rate is reported on the municipality file because DCED tracks rates at the municipality–SD pair level). Each rate field is parsed as a percentage and converted to a numeric in $[0, 5]$.

Municipality identifiers in DCED use the agency's internal `municipality_id` (six-digit code). The crosswalk to Census geography uses a name-based match against Census COUSUBFP (county subdivision FIPS codes) for the 2020 vintage. Matching proceeds by normalized name (lower-case, stripped of suffixes "borough", "township", "city of", etc.) and county. The match rate is 99.6%; the ∼10 unmatched municipalities are smaller boroughs that either dissolved or merged during the panel and are dropped.

DCED's nonresident-rate coverage is incomplete pre-2012: only 55–60% of municipalities report a nonresident rate before Act 32 made the field a legal requirement; 72–80% report it post-2012. I backfill missing pre-2012 nonresident rates with the earliest observed post-2012 value per municipality, on the assumption that a municipality's nonresident rate is sticky over time. Sample-variant robustness checks (§7.1) compare the headline coefficients under strict drop-NA versus this filled-in nonresident-rate panel and show the binding-state classification is stable.

A parallel scraper (`reparse_dced.R`) recovers school district name and `sd_id` on the DCED rate files, which the original Python scraper drops; this is needed for the NCES–DCED school-district crosswalk used to attach Census F-33 fiscal controls. School-district

name normalization uses a Jaro–Winkler fuzzy match with abbreviation expansion ("CO" → "county", "SD" → "school district") and two manual overrides for districts that changed names entirely during the panel (KNOCH ↔ SOUTH BUTLER CO; RIVER VALLEY ↔ BLAIRSVILLE SALTSBURG). The final NCES `LEAID` → DCED `sd_id` crosswalk covers 500/500 Pennsylvania school districts.

E.2 LODES construction

Residence–workplace flows come from the LEHD Origin–Destination Employment Statistics (LODES) Version 8 (Abowd et al., 2009; U.S. Census Bureau, 2024c), the standard administrative source linking each employed worker's home block to their workplace block at annual frequency. LODES is constructed from state Unemployment Insurance wage records combined with the Quarterly Census of Employment and Wages, with synthetic noise infused at the block level to protect respondent confidentiality; the noise mechanism is calibrated so that aggregation to larger geographies recovers ground truth to within validated bounds at the tract level and above (Graham et al., n.d.).

I download the Pennsylvania LODES OD files for 2002–2019 via the `lehdr` R package (`download_lodes.R`). The full state-year file is ∼200 MB compressed; for 18 years and roughly 5,000 blocks per municipality, the raw block-pair-year panel is roughly 4 GB on disk. I trim 2020–2021 from the analysis sample because COVID-19 disrupted both the underlying employment patterns and the LODES coverage methodology in ways that would conflate the pandemic shock with tax-driven sorting.

LODES reports a "main job" indicator (`JT01`) that I use to restrict to primary-job employment relationships; secondary jobs and federal-government employees (which LODES handles separately) are excluded. For each (block, year) the panel records: the count of workers living in the block (`home_count`), the count of workers working in the block (`work_count`), the count of workers both living and working in the block (and in the block's own municipality, after aggregation), and the directional commuter flows used in the municipality- and

SD-level aggregations.

E.3 Block-to-municipality crosswalk

Pennsylvania Census blocks nest in county subdivisions (cousubs) by construction: each block has a unique `COUSUBFP` via its `GEOID20` (U.S. Census Bureau, 2020). I build the block → cousub crosswalk by joining the 2020 TIGER/Line block shapefile against the 2020 cousub shapefile on `STATEFP + COUNTYFP + COUSUBFP`, yielding `pa_block_xwalk.rds` with exactly one cousub per block. This is the assignment used throughout the paper to attach the LODES block-level OD flows to the municipality-level DCED rate panel.

Earlier drafts of this paper attempted a tract-centroid join from LODES tracts to cousubs (joining each tract centroid to the cousub containing it) because LODES Version 7 reported tract-level OD files. That centroid-based assignment silently dropped 42% of PA cousubs whose tract centroids fell outside the cousub polygon (typical of small boroughs surrounded by larger townships). The switch to block-level LODES (Version 8, 2022) plus the exact TIGER block-to-cousub nesting eliminated this artifact and is the basis for the block-level identification on which the paper rests.

E.4 PSD mapping

Pennsylvania assigns each address a six-digit Political Subdivision (PSD) code that uniquely identifies the (municipality, school district) pair the address belongs to. Workers report the PSD code on the Residency Certification Form, and employers use it to determine withholding obligations after Act 32 of 2008 (effective 2012). The PSD code is the operational identifier that ties a worker's home address to the municipality and SD EIT rates that compose the home stack $\tau^{\text{home}} = \tau^{\text{res,muni}} + \tau^{\text{SD}}$ in Equation (1).

I do not observe individual workers' PSD codes; LODES is anonymized at the block level. I instead infer the (municipality, SD) pair for each block by combining the block-to-cousub crosswalk (§E.3) with a primary-SD assignment per municipality. The primary

SD per municipality is the modal `sd_id` across years from the raw DCED rate panel (most municipalities report rates to a single SD across the full panel; 97.6% of municipality-years are single-SD). For the 2.4% of municipalities that span multiple SDs in DCED records, the cousub-level model uses $\max(\tau^{SD})$ across the municipality's SDs (conservative, slightly biases the SD-rate coefficient toward zero); the block-level model assigns each block its municipality's primary SD.

E.5 Fiscal controls

Fiscal controls are share-based and time-varying within unit. The exact control set varies by estimating unit, matching the unit's fiscal authority.

Muni-level Annual Financial Reports. For municipality-level regressions and the municipality rate-setter in block-level regressions, the controls are six expenditure and revenue shares constructed from the PA DCED Annual Financial Report database (Pennsylvania Department of Community and Economic Development, 2024b): `share_police` (police expenditure / total expenditure), `share_streets`, `share_debt_svc`, `share_eit` (EIT revenue / total revenue), `share_re` (real-estate tax revenue / total revenue), and `share_intergov` (intergovernmental aid / total revenue). The DCED AFR panel covers 2002–2023; I trim to 2002–2019 to exclude ARPA-distorted years. Coverage is essentially 100% of PA municipalities post-2006, falling to ∼95% in 2002–2005 due to sparser AFR filings; missing-year imputation is not applied (the ES regressions drop municipality-years with missing controls).

Census F-33 school-district finance. For SD-level regressions and the SD rate-setter in block-level regressions, the controls are eight share variables from the Census Annual Survey of School System Finances (F-33) (U.S. Census Bureau, 2024a): `share_state_rev_sd`, `share_local_rev_sd`, `share_federal_rev_sd`, `share_property_tax_sd` on the revenue side, and `share_instructional_sd`, `share_support_sd`, `share_capital_sd`, `share_debt_int_sd` on the expenditure side. The F-33 panel covers 10,790 Pennsylvania SD–year observations across 500 SDs and 2002–2019. The cross-reference between NCES F-33 `LEAID` and DCED

`sd_id` is described in §E.1.

Block-level regressions use both sets of controls simultaneously, since each block is exposed to both its municipality's fiscal posture (via the resident-rate channel) and its school district's fiscal posture (via the SD-rate channel).

The block-level controls also include a within-SD aggregate of municipality rates (`mean_muni_rate`) when the SD rate-setter portion of the regression is being estimated; this absorbs SD-level variation in the joint municipality resident rate that the SD itself does not control.

E.6 Sample restrictions

The estimation sample is 2002–2019. LODES Version 8 covers 2002–2021; I drop 2020–2021 because COVID-19 disrupted both commuting patterns (LODES coverage of stay-at-home workers is poorly defined) and municipal finances (ARPA transfers distort fiscal-share controls 2020–2023).

I impose four additional sample filters at estimation time. First, a minimum-employment threshold of 50 total jobs per municipality–year is applied to municipality-level regressions because share outcomes are dominated by sampling noise at smaller scales (a single worker moving changes `share_live_in` by $1/N$). Sensitivity to this threshold is documented in §7.1; coefficients are stable across thresholds from 10 to 500. Second, I apply rate outlier filters: DCED reports the Local Services Tax (a flat dollar amount) in the same field as the EIT rate (a percentage) in some pre-2012 years, so any reported rate exceeding 5% is treated as a data-entry error and dropped (the highest standard PA EIT is approximately 1%; only Philadelphia under the Sterling Act exceeds 3%, and Philadelphia is handled separately because it does not participate in the Act 32 framework described in §2.1). Third, for each municipality–year I assign the rate in effect on April 1, aligning with LODES's reference date so that rate-on-paycheck and OD-flow records refer to the same point in time. Rates that take effect mid-year are credited to the year following the effective date. Fourth, the binding-state cell decomposition requires a non-missing nonresident rate to classify the binding state at

$c-1$ and c; municipality–years with missing nonresident rates that have ambiguous binding states are categorized as "Uncertain" and dropped from the analysis (pooling them with NC would contaminate the control trend with treated-but-unclassified units; reporting them in their own cell would mix true binding states with measurement gaps; see §5).

Cluster definitions follow the rate-setting unit: municipality-level regressions cluster at the municipality; SD-level regressions cluster at the SD; block-level regressions two-way cluster at the municipality and SD. The cluster-robust variance is the default `fixest::feols` sandwich estimator with finite-sample corrections.

E.7 Additional validation

This subsection collects three additional validation checks for the data pipeline: pre-treatment balance between rate-changing and never-changing municipalities, the geographic match rates that quantify how much of the universe each step of the crosswalk preserves, and a brief sample-composition diagnostic.

Pre-treatment balance. Table 51 reports normalized differences in fiscal, demographic, and supplementary observables between municipalities that raise the EIT during the panel and never-changed municipalities, computed on pre-treatment years only. Normalized differences are below the conventional 0.15 threshold for all fiscal variables. Treated municipalities are larger in levels (median population 14K versus 3.7K) but the difference is small in normalized terms once the fiscal-share specification is used, which divides out the unit's overall scale. Market value of real estate (normalized diff. 0.211) and building permits (normalized diff. 0.232) exceed the threshold, reflecting that treated municipalities are disproportionately larger jurisdictions with more real estate value and higher absolute permit activity; both are level controls in robustness specs, not in the headline.

Table 51: Pre-Treatment Balance: Treated vs. Never-Treated PA Municipalities

Variable	Treated Mean	Control Mean	Norm. Diff.	N (T / C)
Population	14,174	3,711	0.139	1,265 / 39,688
Total Revenue ($)	25.7M	2.9M	0.077	1,267 / 39,692
EIT Revenue ($)	4.4M	438K	0.062	
RE Tax Revenue ($)	2.9M	538K	0.129	
Police Expenditure ($)	3.1M	424K	0.110	
Streets Expenditure ($)	1.5M	496K	0.121	
Total Debt ($)	28.6M	3.6M	0.089	
Market Value RE ($)	8.2M	2.1M	0.211	523 / 8,774
Total Index Crime	985	133	0.124	816 / 18,322
Building Permits	19.3	7.7	0.232	644 / 25,066
EIT Rate (%)	0.50	0.50	0.055	

Notes: Treated = municipalities that raised EIT during the sample period (pre-treatment years only). Control = municipalities that never raised EIT. Normalized differences < 0.15 for all fiscal variables except market value of real estate (0.211, limited pre-2012 coverage) and building permits (0.232, treated municipalities are larger and issue more permits in levels). Sources: PA DCED Annual Financial Reports, FBI UCR Offenses Known (93.5% agency match rate), Census Building Permits Survey (93.3% MCD match rate).

Crosswalk match rates. Each step of the crosswalk pipeline is reported with its observed match rate, allowing the reader to track how many units drop at each join.

- DCED `municipality_id` → Census 2020 `COUSUBFP`: 99.6% match (name-based, county-keyed, normalized-suffix). Unmatched municipalities are smaller boroughs that dissolved or merged during 2002–2019; they are dropped, not imputed.
- NCES F-33 `LEAID` → DCED `sd_id`: 500/500 PA school districts matched (Jaro–Winkler with abbreviation expansion + two manual overrides for fully-renamed districts). This is the join that attaches Census F-33 SD finance to the DCED rate panel.

- TIGER 2020 Census block → Census 2020 cousub: exact spatial nesting; each block has a unique `COUSUBFP` by construction.
- FBI UCR agency → DCED municipality (for crime controls in pre-treatment balance): 93.5% match via Jacob Kaplan's concatenated UCR file (Kaplan, 2024) and `fips_place_code` keying.
- Census Building Permits Survey place → DCED municipality: 93.3% match via `fips_mcd_code`.

Sample-composition diagnostic. Of the 2,503 DCED-listed Pennsylvania municipalities, the panel classifies ∼2,400 as never-changed (NC) on the municipality resident rate side. The treated tail consists of 40 Inc-bind, 58 Inc-becomes, and 26 Inc-nobind municipalities (the remainder are Dec cells or "Uncertain"; see Section 5). The treated tail spans eight cities, ten boroughs, and twenty-two townships in the Inc-bind cell; magnitudes of the headline coefficient are concentrated in the cities (see Companion Appendix), consistent with high-eventual-dose post-industrial municipalities being the identifying variation.

F Additional Institutional Details

F.1 Piecewise tax examples

This subsection walks through worked numerical examples of the piecewise EIT formula in Equation (1) for the three binding-state cells the paper exploits. Each example uses a worker who lives in municipality j and works in municipality i, with the standard 1% combined cap split 0.5% municipality resident plus 0.5% school-district.

Example 1: Inc-bind (the headline cell). Pre-event, worker lives in a municipality with $\tau_j^{\text{res,muni}} = 0.6\%$, $\tau_j^{\text{SD}} = 0.5\%$, and works in a municipality with $\tau_i^{\text{nr,muni}} = 0.5\%$. The home stack is $0.6 + 0.5 = 1.1\%$; the work nonres rate is 0.5%; the home stack binds. The total bill is $\max(1.1\%, 0.5\%) = 1.1\%$, split as 0.6% to the home municipality and SD (0.6%

home stack − 0% residual to work, actually keep it simple: the home municipality and SD get $\tau^{\text{home}} - \tau^{\text{nr}} = 1.1\% - 0.5\% = 0.6\%$; the work municipality gets 0.5%).

Post-event, the municipality raises its resident rate to $\tau_j^{\text{res,muni}} = 0.9\%$. The home stack becomes 1.4%, still binding (1.4% > 0.5% work nonres). The total bill rises to 1.4%, an increase of 0.3 pp; the home jurisdictions get $1.4 - 0.5 = 0.9\%$. A residential move to a lower-rate municipality j' with $\tau_{j'}^{\text{res,muni}} = 0.5\%$ would lower the home stack to $0.5 + 0.5 = 1.0\%$, still binding, with a new total bill of 1.0%. The worker's savings from moving: $1.4\% - 1.0\% = 0.4\%$ of earned income. This is the cell where the institutional rule predicts the strongest sorting response.

Example 2: Inc-nobind (the within-design placebo). Pre-event, worker lives in a municipality with $\tau_j^{\text{res,muni}} = 0.5\%$, $\tau_j^{\text{SD}} = 0.5\%$, and works in a Home Rule city with $\tau_i^{\text{nr,muni}} = 1.5\%$. The home stack is 1.0%; the work nonres rate is 1.5%; the work nonres binds. Total bill is $\max(1.0\%, 1.5\%) = 1.5\%$, with all of it going to the work municipality. The home municipality and SD collect zero.

Post-event, the home municipality raises its resident rate to 0.7%. The home stack becomes 1.2%, still below the 1.5% work nonres rate; the work nonres still binds. Total bill remains 1.5%. The worker's bill is unchanged by the home-municipality rate hike, so the institutional rule predicts no sorting response. The worker has no incentive to move home in response to this rate change. This is the placebo cell.

Example 3: Inc-becomes (the intermediate case). Pre-event setup identical to Example 2 in some respects: home stack 1.0%, work nonres 1.2%, work nonres binds, total bill 1.2%.

Post-event, the home municipality raises its resident rate to 1.0%, pushing the home stack to 1.5%. The home stack now binds (1.5% > 1.2% work nonres). Total bill rises to 1.5%, an increase of 0.3 pp. But unlike Inc-bind, only the portion above the work nonres rate ($1.5 - 1.2 = 0.3\%$) accrues to the home jurisdictions; the work municipality still collects

its 1.2%. The worker's bill rises by less than the full statutory $\Delta\tau = 0.5\%$ (it rises by only 0.3%, the amount the home stack exceeds the work nonres rate), so the institutional rule predicts an intermediate sorting response.

F.2 Additional institutional details

Three further institutional details affect interpretation but were deferred from the main text for space.

Act 1 of 2006 (school-district EIT authorization). The Taxpayer Relief Act of 2006 authorized school districts to levy or raise an EIT specifically to fund property-tax reduction, above the 0.5% statutory default. This generated a wave of SD rate increases over 2006–2011 that the paper exploits as independent rate variation: the cross-panel correlation between municipality and SD EIT rates is -0.12, so municipality and SD rate changes can be identified separately in a joint regression.

Act 47 distressed municipalities. Approximately 15 Pennsylvania municipalities have been designated "financially distressed" under Act 47 of 1987 (Pennsylvania General Assembly, 1987) and granted authority to exceed the standard 0.5% municipality EIT cap (typically up to 1.5% or higher). These include Pittsburgh, Wilkes-Barre, Hazleton, Johnstown, and others in the Inc-bind cell, and they account for the right tail of the resident-rate distribution shown in Figure 1. Their inclusion in the treated cohort is responsible for the large pre-2012 cumulative doses observed in some Inc-bind cohorts; sensitivity to dropping Act 47 municipalities is documented in the heterogeneity analysis in the companion appendix.

Home Rule municipalities. Pennsylvania's Home Rule Charter and Optional Plans Law (53 Pa. C.S. § 2961) authorizes municipalities to adopt a Home Rule Charter and exceed the standard 0.5% municipality-resident-rate cap without an Act 47 designation. About 80 PA municipalities have a Home Rule Charter, of which around 20 have used the authority to

raise their municipality EIT above the cap. These are typically older industrial boroughs (Bethlehem, Allentown, Reading, Lancaster, York, Erie).

F.3 Philadelphia treatment

Philadelphia operates under the Sterling Act of 1932 (Act of August 5, 1932, P.L. 45), which predates Act 511 and creates a separate EIT framework. Three features differ from the rest of Pennsylvania. First, Philadelphia residents pay the city resident rate ($\sim$3.9%) regardless of where they work, with no offset for work-location nonresident rates. Second, Philadelphia nonresidents working in Philadelphia pay the Philadelphia nonresident rate ($\sim$3.5%), with no residual flowing to their home jurisdiction's EIT. Third, the piecewise max-rule of Act 32 does not apply to Philadelphia-resident or Philadelphia-employed workers.

The main estimates retain Philadelphia in the sample because (a) its $\sim$ 1.5M residents contribute meaningful identifying variation in the residential-sorting outcomes; (b) the cousub crosswalk maps Philadelphia to a single COUSUBFP cleanly; and (c) the binding-state classification still applies internally (Philadelphia is in NC throughout the panel because its rates change only slightly and never trigger an Inc-bind / Inc-becomes / Inc-nobind transition). Sensitivity to dropping Philadelphia entirely is reported in the companion appendix and shows the headline coefficients are essentially unchanged.